

Curriculum Vitae

Education/Awards

Studies at KIT/Univie 2004–2013:

- DIPLOMA MATHEMATICS (1.0)
- DIPLOMA INFORMATICS (1.1)
- PHD SUMMA CUM LAUDE

Honors and Awards:

- BEST STUDENT OF DEPARTMENT OF INFORMATICS (KIT)
- UNISERV DISSERTATION AWARD
- KIT DOCTORAL AWARD IN THE AREA OF COMPETENCE
“INFORMATION, COMMUNICATION AND ORGANIZATION”
- HEINZ BILLING PRIZE
- BEST PAPER AWARD IPDPS’18
- SPP ALGORITHMS FOR BIG DATA: BEST PAPER AWARD
- PACE IMPLEMENTATION CHALLENGE

Curriculum Vitae

Professional Experience

- January 2017 – present UNIVERSITY VIENNA.
Research associate in group of Monika Henzinger
Leading Algorithm Engineering Subgroup
- March 2010 – December 2017 KARLSRUHE INSTITUTE OF TECHNOLOGY.
Research associate in group of Peter Sanders
Leading Parallel Algorithms & Partitioning Subgroup
- since 2013: TU VIENNA visiting scientist/professor
in the group of Jesper Larsson Träff

Curriculum Vitae

Third Party Funding / Professional Service

January 2017 PARTITIONING LARGE GRAPHS
Co-PIs: Dorothea Wagner, Peter Sanders
DFG application. Funded: 580T Euro.

Sept 2017 PROCESS MAPPING
Co-PI: Jesper Larsson Träff
FWF application. Funded: 312T Euro.

Sept 2019 KERNELIZATION
FWF Start application submitted

PC Member:

ESA'17, IPDPS'18, IPDPS'19, IPEC'19, ISC'19, ALENEX'20, CSC'20 [...]

Reviewer:

JEA, ALX, SPAA, IPDPS, DIMACS, ESA, TPDS, JPDC, SODA, TVCG,
TOPC, SISC [...]

Curriculum Vitae

Third Party Funding / Professional Service

January 2017

PARTITIONING LARGE GRAPHS

Co-PIs: Dorothea Wagner, Peter Sanders

DFG application. Funded: 580T Euro.

Sept 2017

PROCESS MAPPING

Co-PI: Jesper Larsson Träff

FWF application. Funded: 312T Euro.

Sept 2019

KERNELIZATION

FWF Start application submitted

Co-organizer Shonan Meeting

“Parameterized algorithms & data reduction: Theory meets practice”

Guest Editor Special Issue of Algorithms

“Graph Partitioning: Theory, Engineering and Applications”.

Invited Speaker Finse Winter School on “Parameterized Complexity”

Curriculum Vitae

Teaching

Supervision:

2010-19 ≥ 30 thesis (BA/MA), 3 currently running

Lecturing:

2014-17	5xGRAPH PARTITIONING AND CLUSTERING, KIT
2016	ALGORITHMS II, KIT
2017-18	ADVANCED ALGORITHMS, UNI VIE
2018	ALGORITHMS AND DATA STRUCTURES, UNI VIE
2018	ALGORITHMS AND DATA STRUCTURES II, UNI VIE
2019	NUMERICAL ALGORITHMS, UNI VIE
2019	ALGORITHMS AND DATA STRUCTURES II, UNI VIE
2019	DISTRIBUTED AND PARALLEL ALGORITHMS, UNI VIE
2019/20	NUMERICAL ALGORITHMS, UNI VIE
2019/20	ADVANCED TOPICS IN ALGORITHMS, UNI VIE

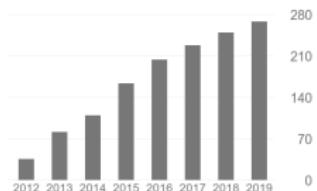
$\Sigma \rightarrow 30$ SWS

Metrics



- six years after PhD
- open source: Ka{HIP, Gen, MIS, Draw, LP}, VieM, VieClus, VieCut
- 58 co-authors
- 57 peer-reviewed publications (+4 currently in submission)
 - Algorithm Engineering
(12x ALX, 1x CSC, 1x ESA, 12x SEA, 3x GECCO, 1x SoCS)
 - Parallel Processing (4x IPDPS, 3x Euro-Par, 1x SPAA)
 - 42 publications as PostDoc
 - 4 book chapters
 - 10 journal papers (JEA, JoH, TPDS, TVCG, JPDC)

	All	Since 2014
Citations	1394	1238
h-index	17	16
i10-index	22	20



Research Overview

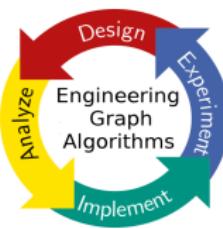
Multilevel Algorithms

Evolutionary Computation

Parallel Programming

Kernelization

shared-, distributed-, external-, internal memory



Algorithms for

graph partitioning
graph clustering
graph generation

process mapping
minimum cuts
independent sets

longest paths
graph drawing
matching

node separators
....

Open Source

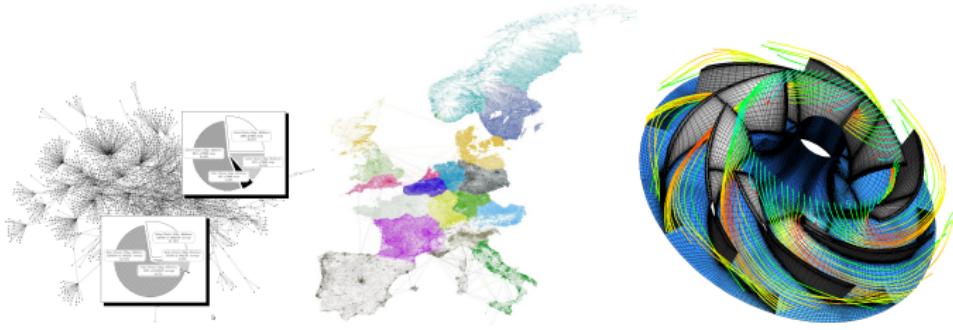
Applications

territory design
large scale simulations
quantum annealing

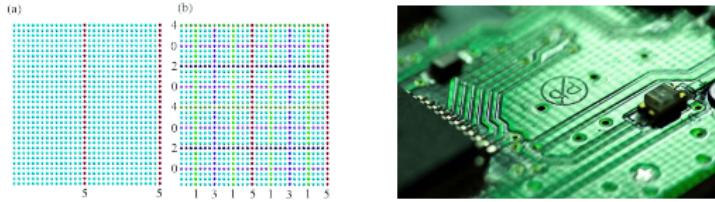
route planning
distributed system design
nuclei segmentation

multiprocessor scheduling
high-throughput DNA sequencing
....

Applications

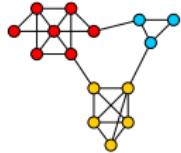


$$\mathbf{R}^{n \times n} \ni Ax = b \in \mathbf{R}^n$$



Highlights

Graph Clustering

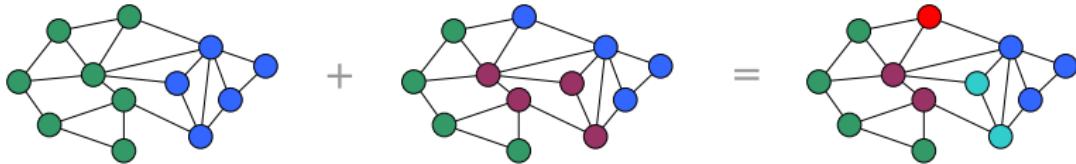


Problem:

- partition graph into tightly connected groups
- clustering paradigm: **internally dense** and **externally sparse**
- quality measure: **modularity**

Recombination Mechanism

- borrowed from ML: ensemble learning → **maximum overlap**



→ results in **overlap clustering**, contract + recluster

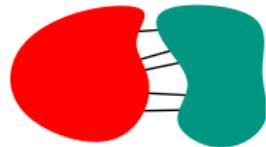
→ + coarse-grained parallelization

**Outperform all DIMACS Challenge results
with ONE solver and in less time**

Minimum Cuts

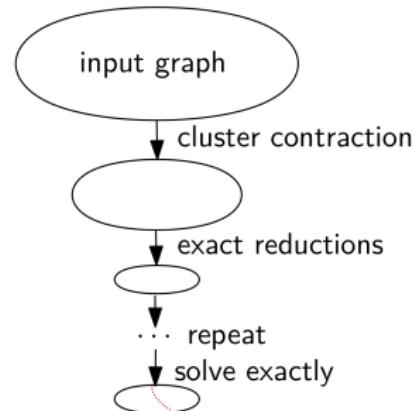
Cut: A **cut** in a multigraph is a partition of $V = C \cup \bar{C}$

Problem: size of minimum cut in G ?



Contributions:

- (in)exact reductions + solve kernel to optimality
→ linear running time, but potentially suboptimal cuts
- NO guarantee, but experiments say likely opt
- reductions depend on bound $\hat{\lambda}$
 - ~~ use $\hat{\lambda}$ in exact (parallel) NOI algorithm
 - ~~ **fastest exact minimum cut algorithm**
- \approx order of magnitude



Graph Generation

Ideas:

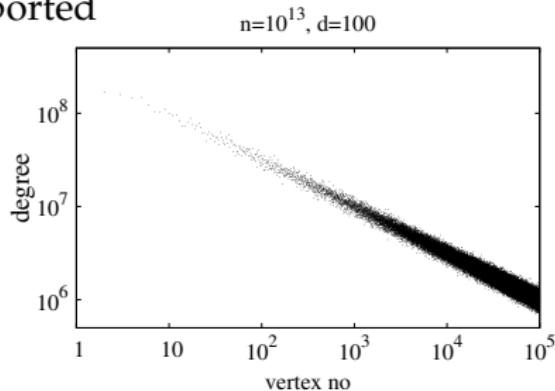
- replace randomness by pseudo-randomness, e.g. hashing
- replace array accesses by recomputation
→ embarrassingly parallel algorithm without communication

Results:

- generate Peta-edge BA graph in < one hour (10^{15} edges, 16K cores)
- 20 000 times larger than previously reported
- 16x faster than RMAT ($m = 50 \cdot 10^9$)
- communication efficient algorithms

more work:

- Barabasi-Albert
- more models:
ER, $G(n,m)$, $G(n,p)$,
RHG, RGG, DEL



Graph Drawing

$$G = (V, E, d)$$
$$d : E \rightarrow \mathbf{R}$$
$$\Rightarrow$$

Maximal Entropy Stress Model:

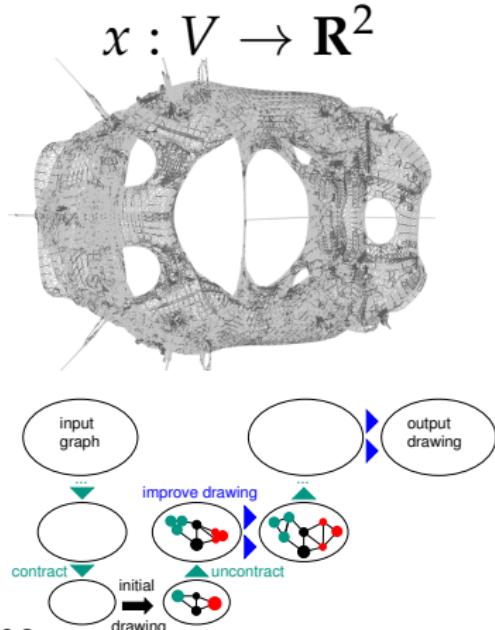
$$\max H(x) := \sum_{\{u,v\} \notin E} \ln ||x_u - x_v||$$

subject to $||x_u - x_v|| = d_{uv}, \{u, v\} \in E$

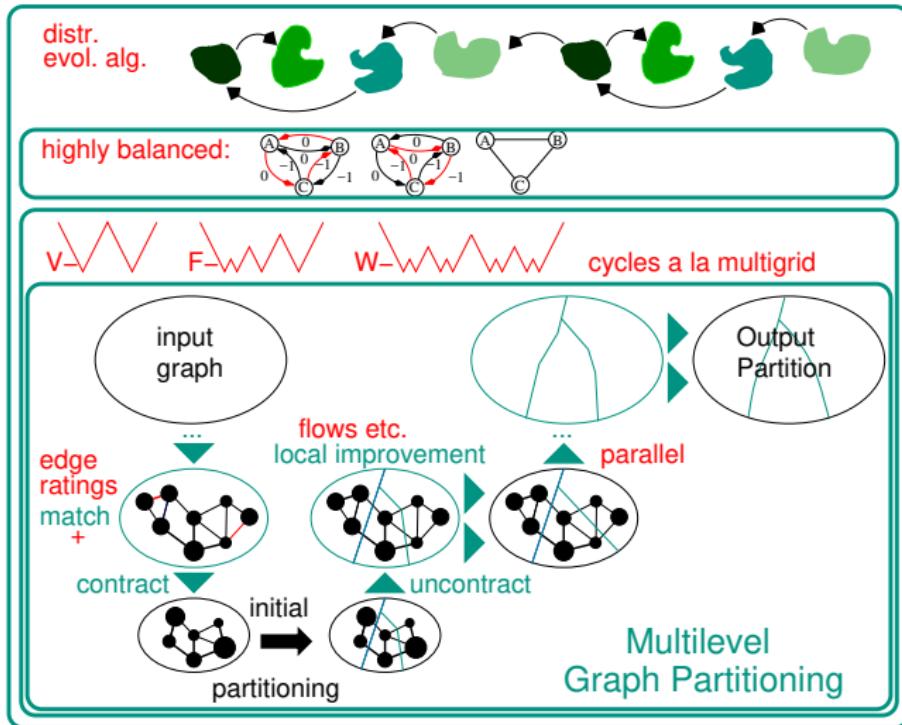
Contributions:

- multilevel integration of iterative scheme
- approximate long-range forces, by cluster contraction
- employ parallelism

Draw graphs with millions of vertices in seconds.



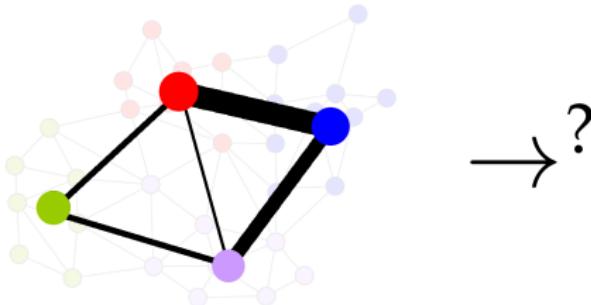
Karlsruhe High Quality Partitioning



Karlsruhe High Quality Partitioning

- 
- pre 2014:
 - multi-level, localized and flow-based local search
 - perfectly balanced,
 - distributed evolutionary algorithms
 - specialized algorithms for road networks
 - social networks
 - distributed memory
 - semi-external memory
 - hypergraph partitioning
 - node separators
 - DAG partitioning
 - shared-memory parallel
 - ILP-based

Process Mapping

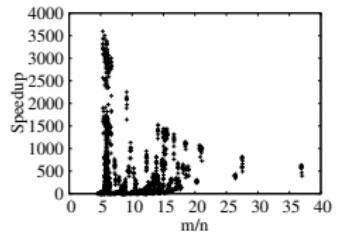


Exploiting Assumptions:

- communication matrix \mathcal{C} is **sparse** → graph $G_{\mathcal{C}}$
- compute system is **hierarchically structured**

Yields:

- three orders of magnitude faster local search
- 50% improved mapping quality over state-of-the-art
→ resulted in a FWF project



Support Vector Machines

- **Binary Classification Problem:**

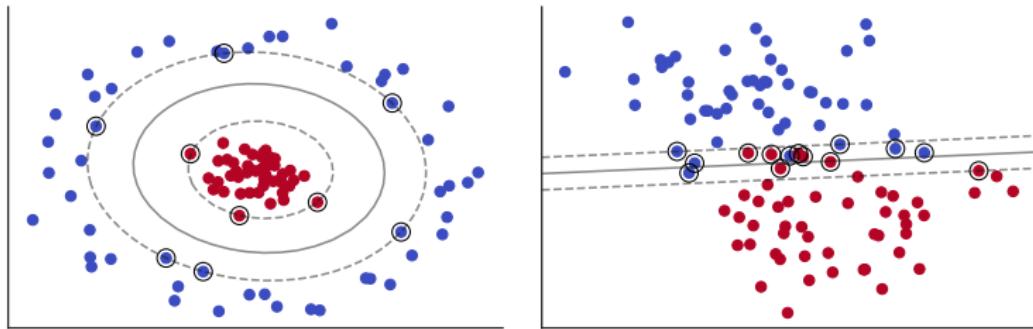
Train classifier on n labeled data points ($x_i \in \mathbf{R}^d, y_i \in \{-1, 1\}$)

Goal: Assign label y_{n+1} to new data points x_{n+1}

- **SVM Optimization Problem:**

$$\text{min.} \quad \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi_i$$

$$\text{s.t.} \quad y_i(w \cdot \phi(x_i) - b) \geq 1 - \xi_i, \quad \xi_i \geq 0.$$



Support Vector Machines

■ Binary Classification Problem:

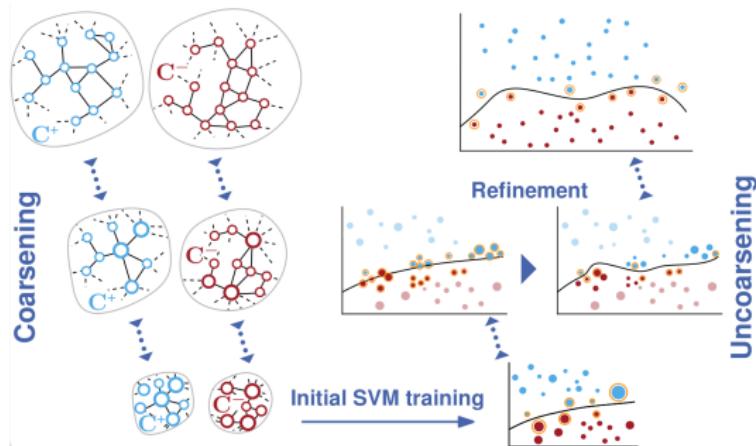
Train classifier on n labeled data points ($x_i \in \mathbf{R}^d, y_i \in \{-1, 1\}$)

Goal: Assign label y_{n+1} to new data points x_{n+1}

■ SVM Optimization Problem:

$$\text{min. } \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi_i$$

$$\text{s.t. } y_i(w \cdot \phi(x_i) - b) \geq 1 - \xi_i, \quad \xi_i \geq 0.$$



Support Vector Machines

■ Binary Classification Problem:

Train classifier on n labeled data points ($x_i \in \mathbf{R}^d, y_i \in \{-1, 1\}$)

Goal: Assign label y_{n+1} to new data points x_{n+1}

■ SVM Optimization Problem:

$$\begin{aligned} \text{min. } & \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi_i \\ \text{s.t. } & y_i(w \cdot \phi(x_i) - b) \geq 1 - \xi_i, \quad \xi_i \geq 0. \end{aligned}$$

multilevel SVM using label propagation
→ comparable classification quality
→ up to two orders of magnitude faster training

currently working on parallelization

Concrete Example

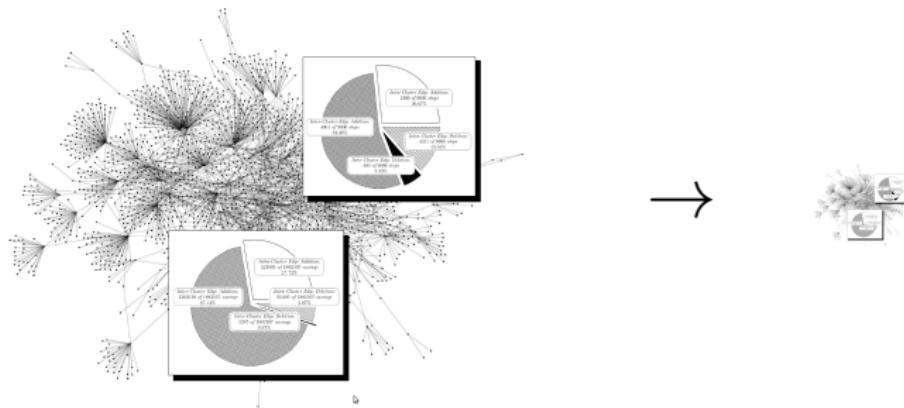
Applied Kernelization for Independent Sets

Kernelization

General Idea

Reductions:

rules to decrease graph size, while maintaining optimality



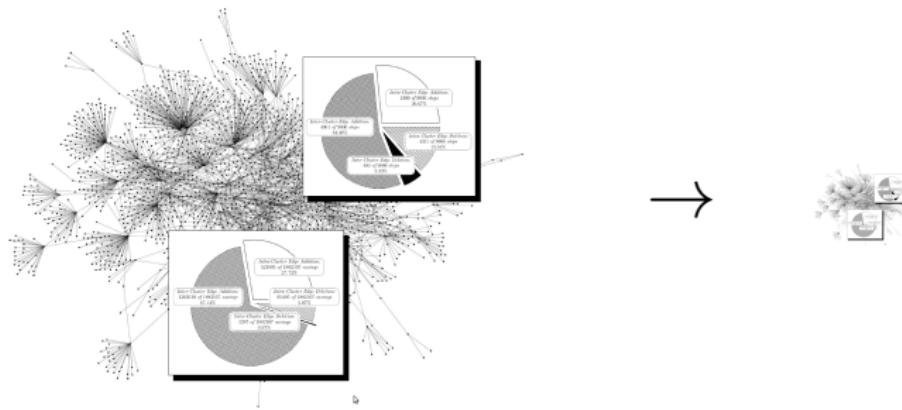
solve problem on **problem kernel**
→ obtain solution on input graph

Kernelization

General Idea

Reductions:

rules to decrease graph size, while maintaining optimality



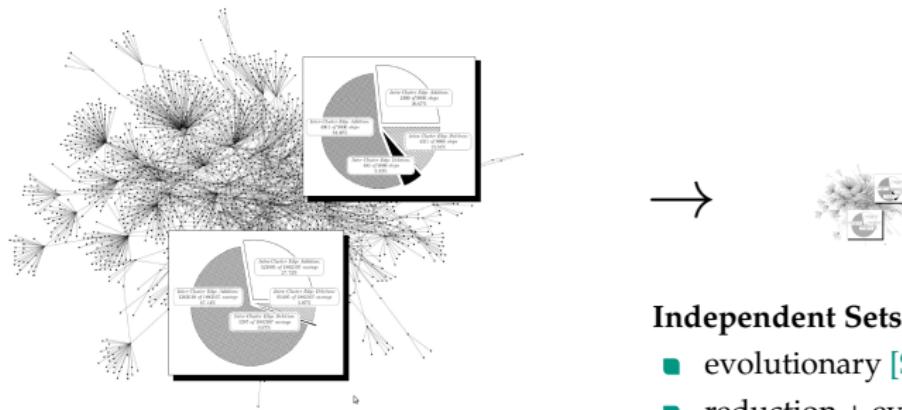
solve problem on **problem kernel** (using a heuristic)
→ obtain solution on input graph **quickly**

Kernelization

General Idea

Reductions:

rules to decrease graph size, while maintaining optimality



solve problem on **problem kernel**
→ obtain solution on input graph

Independent Sets

- evolutionary [SEA'15]
- reduction + evolutionary [ALX'16]
- online reductions + LS [SEA'16]
- shared-mem parallel [ALX'18]
- weighted exact [ALX'19]

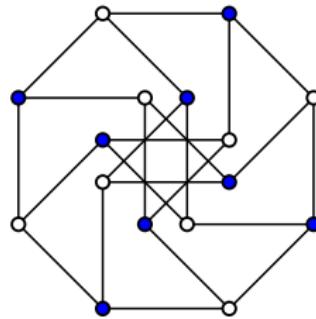
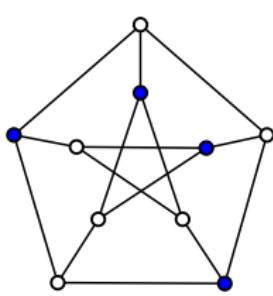
Definitions

Independent Set:

- subset $S \subseteq V$ such that there are no adjacent nodes in S

Maximum Independent Set (MIS):

- maximum cardinality set S



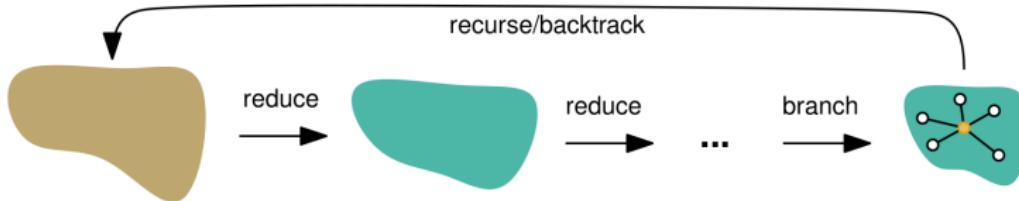
- related to **maximum clique** and **minimum vertex cover**
- finding a MIS is **NP-hard** and hard to approximate

Exact Algorithms

Independent Sets

Exact solution using **branch and bound**:

- Modify and remove subgraphs during recursions → **reductions**
- Branch when the graph can no longer be reduced



→ Running time $O(1.211^n)$ [Akiba, Iwata'16]

Works well, except when it doesn't:

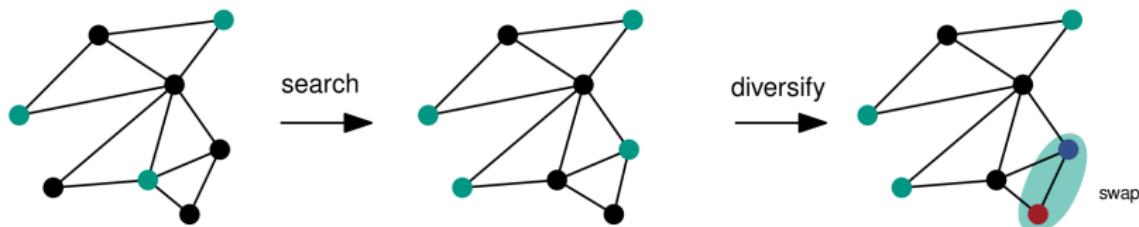
- Many networks with **small kernels** are easy to solve
- Similar instances with **large kernels** remain unsolved

Heuristic Algorithms

Huge real-world networks infeasible for exact algorithms
→ finding **high-quality** approximations in short time

Recent Approaches:

- Local search based on swaps [Andrade et al. 2012]

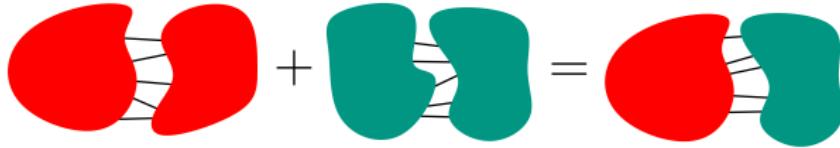


Heuristic Algorithms

Huge real-world networks infeasible for exact algorithms
→ finding **high-quality** approximations in short time

Recent Approaches:

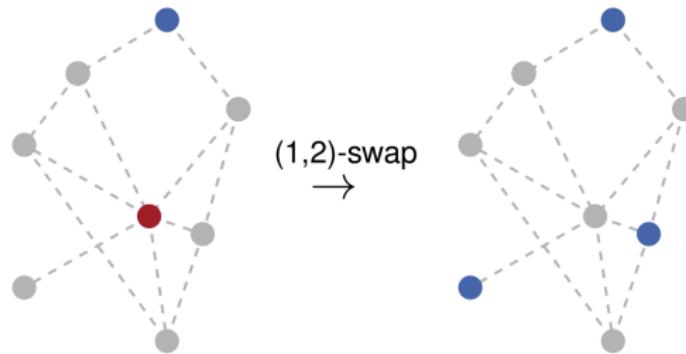
- Local search based on swaps [Andrade et al. 2012]
- Evolutionary algorithms (EvoMIS) [SEA'15]



Iterated Local Search

[Andrade et al.'12]

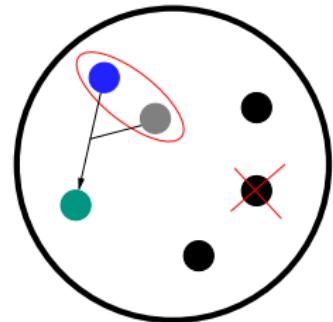
- Compute initial maximal independent set
- Improve using **(1, 2)-swaps**:
 - remove **single** solution node and insert **two** new ones
 - search for (1, 2)-swaps in time $O(m)$
- **Tabu mechanism** for “recently” swapped vertices
- No (1,2)-swap → **perturbation step**



Evolutionary Algorithm

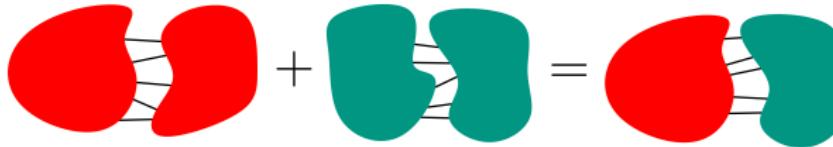
[SEA'15]

```
create initial population  $\mathcal{P}$ 
while stopping criterion not fulfilled
  select parents  $\mathcal{P}_1, \mathcal{P}_2$  from  $\mathcal{P}$ 
  recombine  $\mathcal{P}_1$  with  $\mathcal{P}_2$  to create offspring  $o$ 
  mutate offspring  $o$ 
  evict individual in population using  $o$ 
return fittest individual
```



Recombine Using Graph Partitioning:

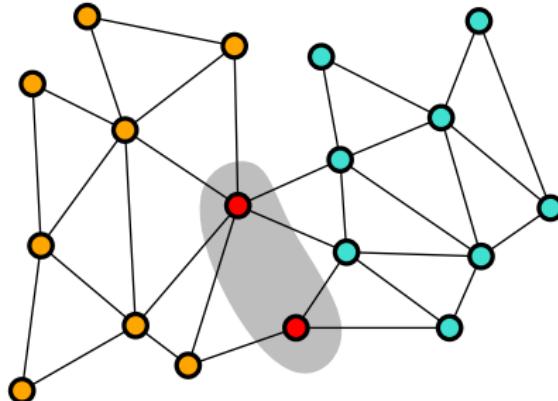
- exchange **whole blocks** of solutions
- small **objective** vital for efficiency



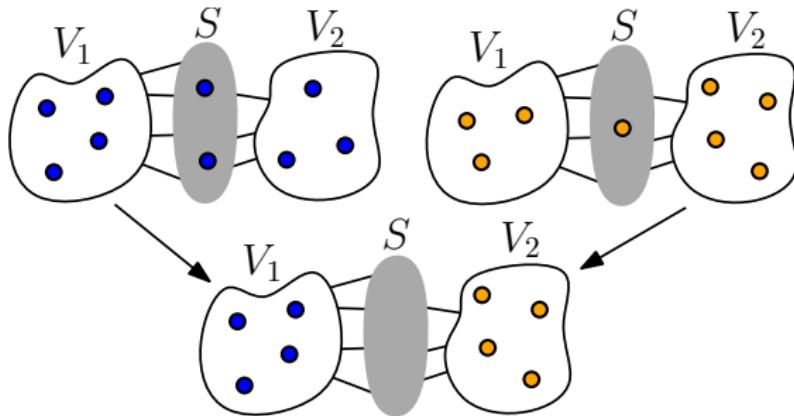
Separators

Partition graph $G = (V, E, c : V \rightarrow \mathbf{R}_{>0}, \omega : E \rightarrow \mathbf{R}_{>0})$
into k disjoint blocks + **node separator** s.t.

- total node weight of each block $\leq \frac{1+\epsilon}{k}$ total node weight
- total size of **node separator** as small as possible

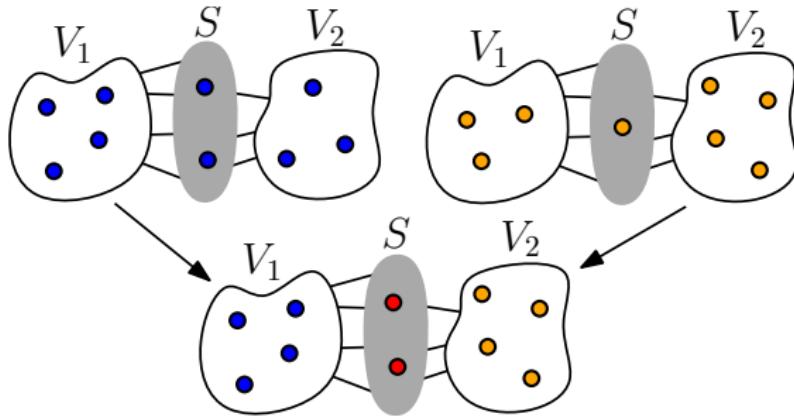


Separator Combine



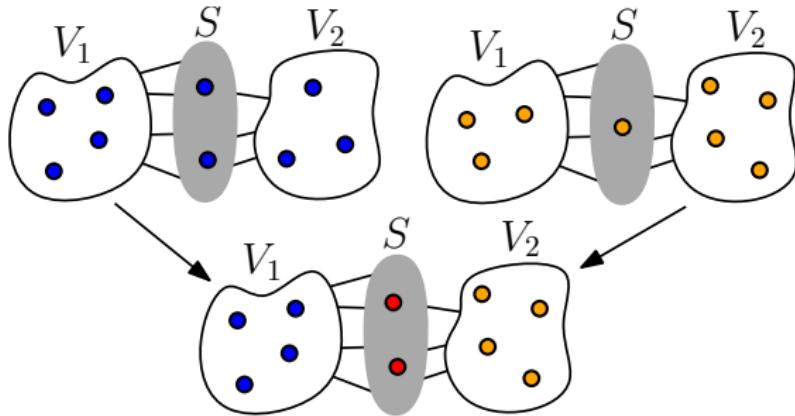
- build **node separator** $V = V_1 \cup V_2 \cup S$
- use node separator as **crossover point**
- combination takes **linear time** $O(n)$
- maximize with local search

Separator Combine



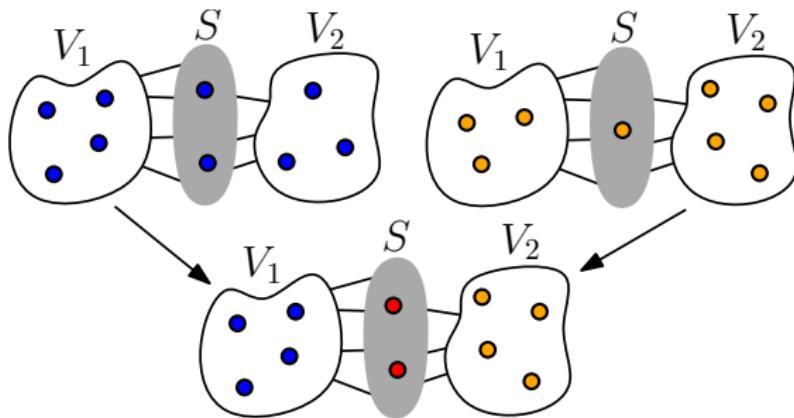
- build **node separator** $V = V_1 \cup V_2 \cup S$
- use node separator as **crossover point**
- combination takes **linear time** $O(n)$
- **maximize** with local search

Separator Combine



Good Solutions, but Slow

Separator Combine



Good Solutions, but Slow



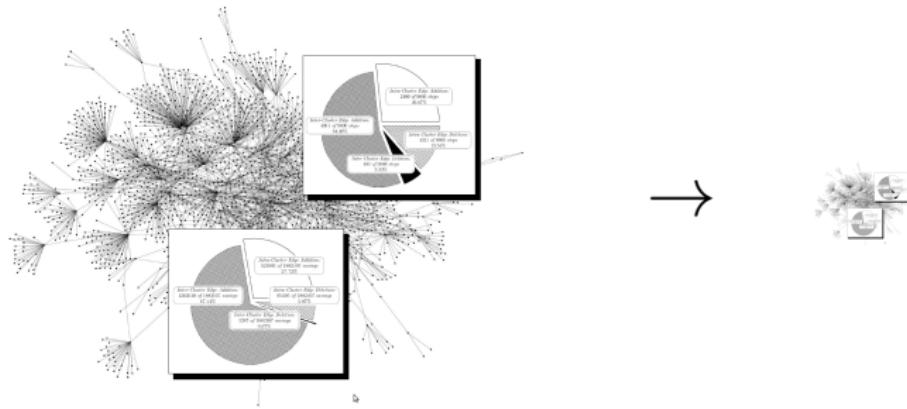
Apply EA on Kernel

Kernelization

[Akiba, Iwata'15]

Reductions:

rules to decrease graph size, while maintaining optimality



solve problem on **problem kernel** (using EA)
→ obtain solution on input graph

Kernelization

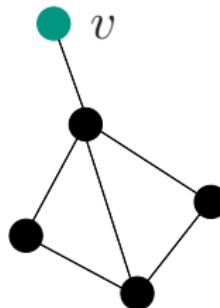
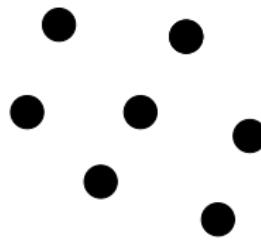
[Akiba, Iwata'15]

Reductions:

rules to decrease graph size, while maintaining optimality

Example:

remove degree 0 or 1 vertices recursively



there is always a MIS that contains v

if neighbor of v in MIS choose v instead; otherwise add v

Kernelization

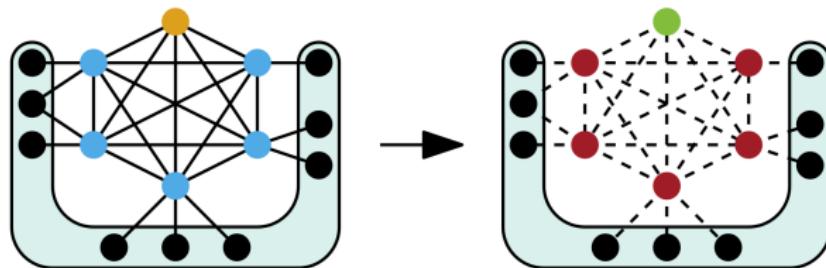
[Akiba, Iwata'15]

Reductions:

rules to decrease graph size, while maintaining optimality

Example:

isolated clique reduction



Kernelization

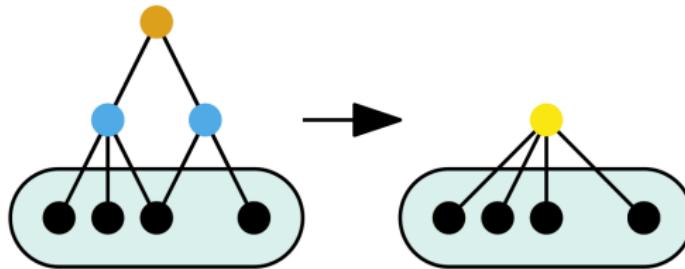
[Akiba, Iwata'15]

Reductions:

rules to decrease graph size, while maintaining optimality

Example:

vertex folding



Kernelization

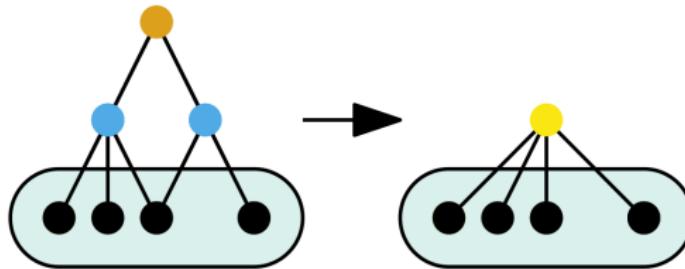
[Akiba, Iwata'15]

Reductions:

rules to decrease graph size, while maintaining optimality

Example:

vertex folding



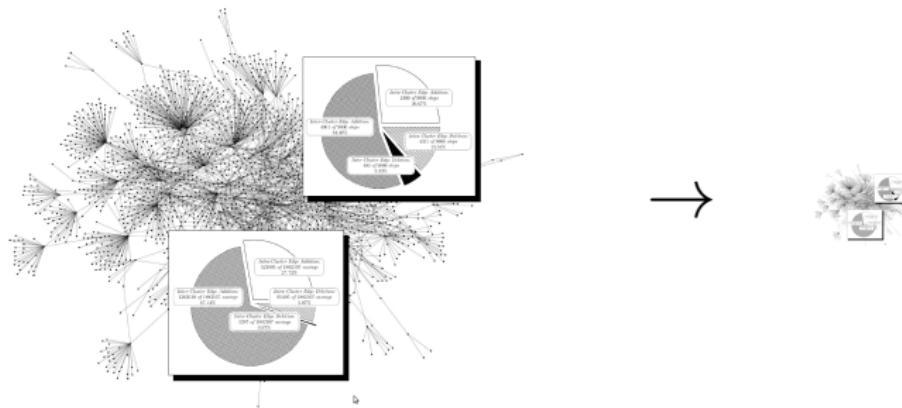
more reductions used in practice → [ALENEX'16]

Kernelization

General Idea

Reductions:

rules to decrease graph size, while maintaining optimality

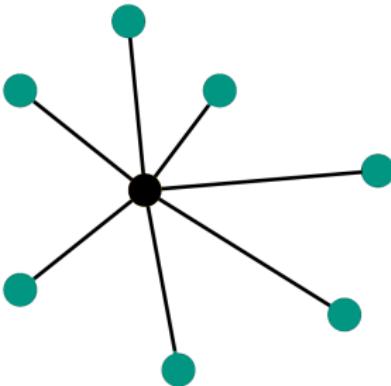


solve problem on **problem kernel**
→ obtain solution on input graph **quickly**

Guess “likely candidates”

can we **guess** vertices that are in MIS?

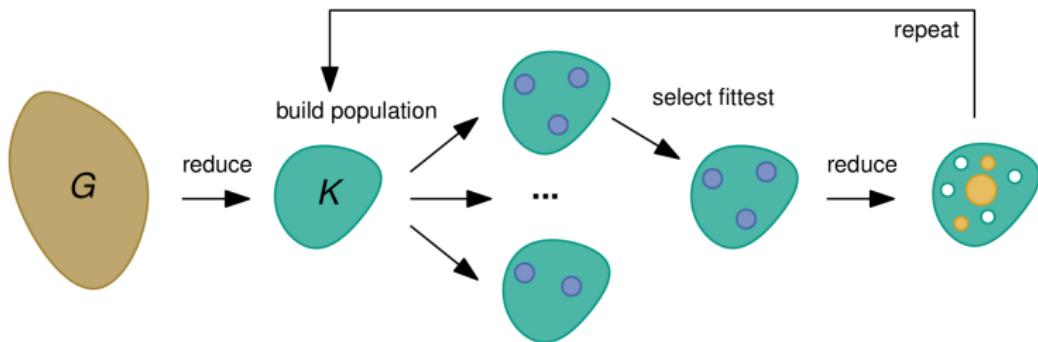
idea: select small-degree vertices from “fittest” independent set
apply more reductions and recurse!



Guess “likely candidates”

can we **guess** vertices that are in MIS?

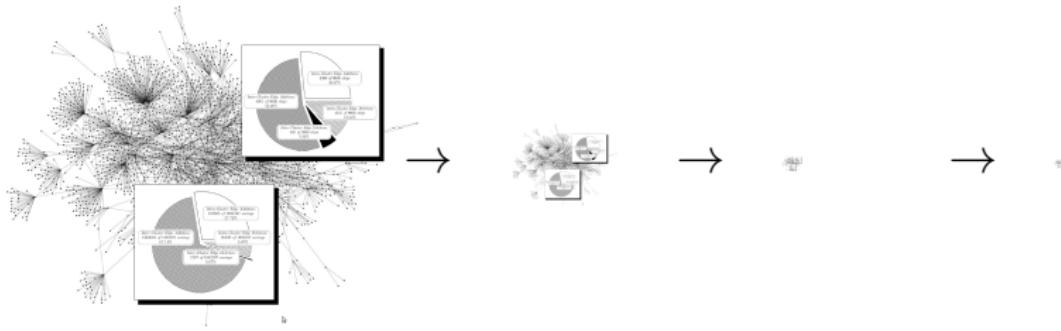
idea: select small-degree vertices from “fittest” independent set
apply more reductions and recurse!



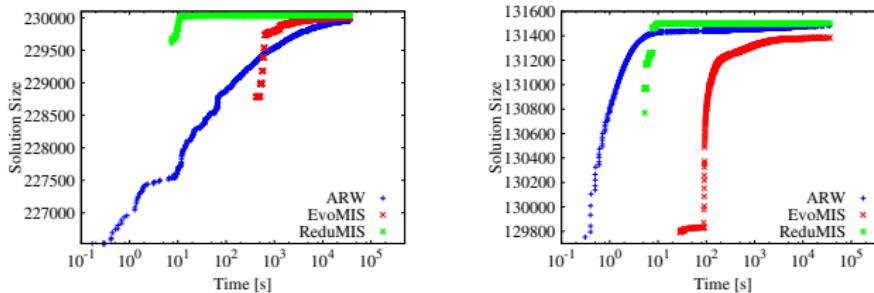
Guess “likely candidates”

can we **guess** vertices that are in MIS?

idea: select small-degree vertices from “fittest” independent set
apply more reductions and recurse!



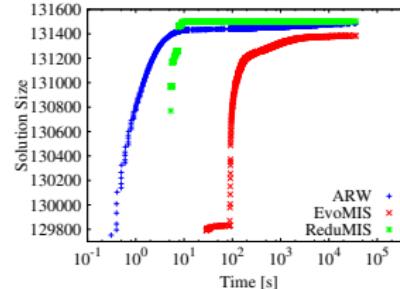
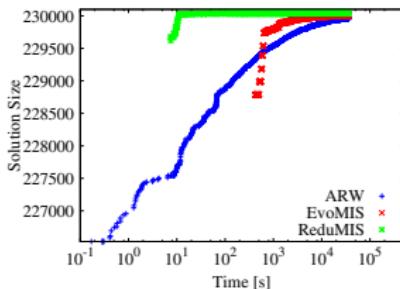
Near-optimal on “difficult networks”



Results:

- finds exact MIS faster, when exact algorithm is slow:
 - Skitter **48 min** → **21 min**
 - Stanford **13 hours** → **5 min**
 - bcsstk30 **8.6 hours** → **2.4 sec**
 - Skitter **2 hours** → **28 sec**, ...
- finds exact MIS, for large networks with known MIS size
- consistently finds larger solutions on social and road networks
- even as we scale to graphs to **10M** to **100M** nodes

Near-optimal on “difficult networks”



Problems:

- Kernelization gives **high-quality** at cost of preprocessing time
- ARW is **fast** but struggles with complex scale-free networks

→ accelerate local search?
→ parallelization?
→ get rid of complicated evolutionary algorithm?

ARW on Complex Networks

Issues

- High degree vertices take long to process ...
... and are unlikely to be in MIS → High-degree cutting
- Many vertices are always in some solution → Kernelization



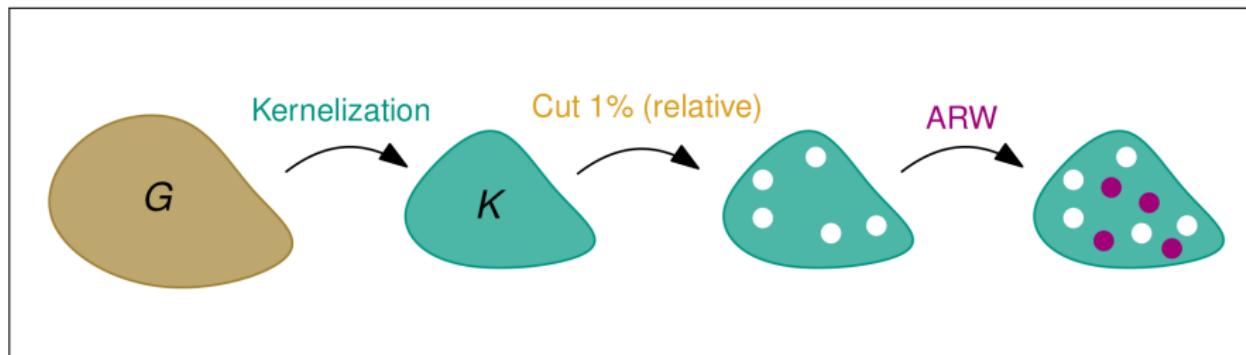
Strategy I

Full Kernelization

- Repeatedly apply reductions → kernel graph

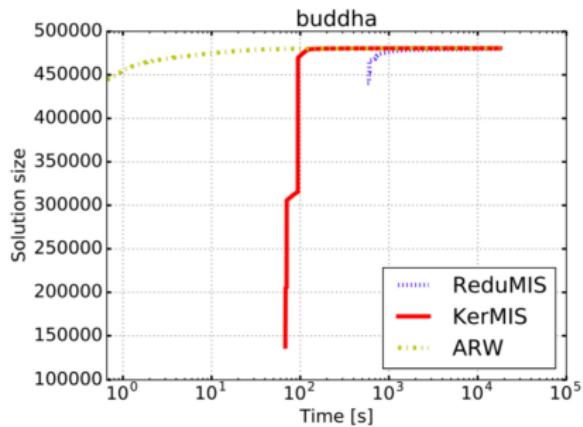
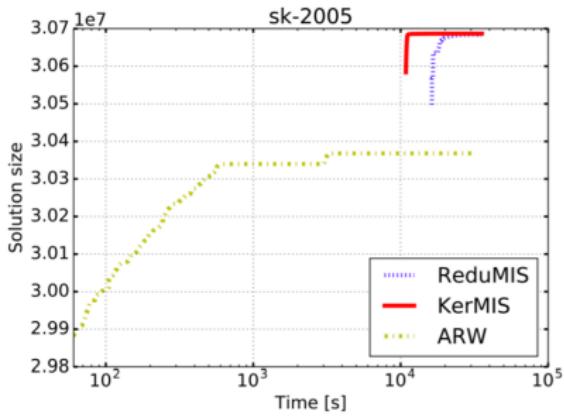


KerMIS:



KerMIS evaluation

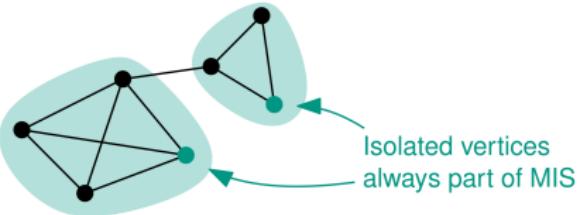
- Still long **preprocessing time** to compute kernel
- Inexact reductions **improve performance** of local search
- Quality comparable to **ReduMIS**
 - → Inexact reductions **don't worsen solution quality**



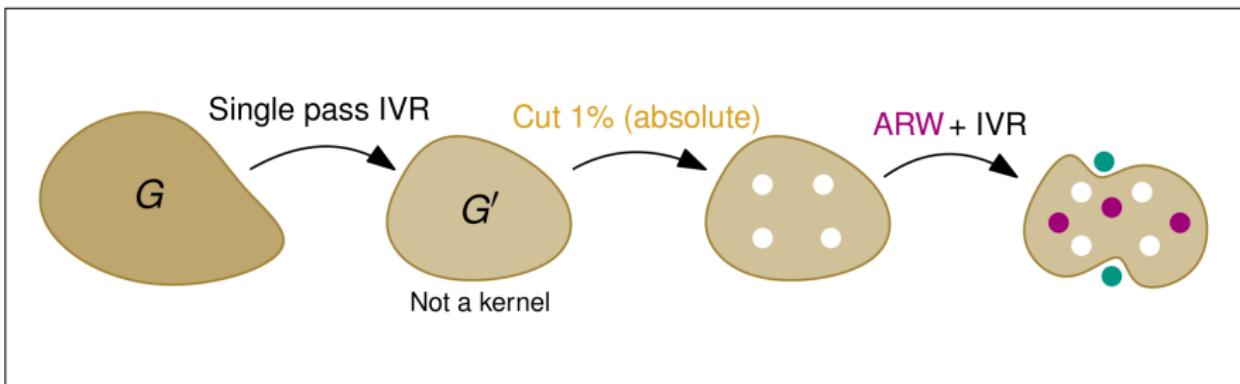
Strategy II

Online “Removal”

- Isolated vertex removal (IVR)

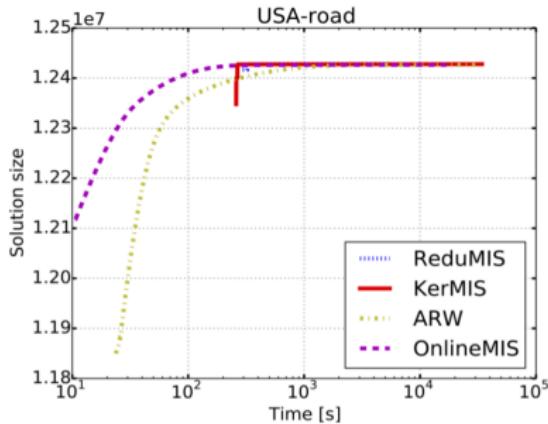
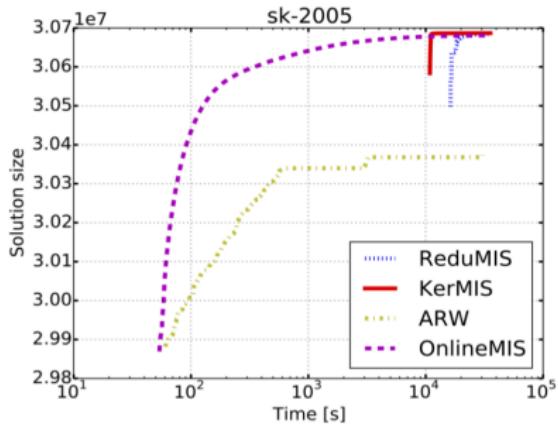


OnlineMIS:



OnlineMIS evaluation

- Significant improvements in **speed and quality**
- Huge scale-free networks can be processed efficiently

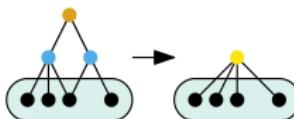


- Maximum speedup ≥ 300 over ReduMIS for 14/24 instances

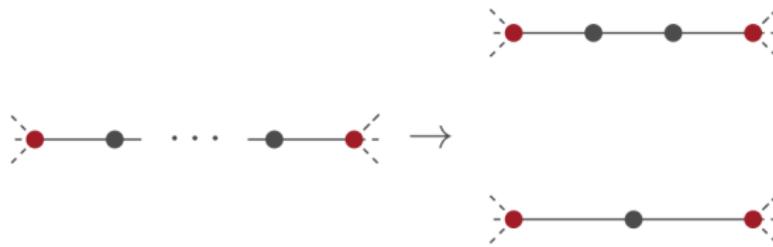
Linear Time Reductions

[Chang et al.'17]

~~ vertex folding is slow with high-degree neighbors



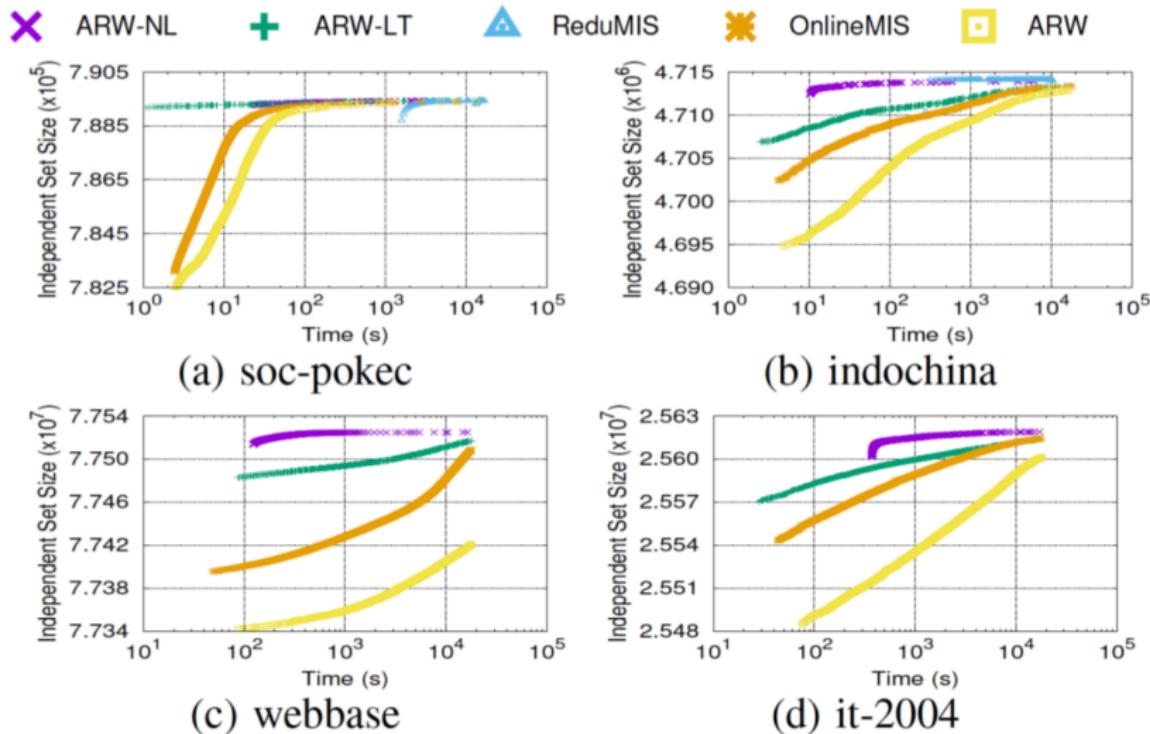
~~ avoid it



Repeat: add small degree vertices to solution + reduce

Linear Time Reductions

[Chang et al.'17]



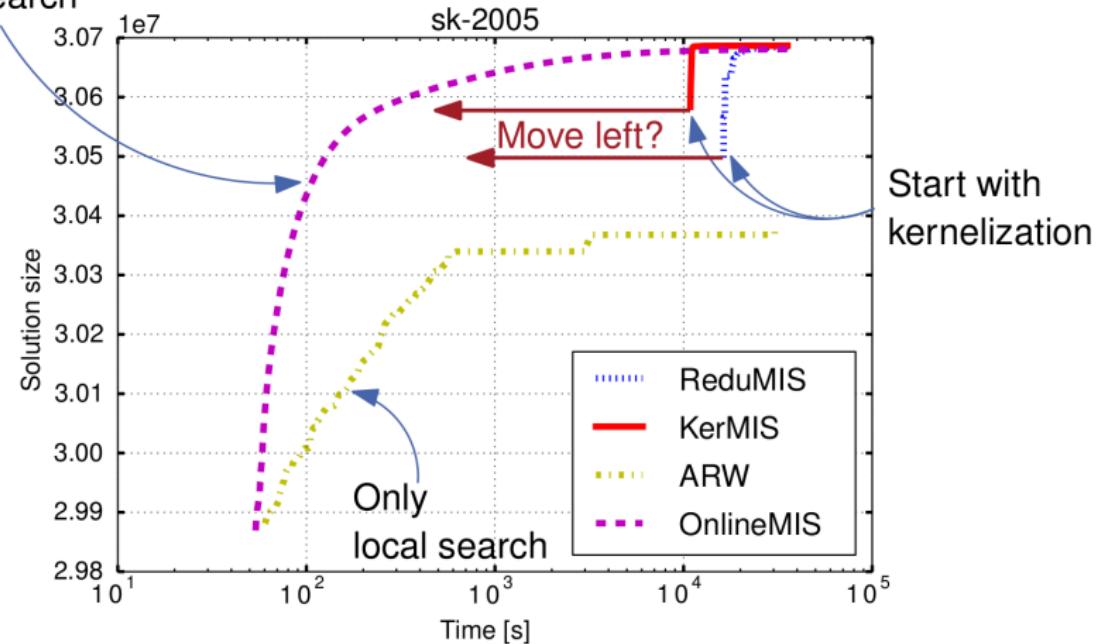
Scalable Reductions

Effective Reductions Are Slow!

Graph		LinearTime		NearLinear		VCSolver	
name	n	K	time	K	time	K	time
uk-2002	19M	11.7M	1.5	4.0M	28.0	0.2M	336.9
arabic-2005	23M	15.6M	2.6	6.7M	246.1	0.6M	1 033.2
gsh-2015-tpd	31M	2.0M	11.6	1.2M	97.4	0.4M	372.3
uk-2005	39M	28.2M	2.5	5.9M	60.5	0.8M	541.4
it-2004	41M	27.1M	3.3	11.3M	1 544.6	1.6M	6 749.0
sk-2005	51M	*	*	*	*	3.2M	10 010.5
uk-2007-05	106M	*	*	*	*	3.5M	18 829.4
webbase-2001	118M	51.7M	13.0	17.3M	121.1	0.7M	4 207.8
asia.osm	12M	626.7K	0.8	594.4K	1.4	15.2K	204.7
road_usa	24M	2.5M	2.5	2.4M	4.1	0.2M	310.0
europe.osm	51M	1 500.0K	4.1	1 329.9K	6.1	8.4K	302.4
rgg26	67M	67.1M	1.0	51.3M	172.6	49.6M	9 887.7
rhg	100M	*	*	*	*	0	124.0
del24	17M	16.8M	0.2	15.6M	12.7	12.4M	4 789.5
del26	67M	67.1M	0.7	62.5M	53.3	49.9M	20 728.7

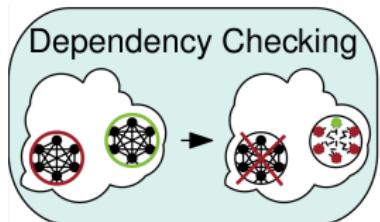
Even Faster??

Reductions during local search

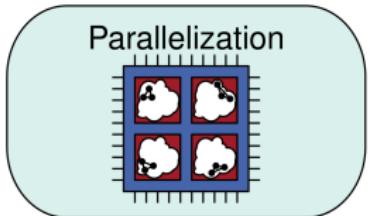


Remaining Ingredients

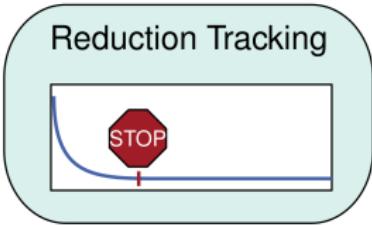
Dependency Checking



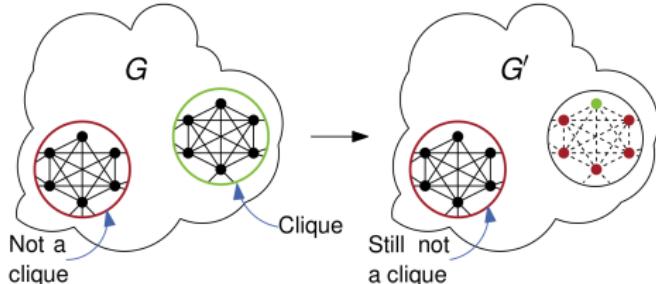
Parallelization



Reduction Tracking



Dependency Checking

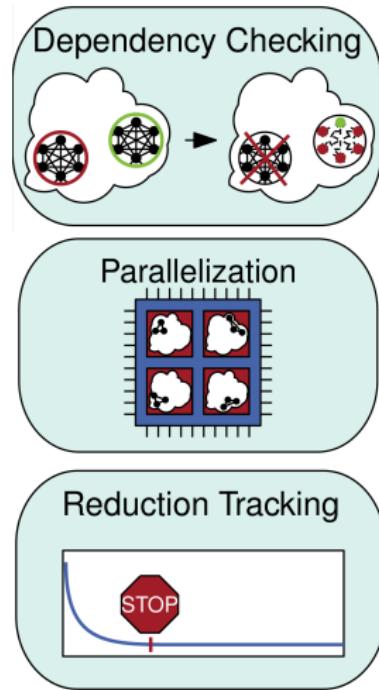


No reduction in G and $N_G(v) = N_{G'}(v) \Rightarrow$
No reduction in G'

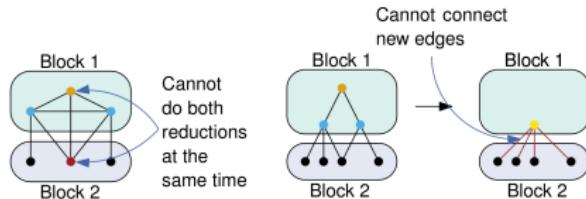
Used for

- Isolated Clique Reduction
- Degree 2 Fold
- Twin Reduction

Remaining Ingredients



Parallelization
partition graph in blocks, reduce separately



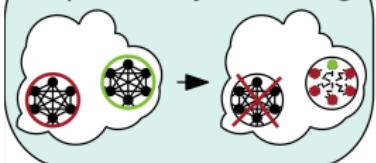
→ deal with conflicts

Ingredients:

- ParHIP for partitioning
- Parallelize LP reduction with parallel maximum bipartite matching
[Azad et al.'17]

Remaining Ingredients

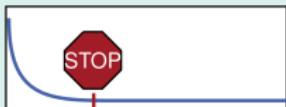
Dependency Checking



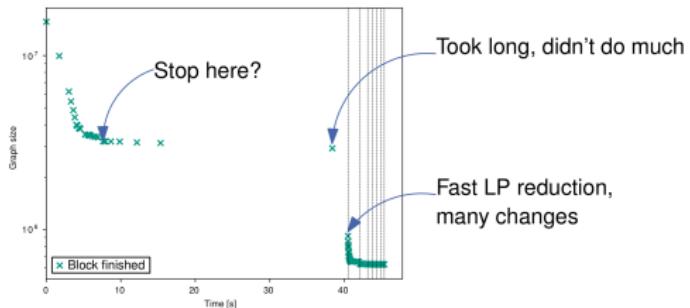
Parallelization



Reduction Tracking



Reduction Tracking stop kernelization early

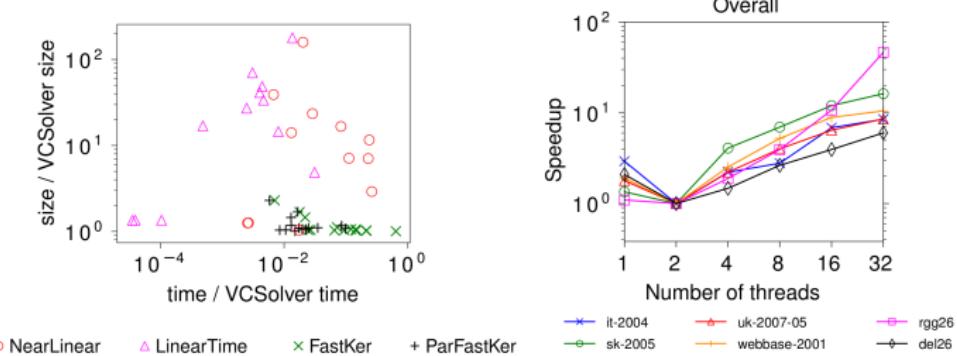


- some blocks take longer than others
- few changes after a while
 - stop “local” reductions early
 - run LP reduction

Scalable Reductions

NearLinear		VCSSolver		ParFastKer		
$ \mathcal{K} $	time	$ \mathcal{K} $	time	$ \mathcal{K} $	time	su
4.0M	28.0	0.2M	336.9	0.3M	11.8	28.4
6.7M	246.1	0.6M	1 033.2	0.6M	25.7	40.2
1.2M	97.4	0.4M	372.3	0.5M	32.0	11.7
5.9M	60.5	0.8M	541.4	0.9M	53.3	10.1
11.3M	1 544.6	1.6M	6 749.0	1.7M	151.8	44.4
*	*	3.2M	10 010.5	3.5M	178.3	56.1
*	*	3.5M	18 829.4	3.7M	372.4	50.6
17.3M	121.1	0.7M	4 207.8	0.9M	54.9	76.6
594.4K	1.4	15.2K	204.7	34.9K	1.2	169.8
2.4M	4.1	0.2M	310.0	0.2M	4.1	76.0
1 329.9K	6.1	8.4K	302.4	14.2K	4.9	61.3
51.3M	172.6	49.6M	9 887.7	49.8M	150.3	65.8
*	*	0	124.0	16	64.6	1.9
15.6M	12.7	12.4M	4 789.5	12.9M	51.5	93.1
62.5M	53.3	49.9M	20 728.7	51.7M	179.0	115.8

Experimental Results



- VCSolver [Akiba and Iwata'16]: slow but small kernels
- LinearTime and NearLinear [Chang et al.'17]: fast but large kernels
→ integrate as preprocessing

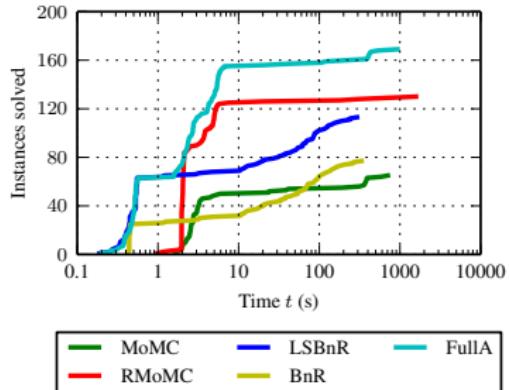
Summary:

- Orders of magnitude smaller than fast methods
- Orders of magnitude faster than algorithm with similar-sized kernels
- Integration into local search (small kernels matter):
 - larger independent sets faster

PACE Challenge 2019

Back to Exact Algorithms ...

- Exact vertex cover solver
→ based on reductions + local search + portfolio of exact solvers
- Exact Solvers:
 - Branch-and-Reduce for vertex cover
 - Clique solver by Li et al.
 - Get the easy instances **fast**, then get the hard instances.
- 24 competing algorithms
- We *won*



PACE Challenge 2019

Back to Exact Algorithms ...



ABOUT - PROBLEMS - GLSL TUTORIAL NEWS



Please note that displayed time is user time on CPU. You may receive **TLE** by exceeding limit in real time, while your program's user time is 0, i.e. sleep(10000000)

#	USER	LANGUAGE	SCORE	TIME [S]	1	2	3	4	5	6	7	8	9	10	11	12	13
1	WeGotYouCovered	Static binary	84.59	31,968.68	0.40	0.00	1.82	0.15	1.56	0.55	0.08	TLE	0.30	0.41	0.07	0.03	0.0
2	peaty	CMake	54.60	43,365.30	0.18	0.00	0.24	0.12	0.30	0.20	0.00	TLE	0.14	0.13	0.05	0.11	0.0
3	vasily_alfarov	C++	30.85	79,474.67	0.10	0.00	0.69	0.07	2.06	0.06	0.00	TLE	0.10	0.15	4.10	0.13	0.0
4	ksimonov	CMake	29.90	119,019.56	0.08	0.00	2.58	0.05	4.22	0.12	0.02	TLE	0.01	0.10	0.02	0.04	0.0
5	opm	Static binary	29.58	122,553.49	0.38	0.12	1.39	0.28	4.88	0.28	0.08	TLE	0.22	0.29	0.18	0.23	0.1
6	tob	Static binary	29.39	91,191.78	3.97	0.09	2.27	0.04	3.13	6.81	0.96	MLE	2.87	6.30	0.20	0.14	0.0
7	hub	Static binary	28.96	124,576.85	0.00	0.00	1.58	0.04	3.58	0.04	0.03	TLE	0.06	0.08	0.02	0.02	0.0
8	sfs	CMake	28.29	119,732.27	2.71	0.00	2.77	0.11	7.00	2.84	0.72	TLE	2.55	3.06	0.00	0.00	0.0
9	bogdan	Static binary	28.13	63,740.73	1.72	0.06	28.47	0.78	29.94	1.50	0.53	TLE	1.48	1.54	0.64	0.46	0.2
10	Vertex_Cover_Solver	C	23.71	128,498.75	2.50	0.14	2.08	0.24	6.63	3.02	0.62	TLE	2.28	3.02	2.69	1.70	0.1
11	pace.challenge	Static binary	21.71	14,694.65	192.37	127.98	122.44	102.44	97.88	174.41	138.37	131.51	117.72	186.72	186.40	120.58	120.3
12	swats	Static binary	18.22	15,075.24	75.35	0.16	124.57	7.99	138.93	88.73	81.62	143.08	88.90	78.46	WA	WA	WA
13	xelmh	CMake	15.02	143,878.07	0.12	0.05	26.87	0.44	85.23	0.15	0.02	TLE	0.14	0.13	2.51	5.92	0.1
14	pptale	CMake	11.93	157,640.11	1.07	0.00	1,074.90	0.03	TLE	1.10	0.39	TLE	1.49	1.94	1.74	0.00	0.0
15	vader95	Python3	11.71	154,823.74	4.69	2.00	629.04	4.87	832.33	4.86	2.59	TLE	4.16	5.03	2.40	50.80	2.4
16	arrighi	CMake	11.56	158,468.48	8.19	0.18	TLE	0.17	TLE	10.08	2.62	TLE	7.82	9.33	12.00	0.14	0.1
17	johannes	CMake	10.98	154,533.16	0.23	0.10	WA	0.32	WA	0.27	0.08	TLE	0.22	0.26	TLE	0.46	0.1
18	tomglint	Static binary	7.99	162,464.67	0.14	0.00	TLE	TLE	TLE	0.18	0.00	TLE	0.16	0.17	TLE	TLE	TLE
19	NTU	C++	3.91	8,253.19	MLE	0.32	WA	0.88	WA	MLE	9.63	WA	MLE	MLE	WA	1.10	WA

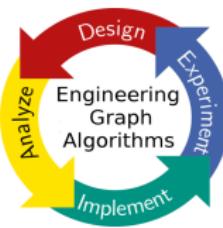
Multilevel Algorithms

Evolutionary Computation

Parallel Programming

Kernelization

shared-, distributed-, external-, internal memory



Algorithms for

graph partitioning
graph clustering
graph generation

process mapping
minimum cuts
independent sets

longest paths
graph drawing
matching

node separators
....

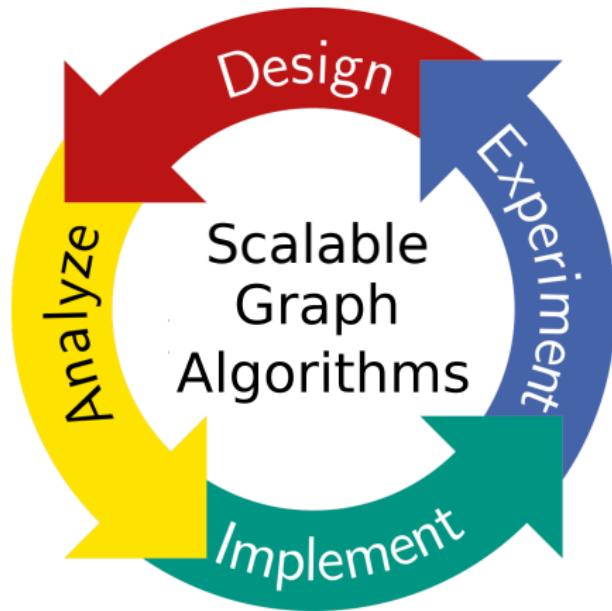
Open Source

Applications

territory design
large scale simulations
quantum annealing

route planning
distributed system design
nuclei segmentation

multiprocessor scheduling
high-throughput DNA sequencing
....



Room 6.03

exists