

# Research Overview

Multilevel Algorithms

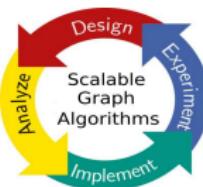
Evolutionary Computation

Parallel Programming

Kernelization

Dynamic Algorithms

shared-, distributed-, external-, internal memory



Algorithms for

graph partitioning  
graph clustering  
graph generation

process mapping  
minimum cuts  
independent sets

longest paths  
graph drawing  
(dyn.) matching

node separators  
dyn. reachability  
....

Open Source

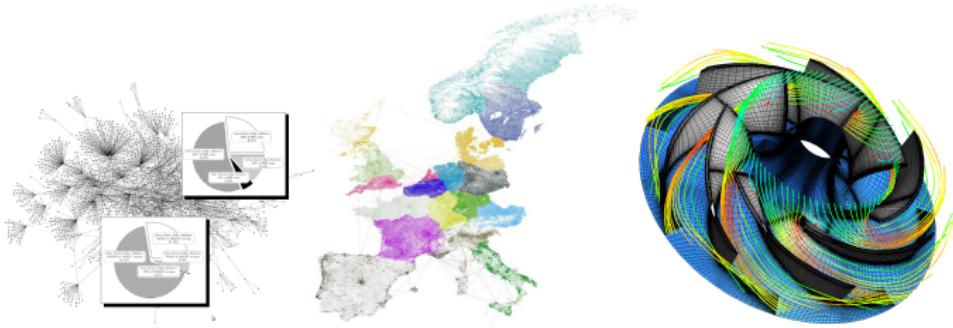
Applications

territory design  
large scale simulations  
quantum annealing

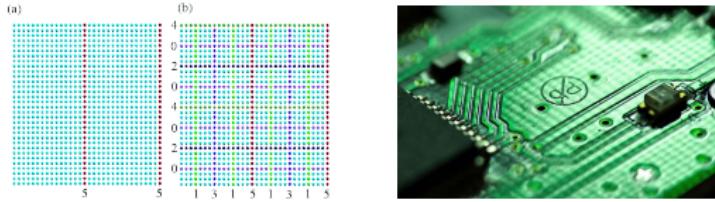
route planning  
distributed system design  
nuclei segmentation

multiprocessor scheduling  
high-throughput DNA sequencing  
....

# Applications

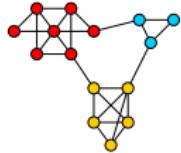


$$\mathbf{R}^{n \times n} \ni Ax = b \in \mathbf{R}^n$$



# Highlights

# Graph Clustering

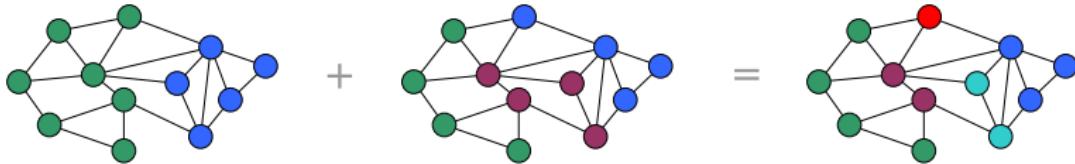


## Problem:

- partition graph into tightly connected groups
- clustering paradigm: **internally dense** and **externally sparse**
- quality measure: **modularity**

## Recombination Mechanism

- borrowed from ML: ensemble learning → **maximum overlap**



→ results in **overlap clustering**, contract + recluster

→ + coarse-grained parallelization

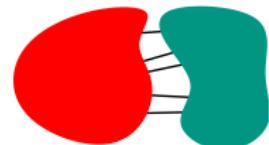
**Outperform all DIMACS Challenge results  
with ONE solver and in less time**

# Minimum Cuts

... and Multiterminal Cuts

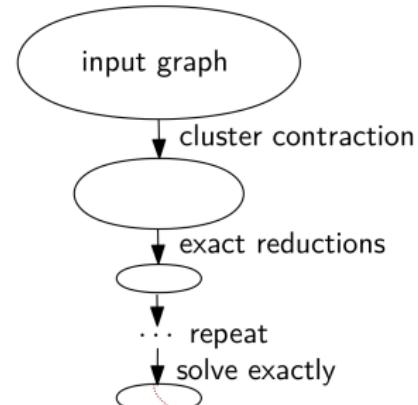
**Cut:** A **cut** in a multigraph is a partition of  $V = C \cup \bar{C}$

**Problem:** size of minimum cut in  $G$ ?



## Contributions:

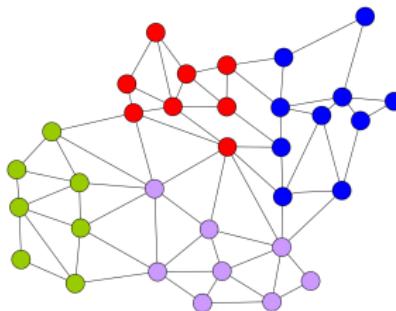
- (in)exact reductions + solve kernel to optimality  
→ linear running time, but potentially suboptimal cuts
- NO guarantee, but experiments say likely opt
- reductions depend on bound  $\hat{\lambda}$ 
  - ~~ use  $\hat{\lambda}$  in exact (parallel) NOI algorithm
  - ~~ **fastest exact minimum cut algorithm**
  - ≈ order of magnitude
- multiterminal cut:  
**first practical solver for large inputs**



# $\epsilon$ -Balanced Graph Partitioning

Partition graph  $G = (V, E, c : V \rightarrow \mathbf{R}_{>0}, \omega : E \rightarrow \mathbf{R}_{>0})$   
into  $k$  disjoint blocks s.t.

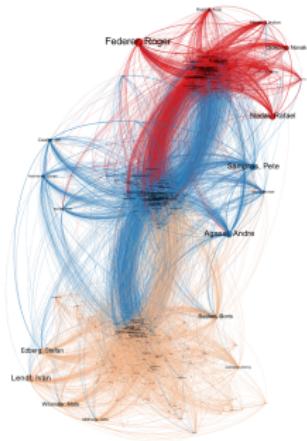
- total **node weight** of each block  $\leq \frac{1 + \epsilon}{k}$  total node weight
- total weight of **cut** edges as small as possible



## Applications:

graph processing frameworks, parallel sparse matrix vector mult., ...

# The Common Parallel Approach



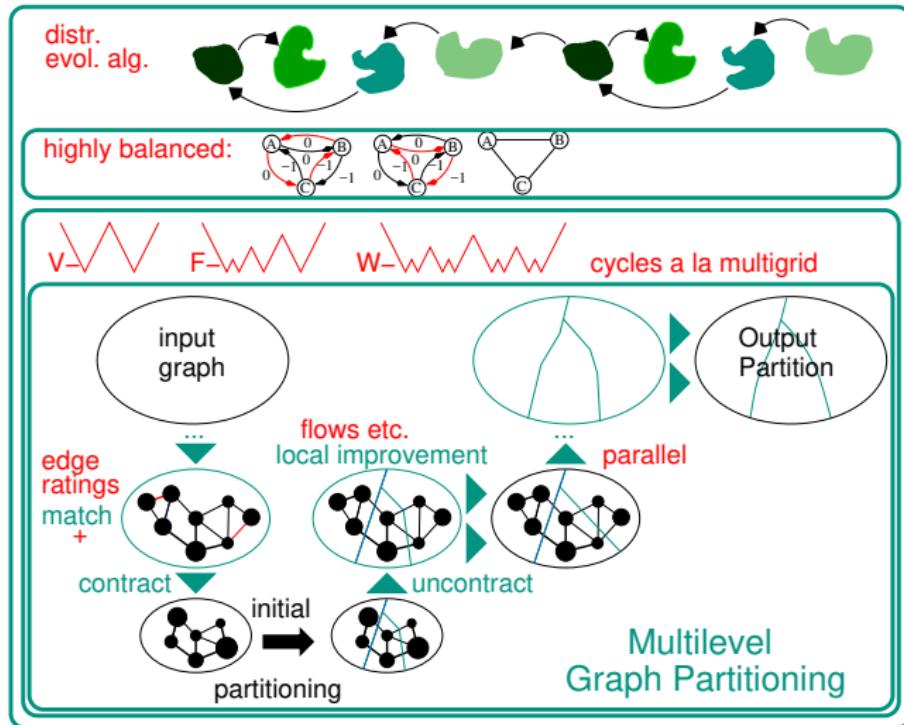
- Mesh partitioned via dual graph
  1. Each volume (data, calculation) represented by a vertex (+edges)
  2. Interdependencies represented by edges
- All PE's get same amount of work
- Communication is expensive

## Graph Partitioning Problem:

Partition a graph into (almost) equally sized blocks, such that the number of edges connecting vertices from different blocks is minimal.

# Karlsruhe High Quality Partitioning

<https://github.com/KaHIP/KaHIP>

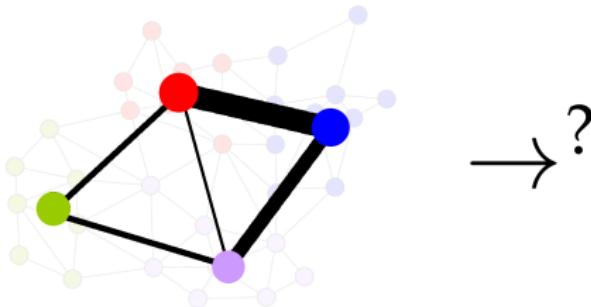


THE solver, when it comes to high solution quality.

# Karlsruhe High Quality Partitioning

- 
- pre 2014:
    - multi-level, localized and flow-based local search
    - perfectly balanced,
    - distributed evolutionary algorithms
    - specialized algorithms for road networks
  - social networks
  - distributed memory
  - semi-external memory
  - hypergraph partitioning
  - node separators
  - (hyper)DAG partitioning
  - shared-memory parallel
  - ILP-based
  - node ordering

# Process Mapping

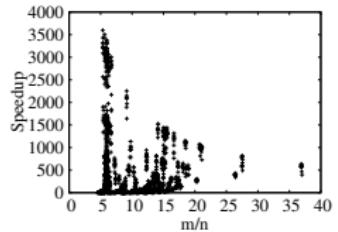


## Exploiting Assumptions:

- communication matrix  $\mathcal{C}$  is **sparse** → graph  $G_{\mathcal{C}}$
- compute system is **hierarchically structured**

## Yields:

- three orders of magnitude faster local search
- 50% improved mapping quality over state-of-the-art



# Graph Generation

## Ideas:

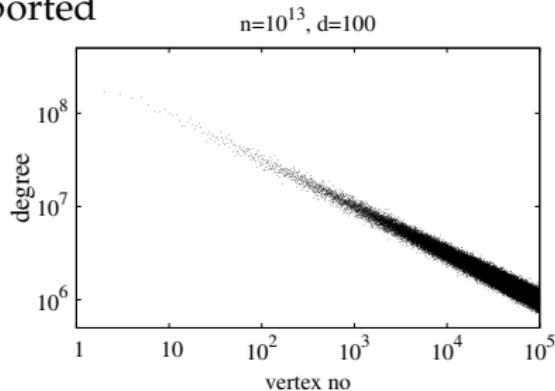
- replace randomness by pseudo-randomness, e.g. hashing
- replace array accesses by recomputation  
→ embarrassingly parallel algorithm without communication

## Results:

- generate Peta-edge BA graph in < one hour ( $10^{15}$  edges, 16K cores)
- 20 000 times larger than previously reported
- 16x faster than RMAT ( $m = 50 \cdot 10^9$ )
- communication efficient algorithms

## more work:

- Barabasi-Albert
- more models:  
ER,  $G(n,m)$ ,  $G(n,p)$ ,  
RHG, RGG, DEL



# Graph Drawing

$$G = (V, E, d)$$
$$d : E \rightarrow \mathbf{R}$$
$$\Rightarrow$$

Maximal Entropy Stress Model:

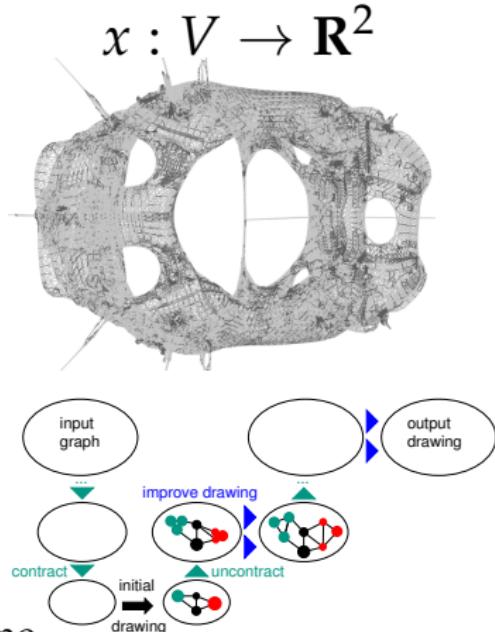
$$\max H(x) := \sum_{\{u,v\} \notin E} \ln ||x_u - x_v||$$

subject to  $||x_u - x_v|| = d_{uv}, \{u, v\} \in E$

Contributions:

- multilevel integration of iterative scheme
- approximate long-range forces, by cluster contraction
- employ parallelism

Draw graphs with millions of vertices in seconds.



# Support Vector Machines

- **Binary Classification Problem:**

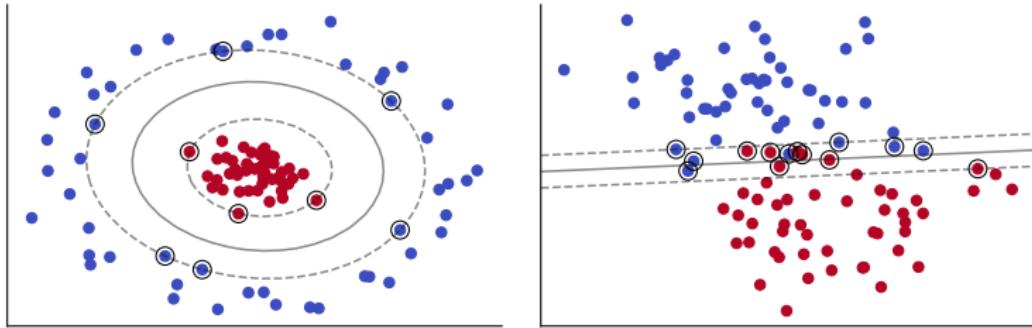
Train classifier on  $n$  labeled data points ( $x_i \in \mathbf{R}^d, y_i \in \{-1, 1\}$ )

Goal: Assign label  $y_{n+1}$  to new data points  $x_{n+1}$

- **SVM Optimization Problem:**

$$\text{min. } \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi_i$$

$$\text{s.t. } y_i(w \cdot \phi(x_i) - b) \geq 1 - \xi_i, \quad \xi_i \geq 0.$$



# Support Vector Machines

## ■ Binary Classification Problem:

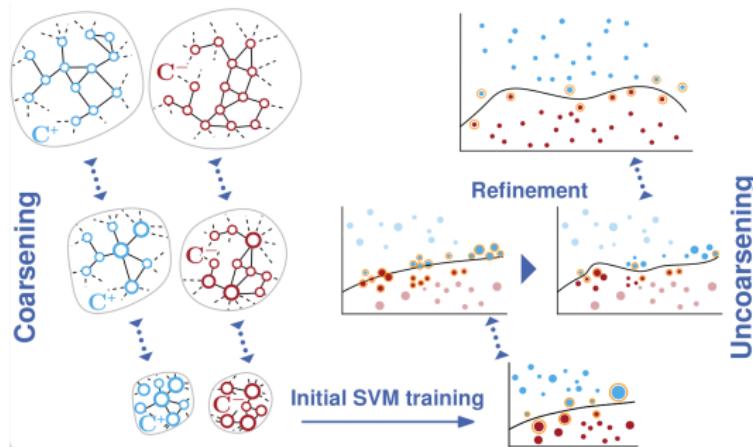
Train classifier on  $n$  labeled data points ( $x_i \in \mathbf{R}^d, y_i \in \{-1, 1\}$ )

Goal: Assign label  $y_{n+1}$  to new data points  $x_{n+1}$

## ■ SVM Optimization Problem:

$$\text{min. } \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi_i$$

$$\text{s.t. } y_i(w \cdot \phi(x_i) - b) \geq 1 - \xi_i, \quad \xi_i \geq 0.$$



# Support Vector Machines

## ■ Binary Classification Problem:

Train classifier on  $n$  labeled data points ( $x_i \in \mathbf{R}^d, y_i \in \{-1, 1\}$ )

Goal: Assign label  $y_{n+1}$  to new data points  $x_{n+1}$

■ SVM Optimization Problem:

$$\begin{aligned} \text{min. } & \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi_i \\ \text{s.t. } & y_i(w \cdot \phi(x_i) - b) \geq 1 - \xi_i, \quad \xi_i \geq 0. \end{aligned}$$

multilevel SVM using label propagation  
→ comparable classification quality  
→ up to two orders of magnitude faster training

currently working on parallelization

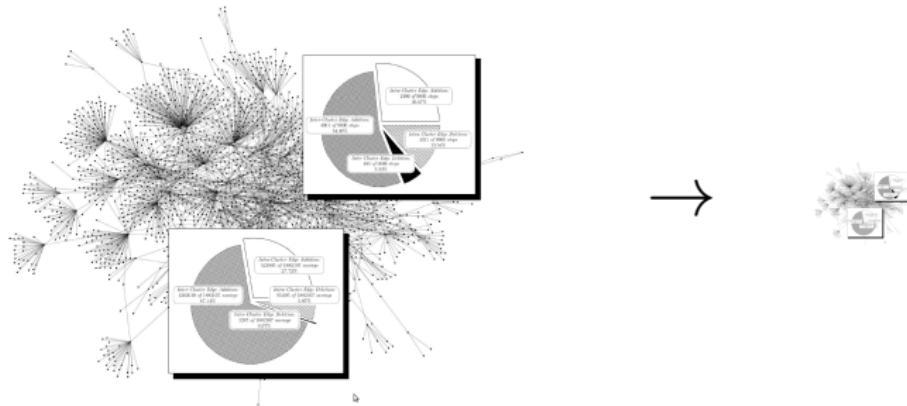
# Kernelization

## Example Independent Sets

find subset  $S \subseteq V$  such that there are no adjacent nodes in  $S$

### Reductions:

rules to decrease graph size, while maintaining optimality



solve problem on **problem kernel**  
→ obtain solution on input graph

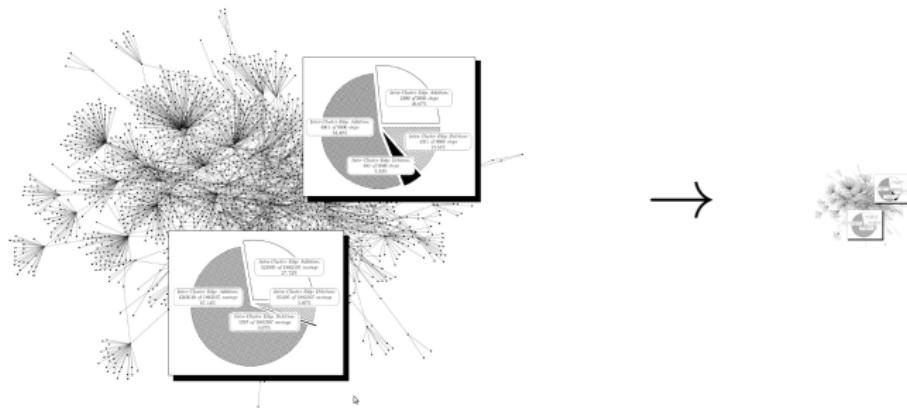
# Kernelization

## Example Independent Sets

find subset  $S \subseteq V$  such that there are no adjacent nodes in  $S$

### Reductions:

rules to decrease graph size, while maintaining optimality



solve problem on **problem kernel** (using a heuristic)  
→ obtain solution on input graph **quickly**

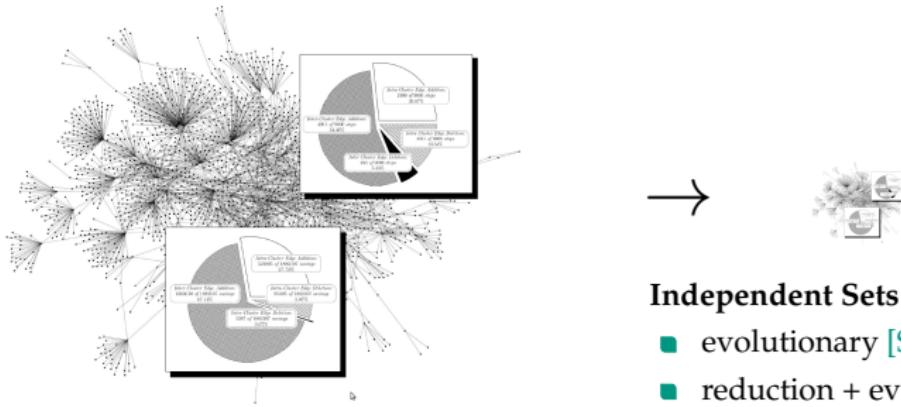
# Kernelization

## Example Independent Sets

find subset  $S \subseteq V$  such that there are no adjacent nodes in  $S$

### Reductions:

rules to decrease graph size, while maintaining optimality



### Independent Sets

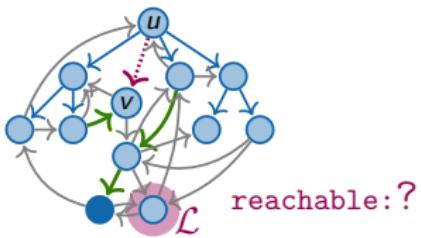
- evolutionary [SEA'15]
- reduction + evolutionary [ALX'16]
- online reductions + LS [SEA'16]
- shared-mem parallel [ALX'18]
- weighted exact [ALX'19]
- ...

solve problem on **problem kernel**  
→ obtain solution on input graph

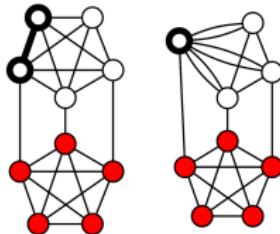
**State-of-the-art for (weighted) independent sets**

# Dynamic Graph Algorithms

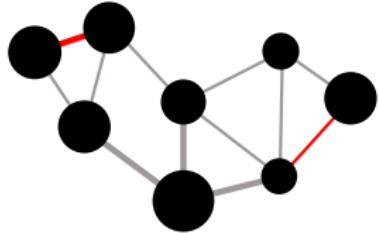
dynamic reachability



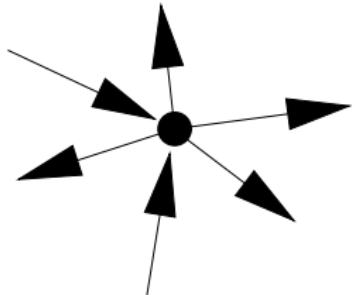
dynamic mincut



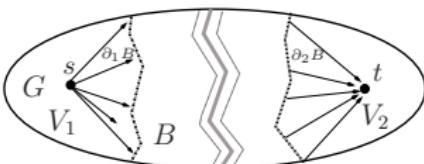
dynamic matching



dynamic  $\delta$ -orientation



dynamic maxflow



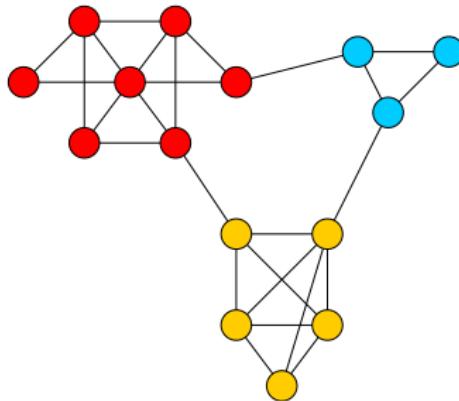
# In Depth Example

## Graph Clustering

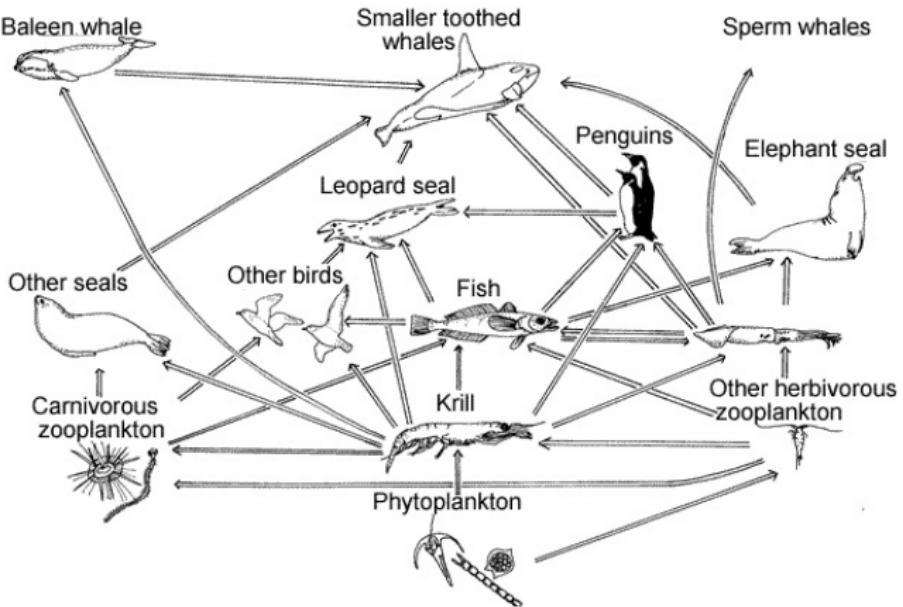
# Graph Clustering

## Basics

- partitioning of a graph into tightly connected groups
- clustering paradigm: **internally dense** and **externally sparse**
- quality measures: many, most popular: **modularity**
- modularity maximization is NP-hard  $\Rightarrow$  use heuristics



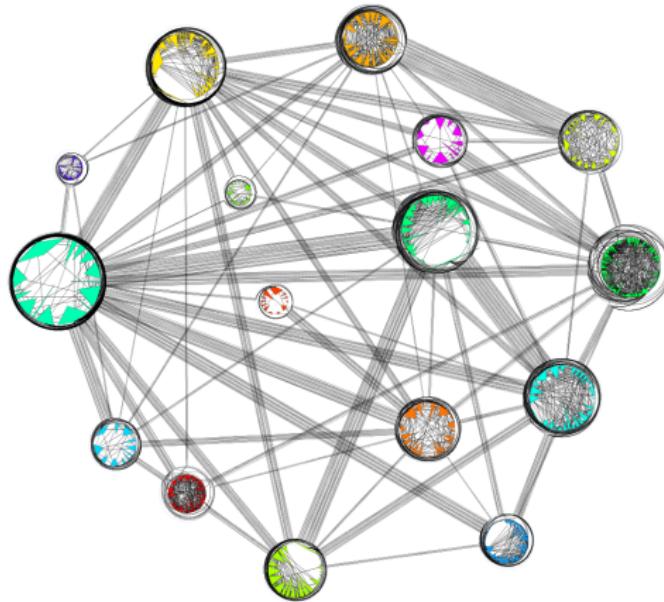
# Clustering Applications



A simplified ecosystem: the antarctic food web  
(source: Antarctica)

cluster  $\approx$  self-sustaining / indivisible subsystem

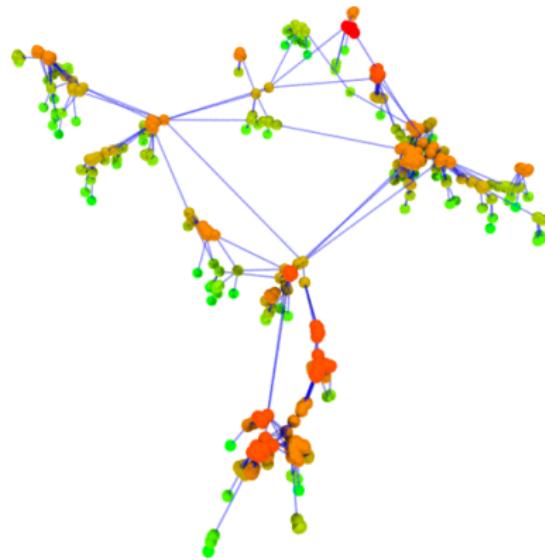
# Clustering Applications



Company-internal email traffic, groups are departments

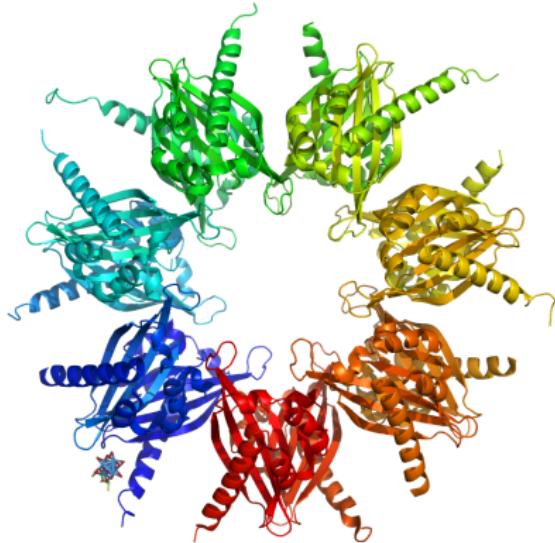
(source: )

# Clustering Applications



Excerpt of the network of Amazon recommendations, around  
"VW Beetle Repairs"  
(source: [Gaertler '07])  
cluster  $\approx$  customer profile

# Clustering Applications



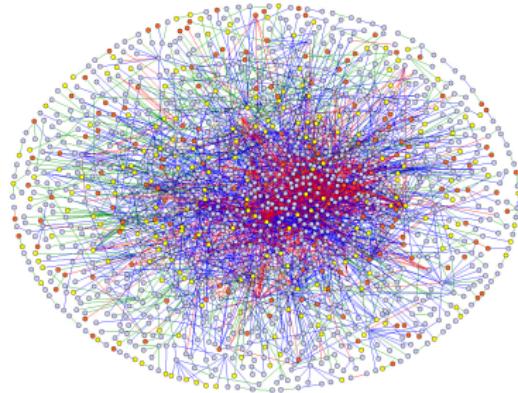
molecular structure of a protein

( $\text{Ca}^{2+}$  / Calmodulin-dependent kinase II (CaMKII))

source: protein database [www.rcsb.org](http://www.rcsb.org))

cluster  $\approx$  functional unit (*domain*) of a protein

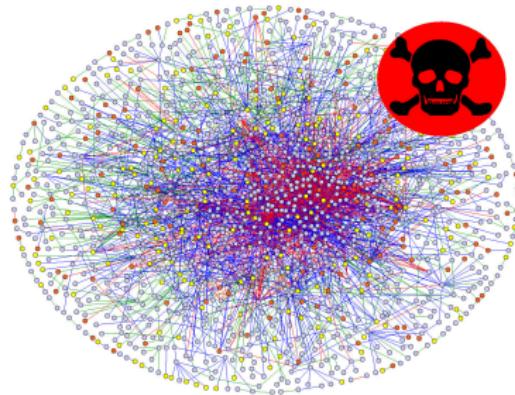
# Clustering Applications



protein interactions

(source: Max-Delbrück-Centre for molecular medicine, [www.mdc-berlin.de](http://www.mdc-berlin.de))

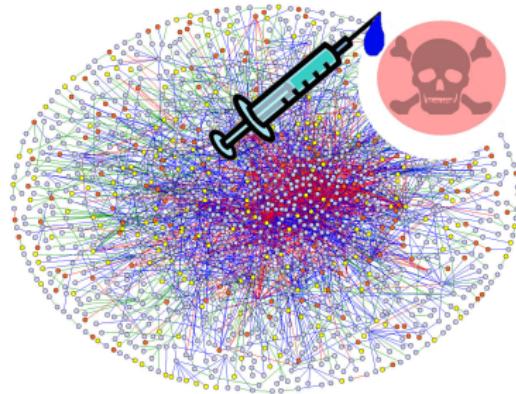
# Clustering Applications



protein interactions

(source: Max-Delbrück-Centre for molecular medicine, [www.mdc-berlin.de](http://www.mdc-berlin.de))

# Clustering Applications

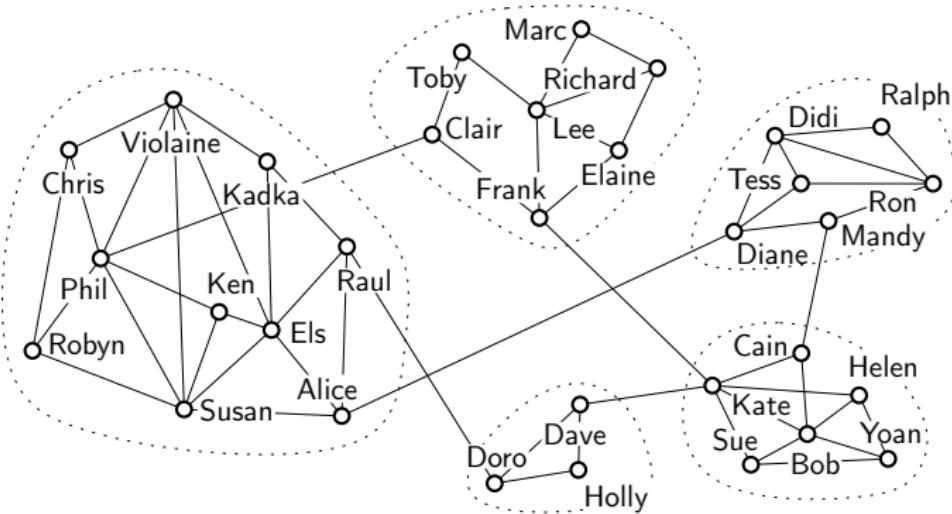


protein interactions

(source: Max-Delbrück-Centre for molecular medicine, [www.mdc-berlin.de](http://www.mdc-berlin.de))

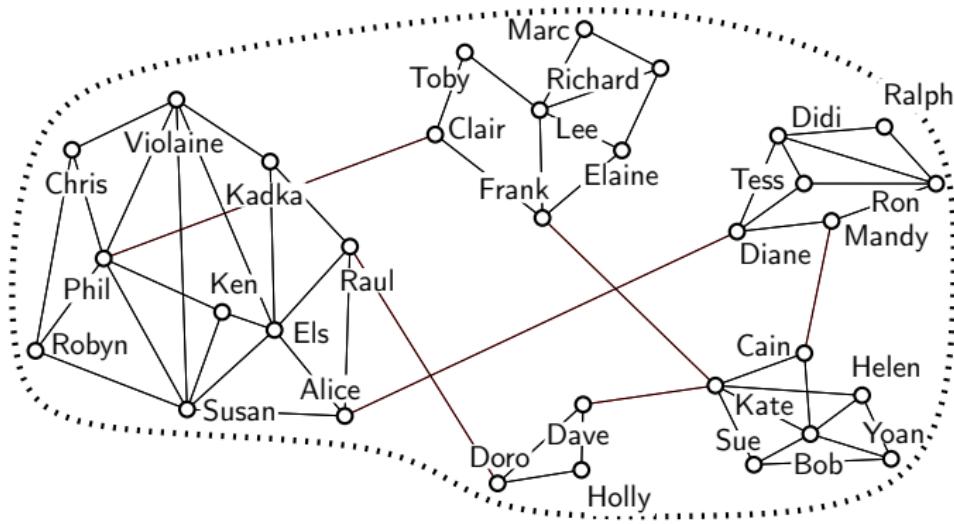
cluster  $\approx$  isolatable seat of disease

# Motivation for Modularity



■  $\text{cov}(\mathcal{C}) = \frac{\# \text{ intra-cluster edges}}{\# \text{ edges}} \approx 0.9$

# Motivation for Modularity



- $\text{cov}(\mathcal{C}) = \frac{\# \text{ intra-cluster edges}}{\# \text{ edges}} \approx 0.9$
- only one cluster  $\Rightarrow \text{cov}(\mathcal{C}') = 1.0$

# A Promising Remedy

Modularity [Girvan and Newman'04]

- subtract from coverage the **expected** value → useful measure

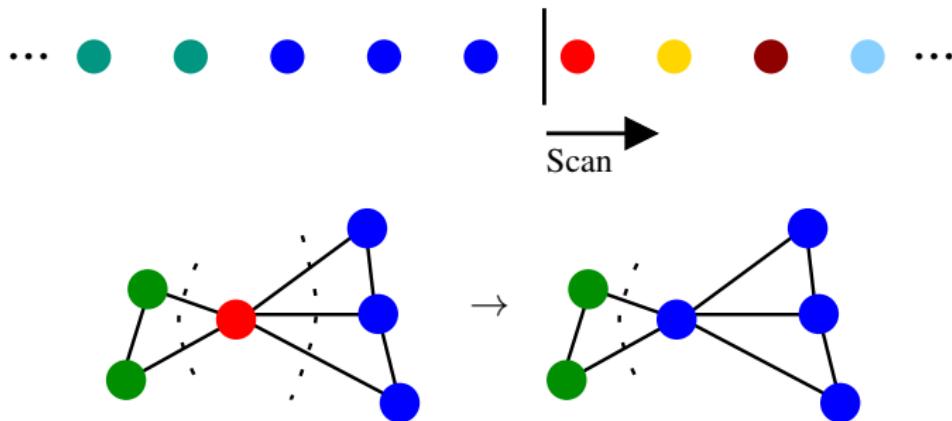
$$\begin{aligned}\text{mod}(\mathcal{C}) &:= \text{cov}(\mathcal{C}) - \mathbb{E}[\text{cov}(\mathcal{C})] \\ &= \frac{m(\mathcal{C})}{m} - \frac{1}{4m^2} \sum_{C \in \mathcal{C}} \left( \sum_{v \in C} \deg(v) \right)^2\end{aligned}$$

- probability model:
  1. random graphs, keep the same clustering
  2. goal: keep *expected* node degrees, randomly throw in edges
  3. edge set multi-set

# Louvain Method

[Blondel et al.'08]

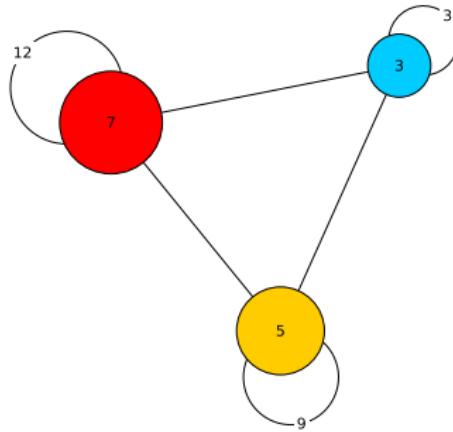
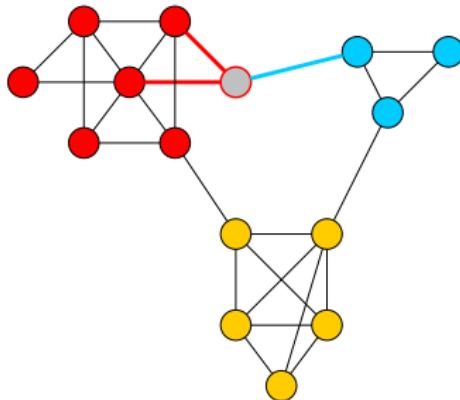
- modularity-based local movement + multi-level clustering
  - start with **singletons**
  - traverse nodes in random order
  - move node to neighboring cluster yielding **highest mod** increase



# Graph Clustering

## Selected Methods: Louvain Method

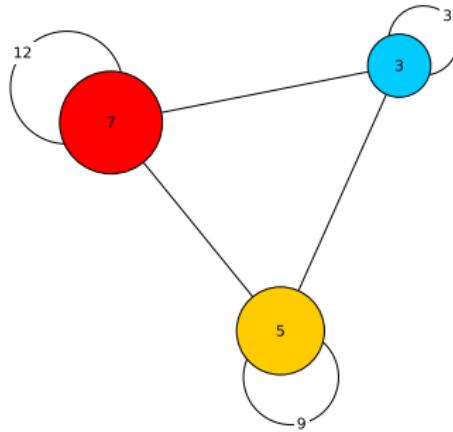
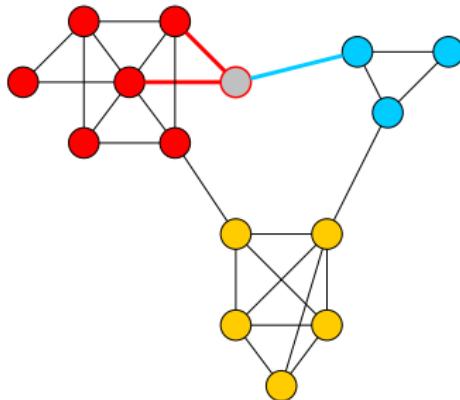
- multilevel modularity maximizing approach
- works in **alternating phases**: local search followed by contraction
- fast and non-deterministic



# Graph Clustering

## Selected Methods: Louvain Method

- multilevel modularity maximizing approach
- works in **alternating phases**: local search followed by contraction
- fast and non-deterministic

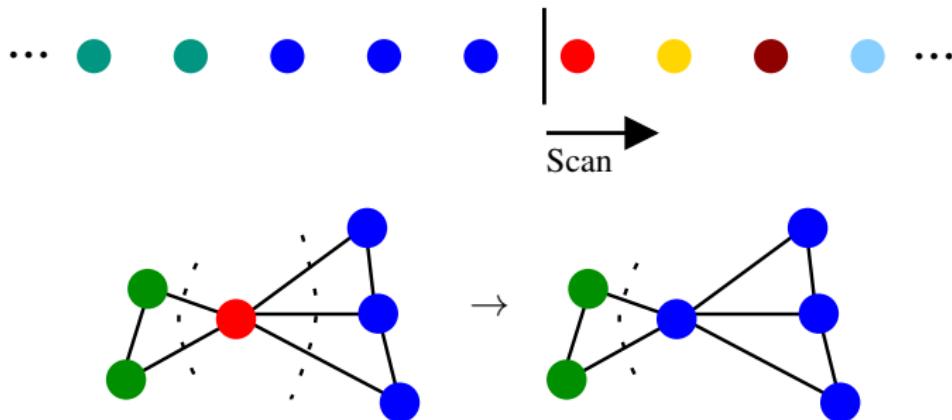


- **uncoarsening**: transfer clustering + local movement every level

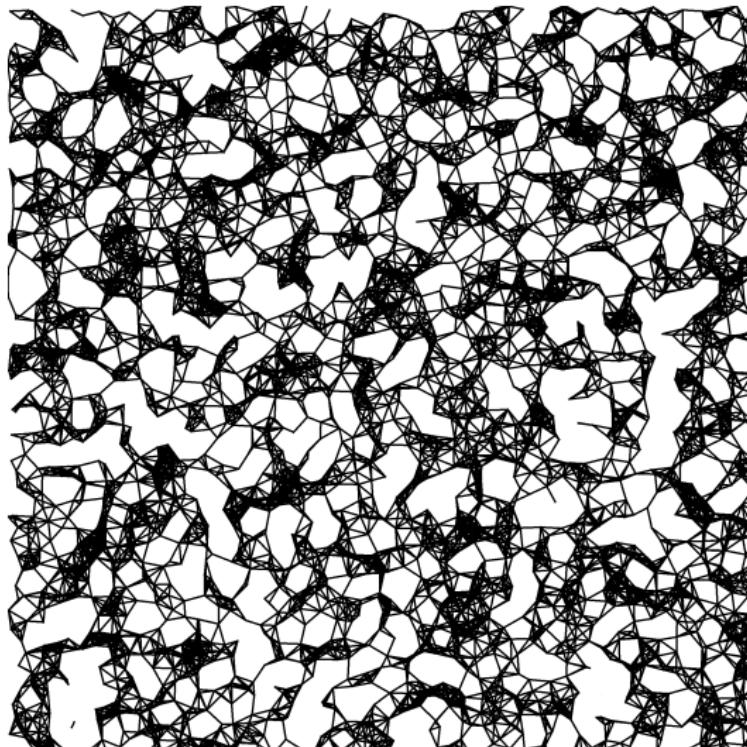
# Label Propagation

Cut-based, Linear Time Clustering Algorithm

- **cut-based** clustering using size-constraint label propagation
  - start with **singletons**
  - traverse nodes in random order or smallest degree first
  - move node to cluster having **strongest eligible** connection
  - *eligible*: w.r.t size constraint  $U$

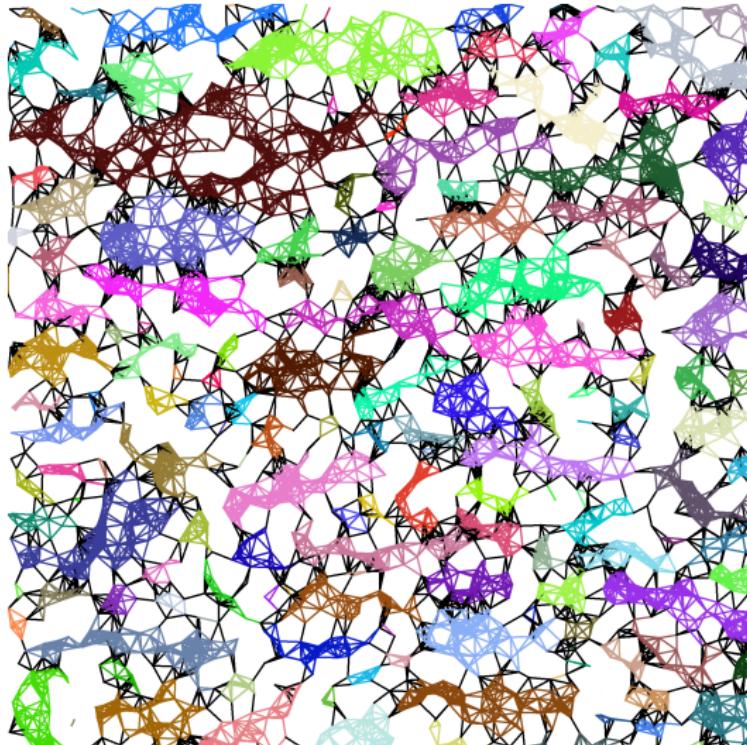


# Label Propagation



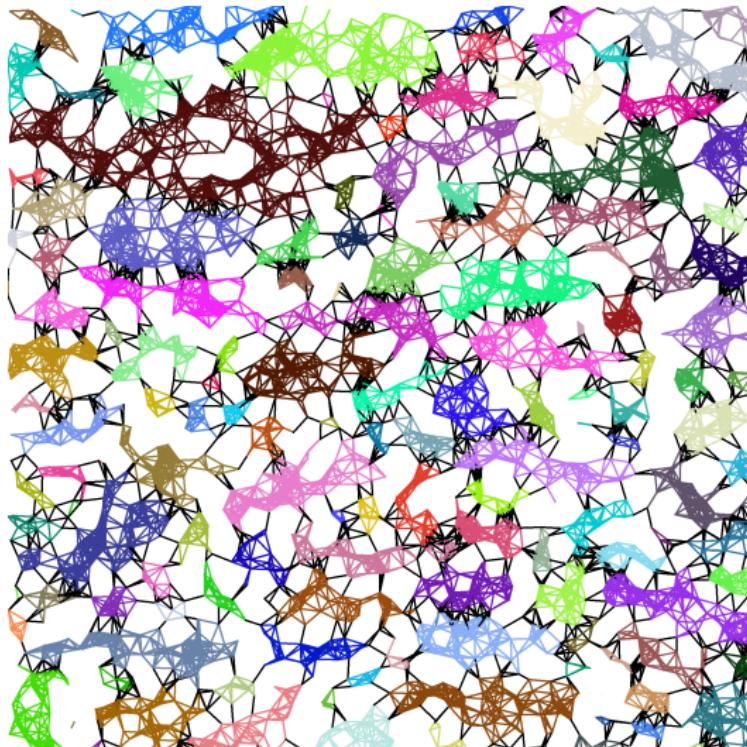
Iteration	Cut [%]
0	100
1	8.96
2	6.15
3	5.66
4	5.44
5	5.28
6	5.25
7	5.21
8	5.18
...	5.09

# Label Propagation



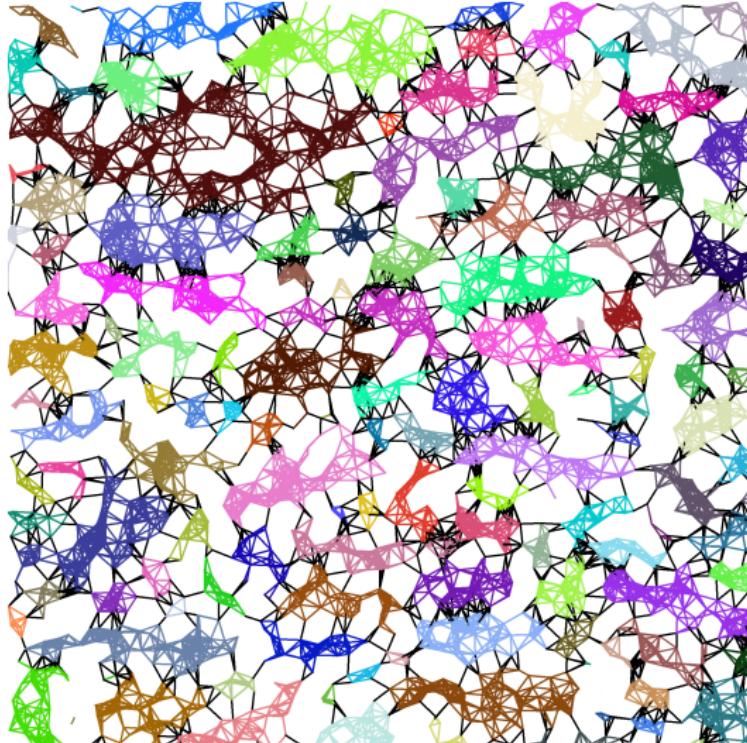
Iteration	Cut [%]
0	100
1	8.96
2	6.15
3	5.66
4	5.44
5	5.28
6	5.25
7	5.21
8	5.18
...	5.09

# Label Propagation



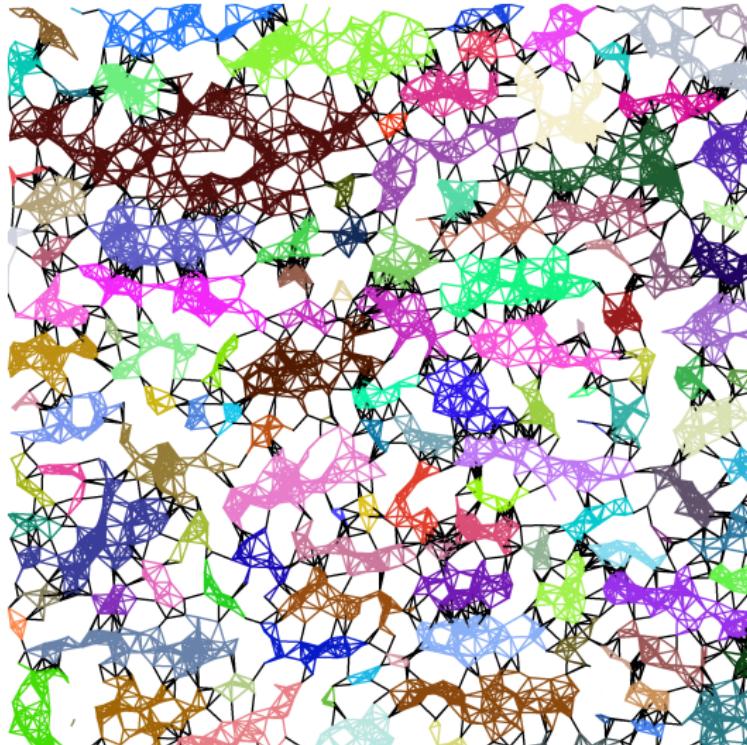
Iteration	Cut [%]
0	100
1	8.96
2	6.15
3	5.66
4	5.44
5	5.28
6	5.25
7	5.21
8	5.18
...	5.09

# Label Propagation



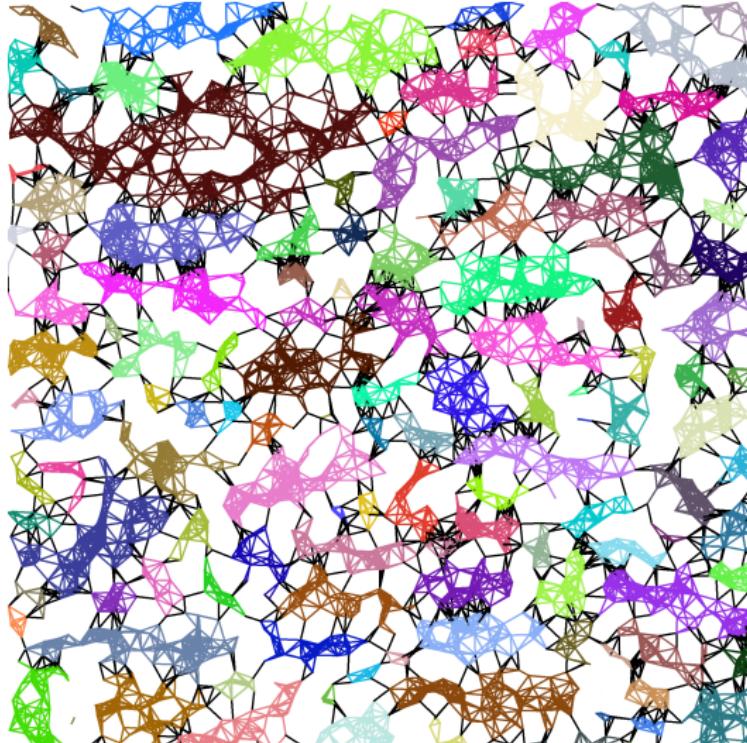
Iteration	Cut [%]
0	100
1	8.96
2	6.15
3	5.66
4	5.44
5	5.28
6	5.25
7	5.21
8	5.18
...	5.09

# Label Propagation



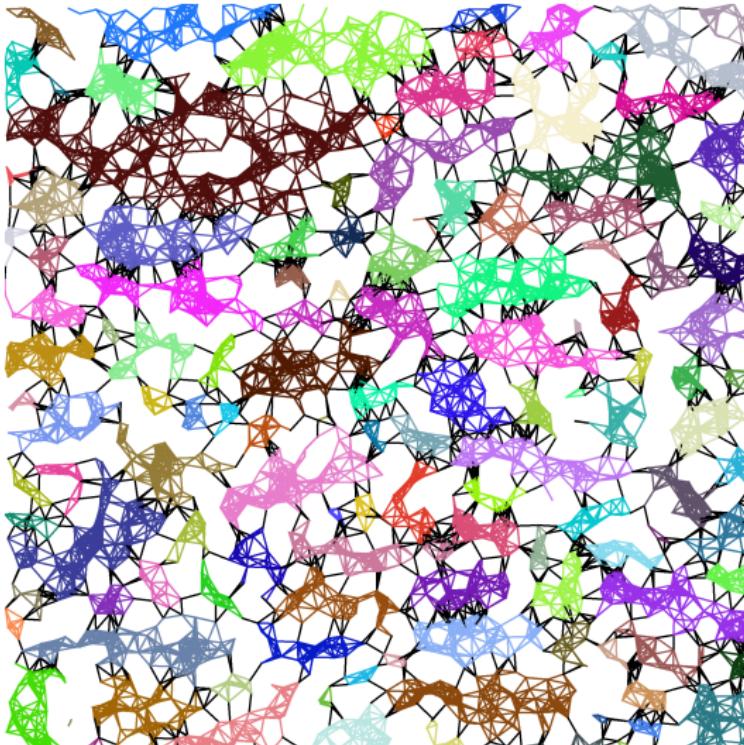
Iteration	Cut [%]
0	100
1	8.96
2	6.15
3	5.66
<b>4</b>	<b>5.44</b>
5	5.28
6	5.25
7	5.21
8	5.18
...	5.09

# Label Propagation



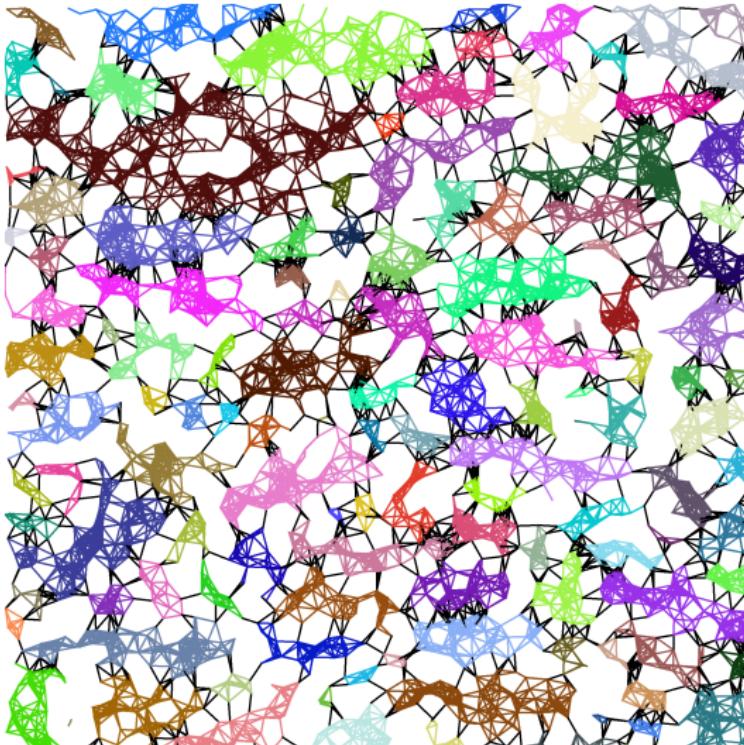
Iteration	Cut [%]
0	100
1	8.96
2	6.15
3	5.66
4	5.44
5	5.28
6	5.25
7	5.21
8	5.18
...	5.09

# Label Propagation



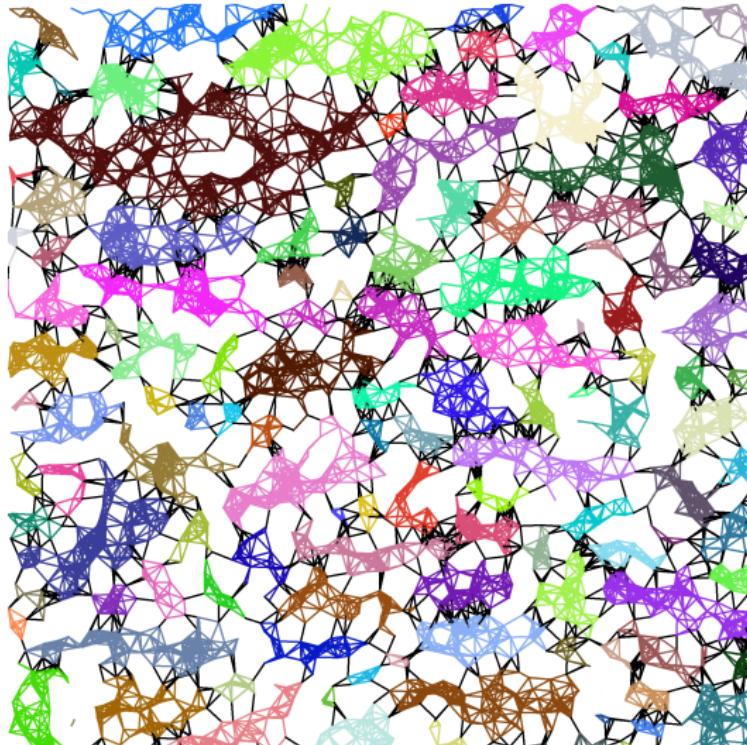
Iteration	Cut [%]
0	100
1	8.96
2	6.15
3	5.66
4	5.44
5	5.28
6	5.25
7	5.21
8	5.18
...	5.09

# Label Propagation



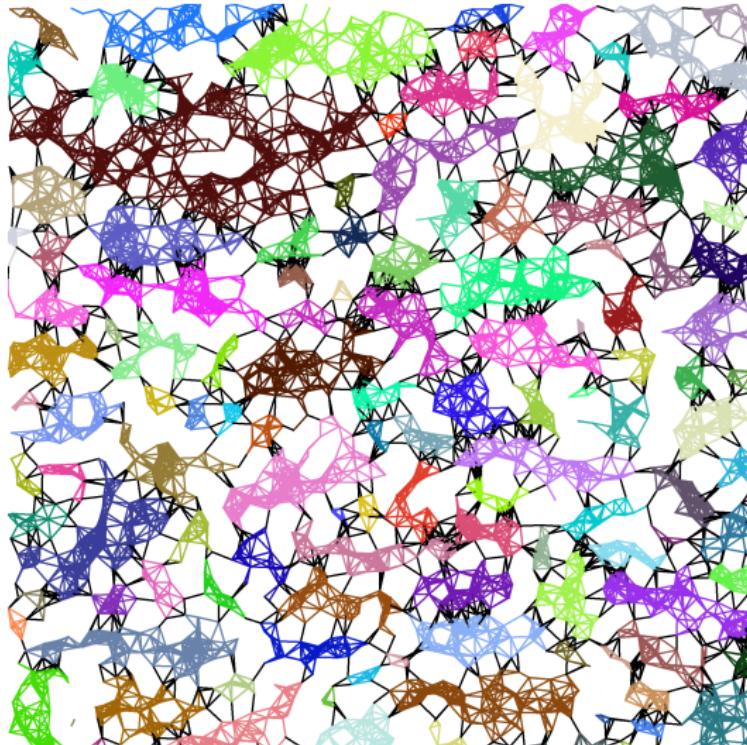
Iteration	Cut [%]
0	100
1	8.96
2	6.15
3	5.66
4	5.44
5	5.28
6	5.25
7	5.21
8	5.18
...	5.09

# Label Propagation



Iteration	Cut [%]
0	100
1	8.96
2	6.15
3	5.66
4	5.44
5	5.28
6	5.25
7	5.21
8	5.18
...	5.09

# Label Propagation

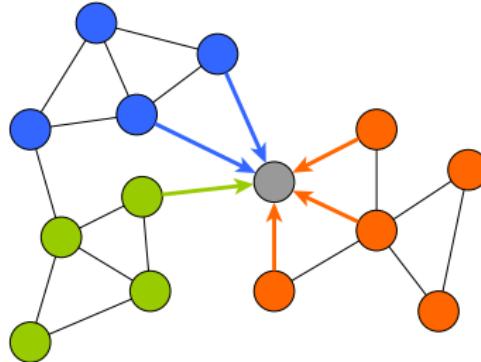


Iteration	Cut [%]
0	100
1	8.96
2	6.15
3	5.66
4	5.44
5	5.28
6	5.25
7	5.21
8	5.18
...	5.09

# Graph Clustering

Selected Methods: Louvain Method with Preprocessing

- fast preprocessing step:  
run label propagation to obtain a preliminary clustering
- speeds up Louvain considerably on large instances
- results in more diverse clusterings between runs



# Memetic Graph Clustering

## Introduction: VieClus

- memetic algorithms (MA): genetic algorithms (GA) + local search
- quality measure agnostic: focus **modularity**
- GAs do like nature does:
  - solutions selected from population by fitness (quality)
  - solutions are combined → offspring
  - offspring may be mutated
  - offspring inserted into population



*Natural selection does not grant organisms what they "need".*

# Memetic Graph Clustering

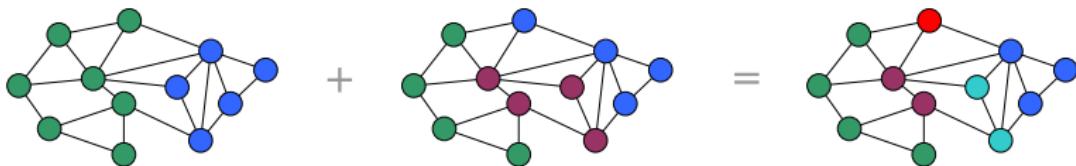
## Initial Population and Recombination Mechanism

### Initial Population

- initial population generated by Louvain method (random seeds)
- both with and without Label Propagation preprocessing
- results in a reasonably **diverse base population**

### Recombination Mechanism

- borrowed from ML: ensemble learning → **maximum overlap**



→ results in **overlap clustering**

# Memetic Graph Clustering

## Recombination Operators

- two kinds: flat and multilevel
- fundamental aim: **high diversification**
- diversity is good: prevents premature convergence

### Flat Recombination

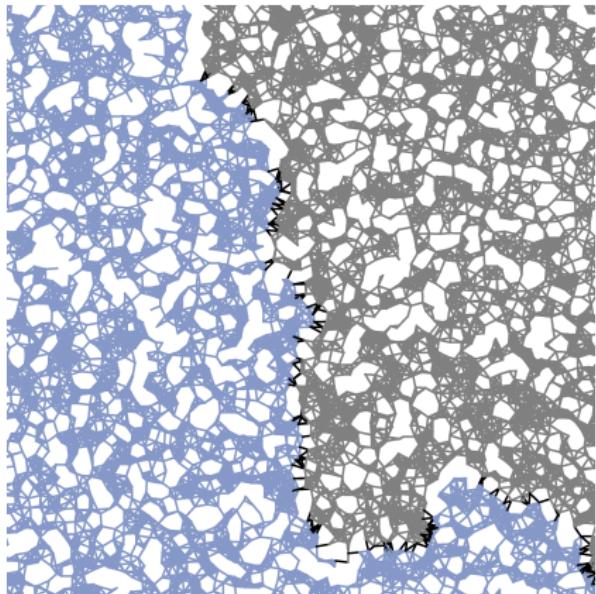
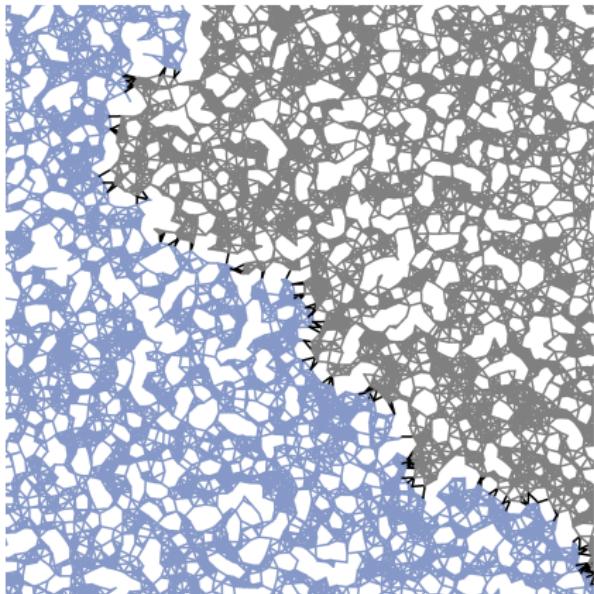
1. calculate overlap of two clusterings
2. contract overlap
3. run Louvain method on the contracted overlap

### Variations on Flat Recombination

- start Louvain with the better parent clustering
- generate second parent using graph partitioning or LP
- this introduces **a lot of diversity**

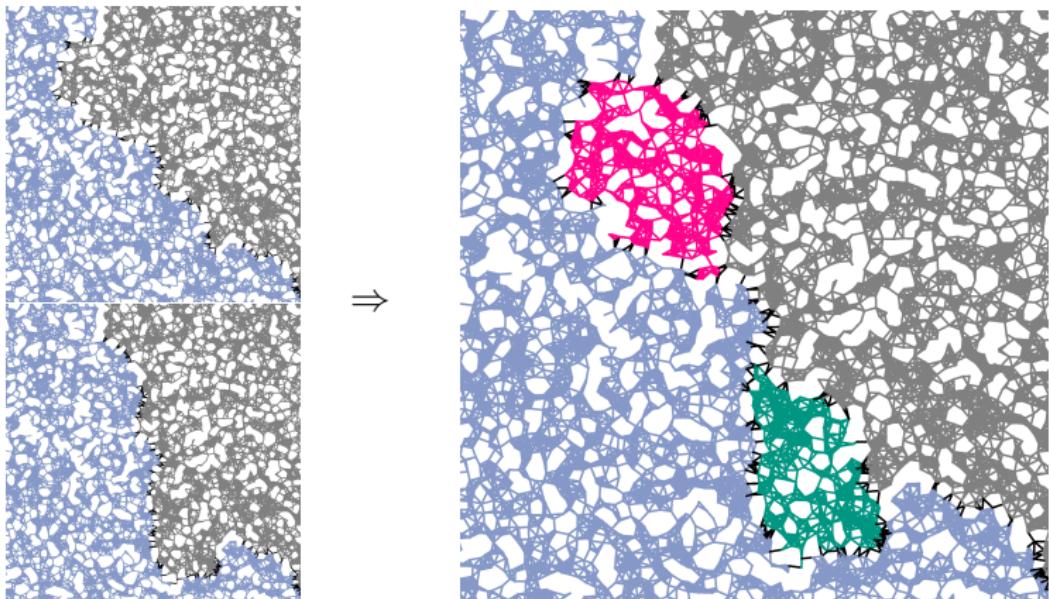
# Example

Two Individuals  $\mathcal{P}_1, \mathcal{P}_2$



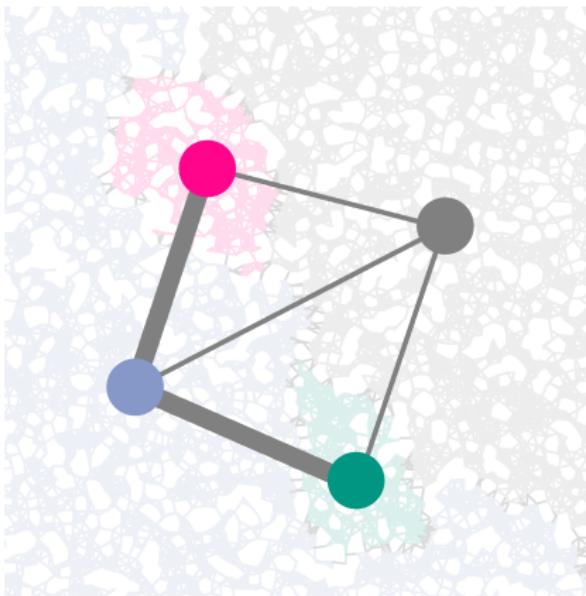
# Example

## Overlay of $\mathcal{P}_1, \mathcal{P}_2$



# Example

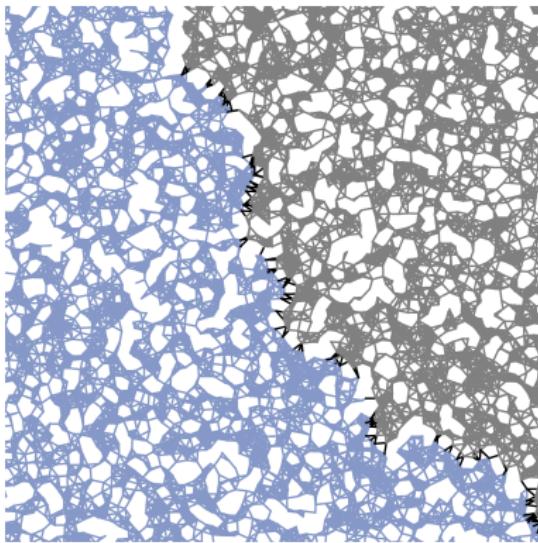
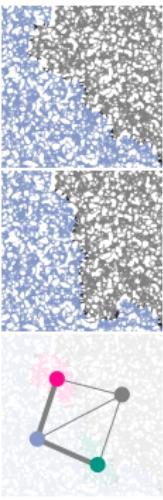
Flat Recombination of  $\mathcal{P}_1, \mathcal{P}_2$



- run Louvain on this graph
- variation: start with better parent clustering

# Example

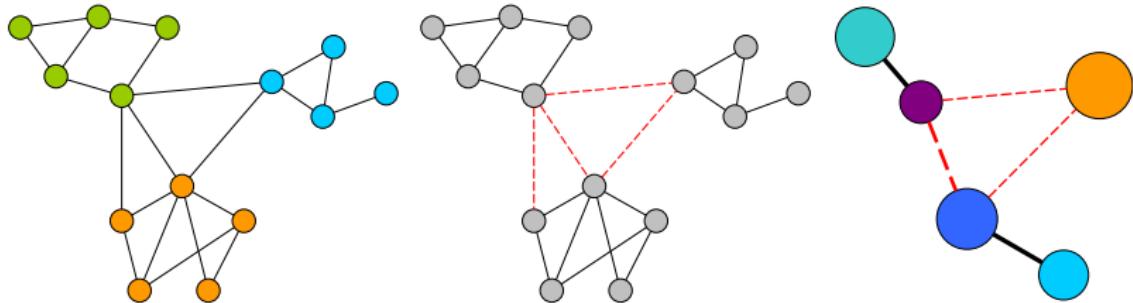
Result of  $\mathcal{P}_1, \mathcal{P}_2$



# Memetic Graph Clustering

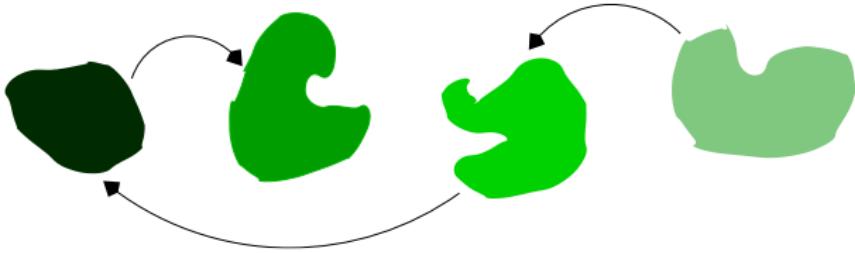
## Recombination Operator: Multilevel

- run the Louvain method, but forbid LS to consider cut edges  
→ cut edges are **not contracted**
- allows for finer, iterative changes compared to running Louvain on the contracted overlap



- coarsest graph is the contracted overlap,  
contraction is usually stopped earlier

# Parallelization



- each PE has its own **island** (a local population)
  - **locally**: perform combine and mutation operations
  - communicate analog to *randomized rumor spreading*
    1. **rumor**  $\leftrightarrow$  currently **best** local clustering
    2. local best partition *changed*  $\rightarrow$  send it to  $\mathcal{O}(\log P)$  random PEs
    3. **asynchronous** communication (MPI Isend)
- **quality records in a few minutes** for small graphs

# Benchmark Instances

Graph	n	m
Small Instances		
celegans_metabolic	453	2025
email	1133	5451
polblogs	1490	16715
power	4941	6594
PGPgiantcompo	10680	24316
astro-ph	16706	121251
memplus	17758	54196
as-22july06	22963	48436
cond-mat-2005	40421	175691
G.pin.pout	100000	501198
smallworld	100000	499998
luxembourg.osm	114599	119666
rgg_n_2_17_s0	131072	728753
caidaRouterLevel	192244	609066
Large Instances		
prefAttachment	100000	499985
coAuthorsCiteseer	227320	814134
citationCiteseer	268495	1156647
coPapersDBLP	540486	15245729
eu-2005	862664	16138468
ldoor	952203	22785136
in-2004	1382908	13591473
belgium.osm	1441295	1549970
333SP	3712815	11108633

- 10th DIMACS implementation challenge [competition instances](#)

# Results

## Small instances

Graph	$\bar{t}$ [m]	Max. $\mathcal{Q}$	DIMACS $\mathcal{Q}$	$t_{\text{sol}}$ [m]	Solver
as-22july06	< 1	<b>0.679 396</b>	0.678 267	6.6	CGGC
astro-ph	< 1	<b>0.746 292</b>	0.744 621	11.9	VNS
celegans_metabol	< 1	0.453 248	0.453 248	< 1	VNS
cond-mat-2005	< 1	<b>0.750 171</b>	0.746 254	40.9	CGGC
email	< 1	0.582 829	0.582 829	< 1	VNS
PGPgiantcompo	< 1	<b>0.886 853</b>	0.886 564	1.9	CGGC
polblogs	< 1	0.427 105	0.427 105	< 1	VNS
power	< 1	<b>0.940 977</b>	0.940 851	< 1	VNS
smallworld	< 1	<b>0.793 187</b>	0.793 042	16.8	VNS
memplus	2.6	<b>0.701 275</b>	0.700 473	3.2	CGGC
G_n_pin_pout	3.9	<b>0.500 466</b>	0.500 098	64.8	CGGC
caidaRouterLevel	5.0	<b>0.872 828</b>	0.872 042	81.0	CGGC
rgg_n17	5.0	<b>0.978 454</b>	0.978 324	37.5	VNS
luxembourg.osm	7.3	<b>0.989 672</b>	0.989 621	40.9	VNS

- $\bar{t}$  [m]: time in minutes to reach DIMACS result
- $t_{\text{sol}}$  [m]: time in minutes to reach DIMACS result of prev. solver

# Results

## Large Instances

Graph	$\bar{t}$ [m]	Max. $\mathcal{Q}$	DIMACS $\mathcal{Q}$	$t_{\text{sol}}$ [m]	Solver
coAuthorsCiteseer	3.9	<b>0.906 830</b>	0.905 297	91.3	CGGC
citationCiteseer	12.9	<b>0.825 545</b>	0.823 930	77.6	CGGC
coPapersDBLP	20.5	<b>0.868 058</b>	0.866 794	603.3	CGGC
belgium.osm	29.5	<b>0.995 064</b>	0.994 940	102.9	CGGC
ldoor	35.1	<b>0.970 555</b>	0.969 370	485.6	ParMod
eu-2005	65.8	<b>0.941 575</b>	0.941 554	341.5	CGGC
in-2004	237.4	<b>0.980 690</b>	0.980 622	244.0	CGGC
333SP	297.1	<b>0.989 356</b>	0.989 095	976.9	ParMod
prefAttachment	*	<b>0.316 089</b>	0.315 994	1 353.1	VNS

- $\bar{t}$  [m]: time in minutes to reach DIMACS result
- $t_{\text{sol}}$  [m]: time in minutes to reach DIMACS result of prev. solver

# More Results

- better than all solvers reported in the recent literature
    - Karypis and LaSalle
    - Lu et al.
    - Ryu and Kim
    - Džamić et al. (ascent-decent VNS)
    - ...
- paper

Vienna Graph Clustering:  
<https://github.com/VieClus/VieClus>



# In Depth Example

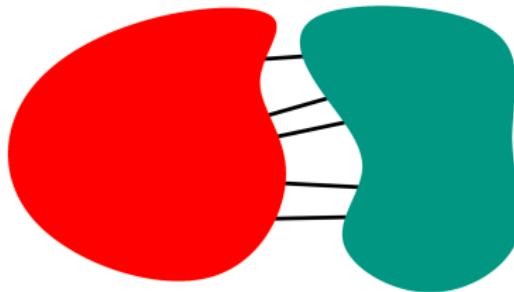
## “(In)exact Reductions” in Minimum Cuts

# Minimum Cuts

**Cut:** A **cut** in a multigraph is a partition of  $V = C \cup \bar{C}$   
→ size of the cut is weight of edges between  $C$  and  $\bar{C}$

**Minimum Cut Problem:**

what is the size of the minimum cut in  $G$ ?



# Basics

If the size of the minimum cut is  $\lambda$ , then it follows

- $\forall v \in V : \deg(v) \geq \lambda$
- number of edges  $m \geq n\lambda/2$

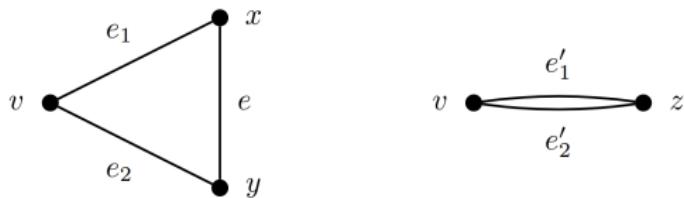
**Proof:** Assume  $\exists v \in V : \deg(v) < \lambda$ ,  
then  $C = \{v\}$  is a cut whose size is  $< \lambda$ . Contradiction.  
The second claim follows from the first one.

# Contraction

In a multigraph  $G$ , let  $u$  and  $v$  be connected by an edge  $e = \{x, y\}$

Create  $G/e = (V', E')$  by **contracting**  $e$ :

- set  $V'$  to  $V \setminus \{x, y\} \cup \{z\}$  ( $z$  is new)
- build  $E'$  from  $E$  by
  - remove all edges between  $u$  and  $v$
  - replace every edge between  $v \in V \setminus \{x, y\}$  and  $x$  or  $y$  by an edge between  $v$  and  $z$
  - keep all other edges from  $E$



→ multi-edges can be created ( $\rightsquigarrow$  practice use weights)!

# Minimum Cut $\leftrightarrow$ Contraction

A minimum cut in  $G/e$  is at least as large as a minimum cut in  $G$ .

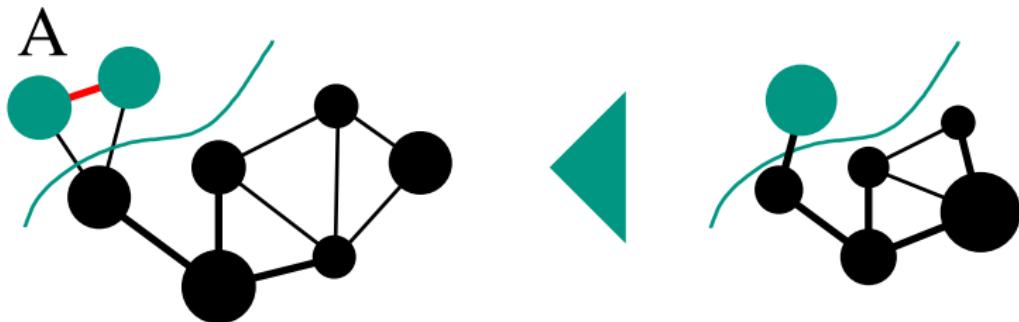
**Proof:**

Let  $(K, \bar{K})$  be a minimum cut in  $G/e$ .

Let the size of the cut be  $\lambda$ .

Wlog let  $x$  and  $y$  be the vertices of  $e$ , and  $z \in K$

Unpack  $z$  and leave  $x$  and  $y$  in  $K \rightarrow$  cut in  $G$  of size  $\lambda$



# Algorithm

## Exhibit A

$H \leftarrow G$

**while**  $H$  has more than 2 nodes **do**

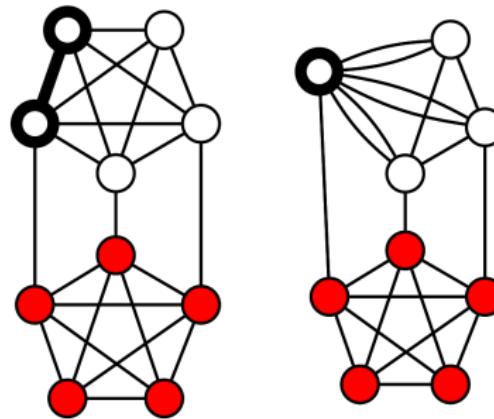
$e \leftarrow$  edge of  $H$  picked uniformly at random

$H \leftarrow \text{contract}(H, e)$

**done**

$(C, \bar{C}) \leftarrow$  vertex set in  $G$  that correspond to the vertices in  $H$

---



# Algorithm

## Exhibit A

$H \leftarrow G$

**while**  $H$  has more than 2 nodes **do**

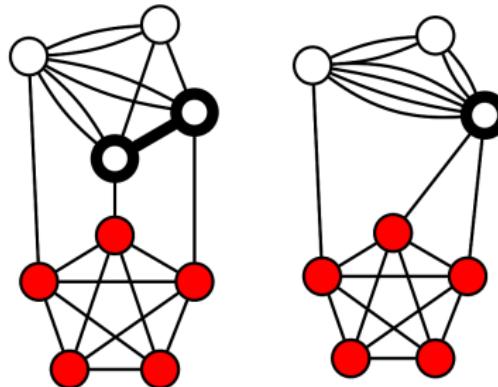
$e \leftarrow$  edge of  $H$  picked uniformly at random

$H \leftarrow \text{contract}(H, e)$

**done**

$(C, \bar{C}) \leftarrow$  vertex set in  $G$  that correspond to the vertices in  $H$

---



# Algorithm

## Exhibit A

$H \leftarrow G$

**while**  $H$  has more than 2 nodes **do**

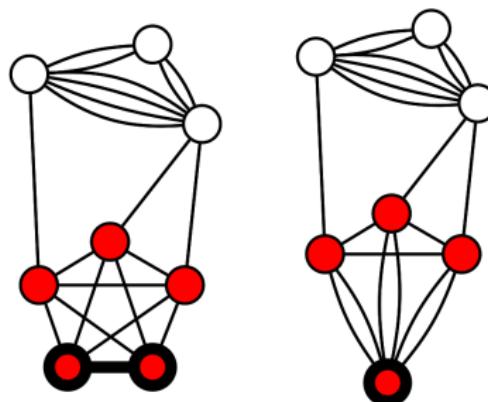
$e \leftarrow$  edge of  $H$  picked uniformly at random

$H \leftarrow \text{contract}(H, e)$

**done**

$(C, \bar{C}) \leftarrow$  vertex set in  $G$  that correspond to the vertices in  $H$

---



# Algorithm

## Exhibit A

$H \leftarrow G$

**while**  $H$  has more than 2 nodes **do**

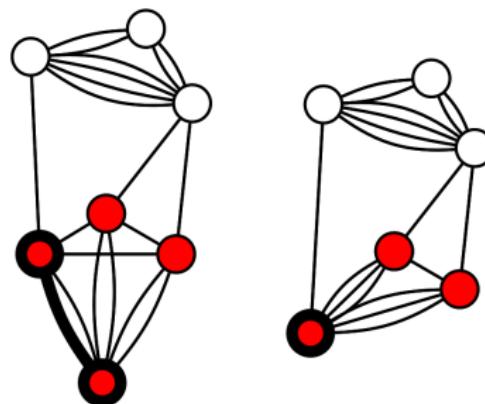
$e \leftarrow$  edge of  $H$  picked uniformly at random

$H \leftarrow \text{contract}(H, e)$

**done**

$(C, \bar{C}) \leftarrow$  vertex set in  $G$  that correspond to the vertices in  $H$

---



# Algorithm

## Exhibit A

$H \leftarrow G$

**while**  $H$  has more than 2 nodes **do**

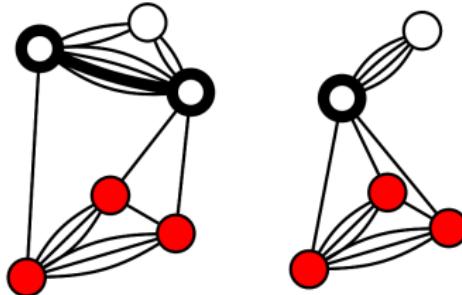
$e \leftarrow$  edge of  $H$  picked uniformly at random

$H \leftarrow \text{contract}(H, e)$

**done**

$(C, \bar{C}) \leftarrow$  vertex set in  $G$  that correspond to the vertices in  $H$

---



# Algorithm

## Exhibit A

$H \leftarrow G$

**while**  $H$  has more than 2 nodes **do**

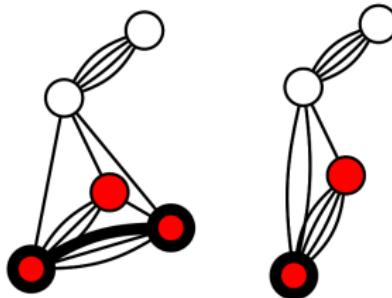
$e \leftarrow$  edge of  $H$  picked uniformly at random

$H \leftarrow \text{contract}(H, e)$

**done**

$(C, \bar{C}) \leftarrow$  vertex set in  $G$  that correspond to the vertices in  $H$

---



# Algorithm

## Exhibit A

$H \leftarrow G$

**while**  $H$  has more than 2 nodes **do**

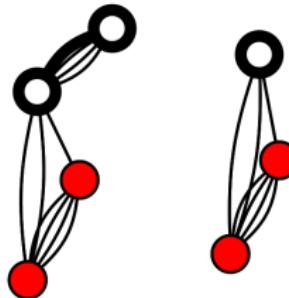
$e \leftarrow$  edge of  $H$  picked uniformly at random

$H \leftarrow \text{contract}(H, e)$

**done**

$(C, \bar{C}) \leftarrow$  vertex set in  $G$  that correspond to the vertices in  $H$

---



# Algorithm

## Exhibit A

$H \leftarrow G$

**while**  $H$  has more than 2 nodes **do**

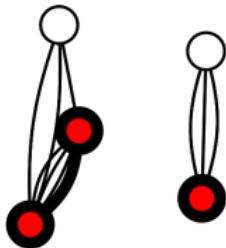
$e \leftarrow$  edge of  $H$  picked uniformly at random

$H \leftarrow \text{contract}(H, e)$

**done**

$(C, \bar{C}) \leftarrow$  vertex set in  $G$  that correspond to the vertices in  $H$

---



# Algorithm

## Exhibit A

$H \leftarrow G$

**while**  $H$  has more than 2 nodes **do**

$e \leftarrow$  edge of  $H$  picked uniformly at random

$H \leftarrow \text{contract}(H, e)$

**done**

$(C, \bar{C}) \leftarrow$  vertex set in  $G$  that correspond to the vertices in  $H$

---

The runtime of the simple minimum cut algorithm is  $O(n^2)$

**Proof:**

- every call  $\text{contract}(H, e)$  is done in  $O(n)$
- every loop iteration reduces  $n$  by 1  $\rightarrow n - 2$  iterations

# Algorithm

## Exhibit A

$H \leftarrow G$

**while**  $H$  has more than 2 nodes **do**

$e \leftarrow$  edge of  $H$  picked uniformly at random

$H \leftarrow \text{contract}(H, e)$

**done**

$(C, \bar{C}) \leftarrow$  vertex set in  $G$  that correspond to the vertices in  $H$

---

The algorithm finds a minimum cut with probability  $\Omega(n^{-2})$

**Proof (sketch):**

- let minimum cut size be  $\lambda$
- probability to select a cut edge  $\frac{\lambda}{|E|} \leq \frac{\lambda}{n\lambda/2} = 2/n$
- $p_n$  probability that  $n$ -vertex graph avoids cut edges

$$p_n \geq (1 - 2/n)p_{n-1} \geq \dots = \binom{n}{2}^{-1}$$

# Algorithm

## Exhibit A

$H \leftarrow G$

**while**  $H$  has more than 2 nodes **do**

$e \leftarrow$  edge of  $H$  picked uniformly at random

$H \leftarrow \text{contract}(H, e)$

**done**

$(C, \bar{C}) \leftarrow$  vertex set in  $G$  that correspond to the vertices in  $H$

---

### Standard Trick: Multiple Repetitions

- non-error probability  $1/n^2$  very low
- smallest out of  $n^2/2$  is minimum with probability  $1 - 1/e$ :

$$(1 - 2/n^2)^{n^2/2} < 1/e$$

$\rightsquigarrow$  runtime  $O(n^4)$

# Better Algorithm

IterContract

$H \leftarrow G$

**while**  $H$  has more than  $t$  nodes **do**

$e \leftarrow$  edge of  $H$  picked uniformly at random  
 $H \leftarrow \text{contract}(H, e)$

**done**

**return**  $H$

$H$  still contains minimum cut with probability at least

$$\binom{t}{2} / \binom{n}{2}$$

# Karger-Stein

**if**  $|V| \leq 6$  **then**  $C \leftarrow$  optimial cut by deterministic algorithm  
**else**

$t \leftarrow \lceil 1 + n / \sqrt{2} \rceil$

$H_1 \leftarrow \text{IterContract}(G, t)$

$H_2 \leftarrow \text{IterContract}(G, t)$

$C_1 \leftarrow \text{CallRecursive}(H_1)$

$C_2 \leftarrow \text{CallRecursive}(H_2)$

$C \leftarrow \min(C_1, C_2)$

**done**

**return**  $C$

$\rightsquigarrow$  running time  $O(n^2 \log n)$

$\rightsquigarrow$  minimum cut with probability  $\Omega(1 / \log n)$

$\rightsquigarrow$  repeat  $\log^2 n$  to error probability  $O(1/n)$

# Main Questions

how can kernelization help?

# Main Questions

something better than contracting random edges?

# Main Questions

can we still obtain good cuts in practice?

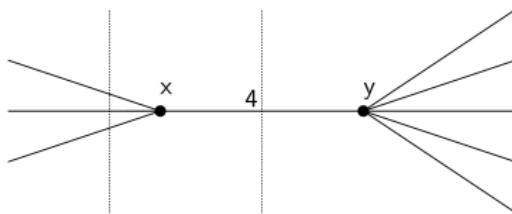
# Main Questions

can we then use this to obtain better kernels?

# Kernelization

## Padberg-Rinaldi Tests

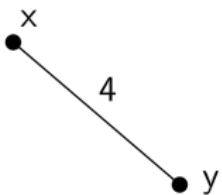
$$\deg(x) \leq 2\omega(x, y)$$



# Kernelization

## Padberg-Rinaldi Tests

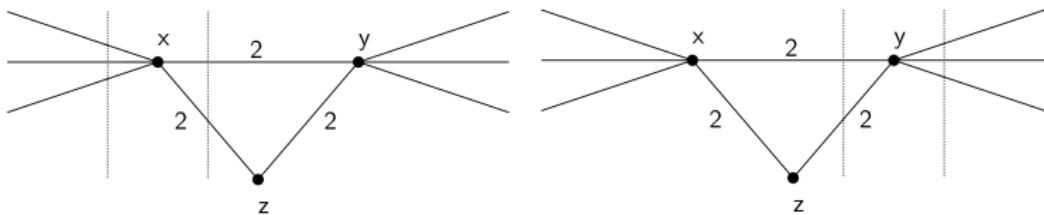
$$\omega(x, y) \geq \hat{\lambda}$$



# Kernelization

## Padberg-Rinaldi Tests

$\exists z : \deg(x) \leq 2\{\omega(x, y) + \omega(x, z)\}$  and  $\deg(y) \leq 2\{\omega(x, y) + \omega(y, z)\}$



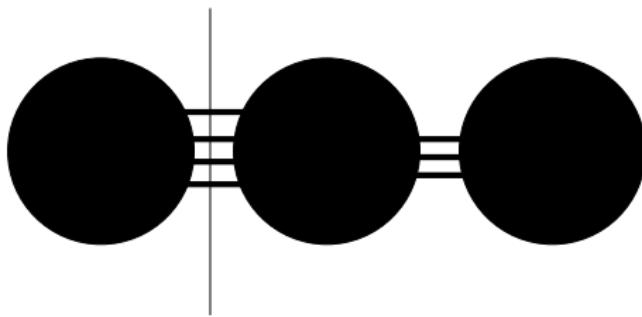
# Kernelization

Nagamochi, Ono, and Ibaraki

**Key Idea:** a spanning tree contains at least one edge from any cut

- Let  $\hat{\lambda}$  be your current bound for minimum cut
- Want: smaller minimum cut
- Compute  $\hat{\lambda} - 1$  maximal spanning forests (iteratively)
  - ~~ edges not in forests connect vertices with connectivity  $\geq \hat{\lambda}$
  - ~~ contract all of them

Example:  $\hat{\lambda} = 4$  ~~ compute 3 edge-disjoint **spanning forests**



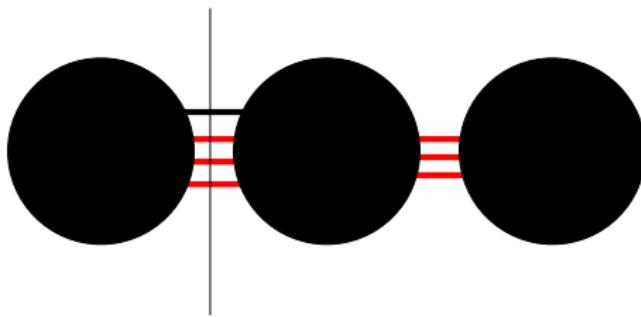
# Kernelization

Nagamochi, Ono, and Ibaraki

**Key Idea:** a spanning tree contains at least one edge from any cut

- Let  $\hat{\lambda}$  be your current bound for minimum cut
- Want: smaller minimum cut
- Compute  $\hat{\lambda} - 1$  maximal spanning forests (iteratively)
  - ~~ edges not in forests connect vertices with connectivity  $\geq \hat{\lambda}$
  - ~~ contract all of them

Example:  $\hat{\lambda} = 4$  ~~ compute 3 edge-disjoint **spanning forests**



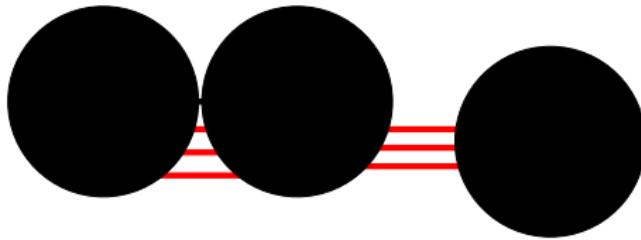
# Kernelization

Nagamochi, Ono, and Ibaraki

**Key Idea:** a spanning tree contains at least one edge from any cut

- Let  $\hat{\lambda}$  be your current bound for minimum cut
- Want: smaller minimum cut
- Compute  $\hat{\lambda} - 1$  maximal spanning forests (iteratively)
  - ~~ edges not in forests connect vertices with connectivity  $\geq \hat{\lambda}$
  - ~~ contract all of them

Example:  $\hat{\lambda} = 4$  ~~ compute 3 edge-disjoint **spanning forests**



# Kernelization

Nagamochi, Ono, and Ibaraki

**Key Idea:** a spanning tree contains at least one edge from any cut

- Let  $\hat{\lambda}$  be your current bound for minimum cut
- Want: smaller minimum cut
- Compute  $\hat{\lambda} - 1$  maximal spanning forests (iteratively)
  - ~~ edges not in forests connect vertices with connectivity  $\geq \hat{\lambda}$
  - ~~ contract all of them

NOI define modified BFS to detect contractable edges  
(more later)

# Kernelization

Nagamochi, Ono, and Ibaraki

**Key Idea:** a spanning tree contains at least one edge from any cut

- Let  $\hat{\lambda}$  be your current bound for minimum cut
- Want: smaller minimum cut
- Compute  $\hat{\lambda} - 1$  maximal spanning forests (iteratively)
  - ~~ edges not in forests connect vertices with connectivity  $\geq \hat{\lambda}$
  - ~~ contract all of them

**Note:** initial  $\hat{\lambda}$  comes from minimum degree

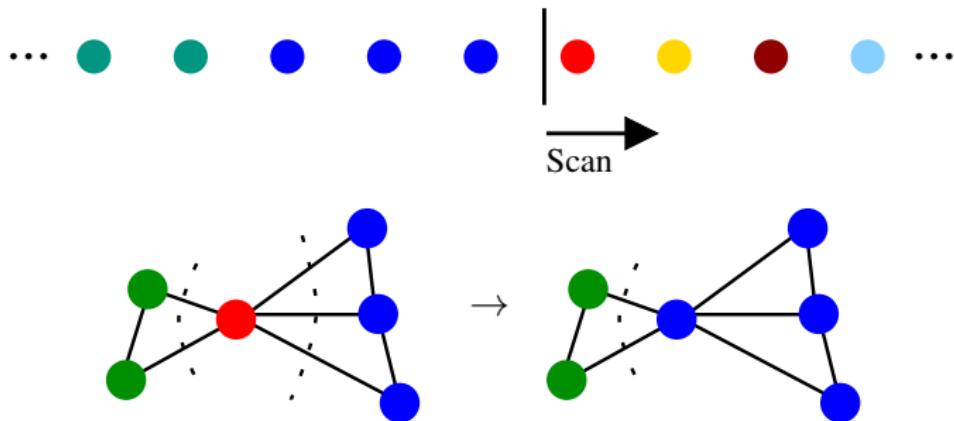
Some of the reductions depend **heavily** on  $\hat{\lambda}$

# Label Propagation

Cut-based, Linear Time Clustering Algorithm [Raghavan et. al]

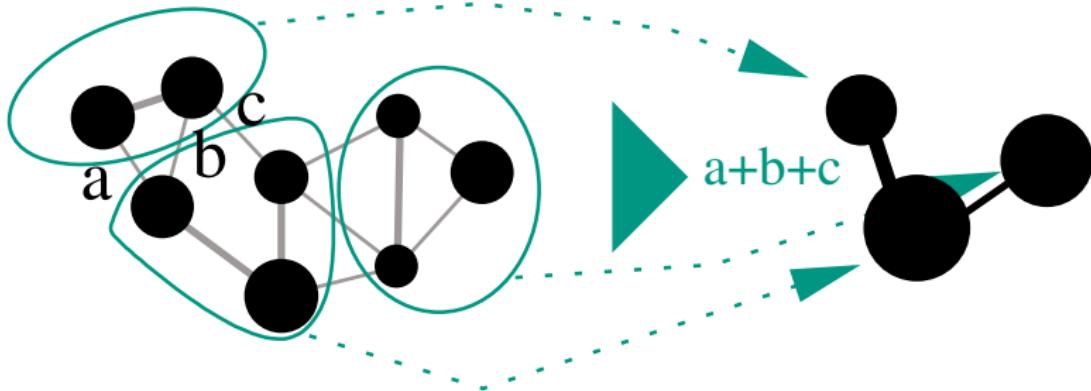
- **cut-based** clustering using label propagation

- start with **singletons**
- traverse nodes in random order or **smallest degree first**
- move node to cluster having **strongest** connection



# Basic Idea

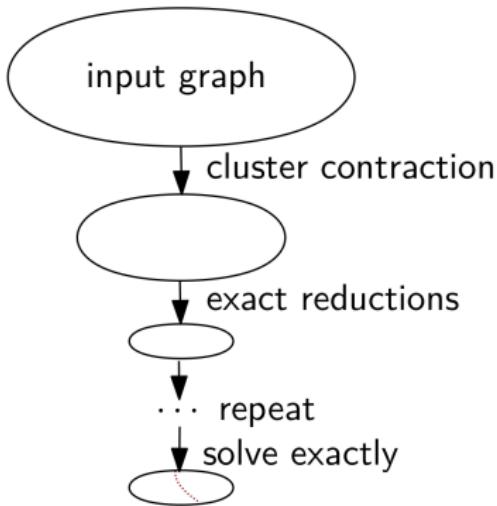
## Contraction of Clusterings



- cluster paradigm: internally dense, externally sparse
- “unlikely” to contract minimum cut edges
- clustering not main goal: only perform a couple of iterations

# Fast Inexact Minimum Cuts

- (Inexact) Cluster reduction + Exact reductions
- Solve kernel to optimality  
using Nagamochi, Ono and Ibaraki's algorithm  
→ overall linear running time, but potentially suboptimal cuts



# Parallelization

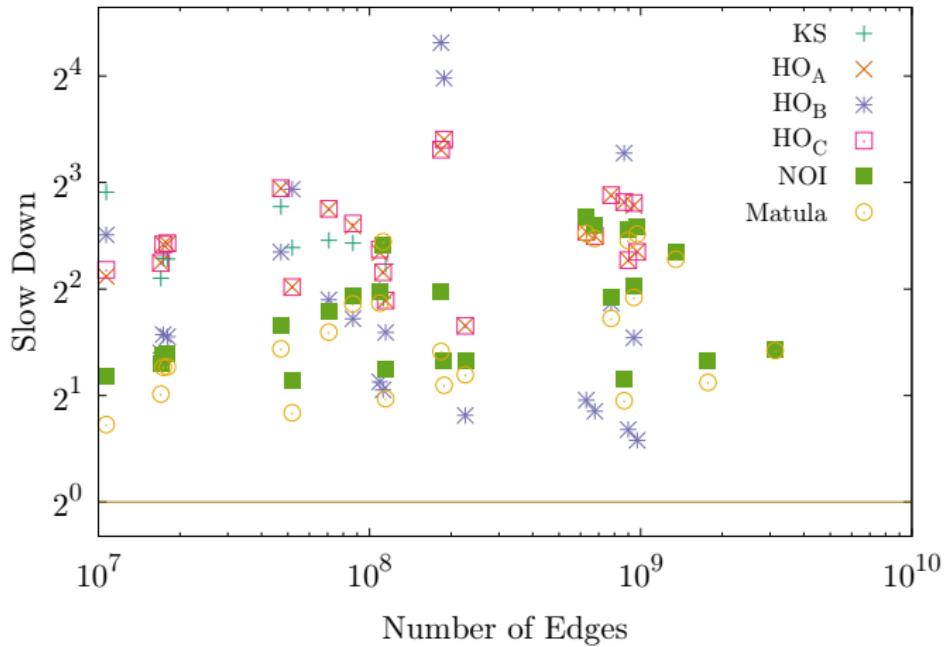
## Shared-memory with OpenMP

- Parallel label propagation



- as brutal as: pragma openmp for and ignore conflicts on labels
- Parallel Padberg-Rinaldi:
  - check edges independently  $\leadsto$  embarrassingly parallel
  - collect edges then contract  
 $\rightarrow$  essentially linear time
- Parallel contraction (not here)
- run Nagamochi, Ono and Ibarakis algorithm sequentially

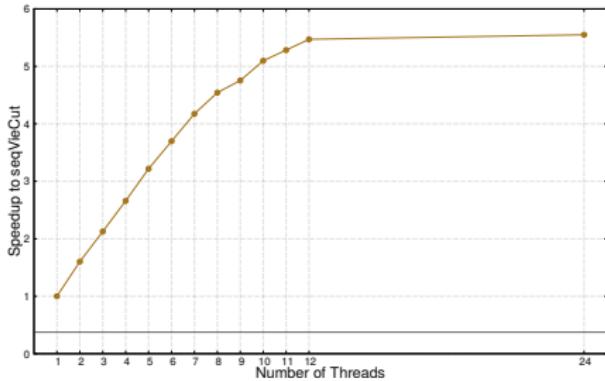
# Real-world Networks



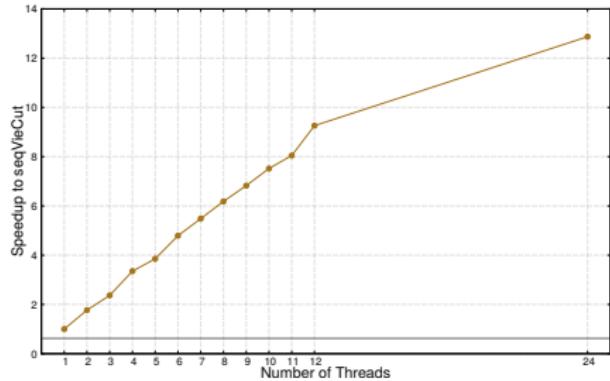
- No incorrect results (except Karger-Stein in 36% of the cases)

# Parallelization

uk-2007-05 k=10



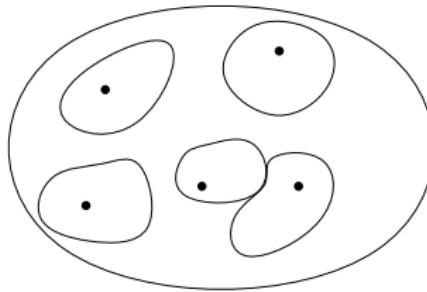
n=200K, d=10%, k=2



- Average speedup using 12 cores: 6.3 (24: 7.9)
- Average speedup to next fastest (Matula): 13.2 (24: 15.8)

# Stating the Obvious

- now  $\approx 16$  times faster than Matula
- NO guarantee for minimum cut, but experiments say very likely
- reductions depend on bound  $\hat{\lambda}$ 
  - ~~ PLUG IN our result into exact NOI algorithm + parallelization
    - ~~ currently fastest exact minimum cut algorithm  
(by an order of magnitude)



- also: extension to ALL minimum cuts, and dynamic minimum cuts

Multilevel Algorithms

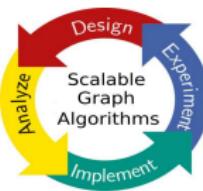
Evolutionary Computation

Parallel Programming

Kernelization

Dynamic Algorithms

shared-, distributed-, external-, internal memory



Algorithms for

graph partitioning  
graph clustering  
graph generation

process mapping  
minimum cuts  
independent sets

longest paths  
graph drawing  
(dyn.) matching

node separators  
dyn. reachability  
....

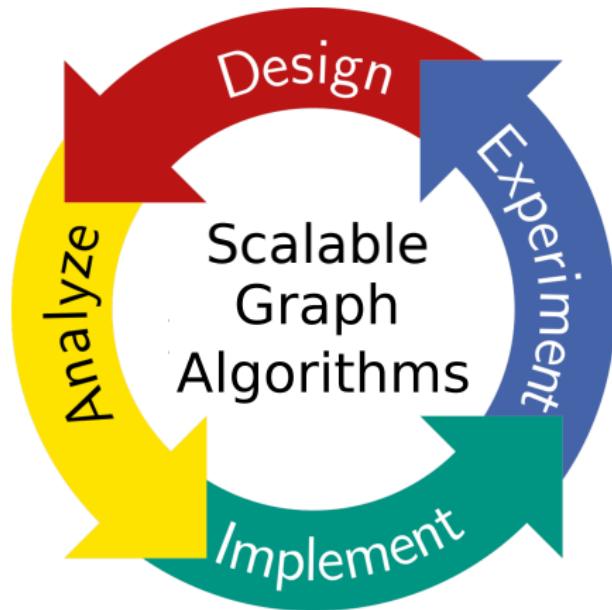
Open Source

Applications

territory design  
large scale simulations  
quantum annealing

route planning  
distributed system design  
nuclei segmentation

multiprocessor scheduling  
high-throughput DNA sequencing  
....



backup slides for other problems

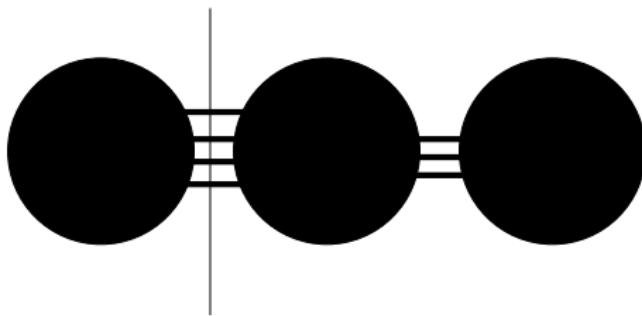
# Kernelization

Nagamochi, Ono, and Ibaraki

**Key Idea:** a spanning tree contains at least one edge from any cut

- Let  $\hat{\lambda}$  be your current bound for minimum cut
- Want: smaller minimum cut
- Compute  $\hat{\lambda} - 1$  maximal spanning forests (iteratively)
  - ~~ edges not in forests connect vertices with connectivity  $\geq \hat{\lambda}$
  - ~~ contract all of them

Example:  $\hat{\lambda} = 4$  ~~ compute 3 edge-disjoint **spanning forests**



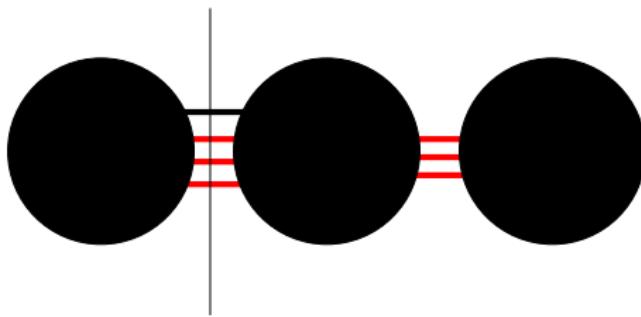
# Kernelization

Nagamochi, Ono, and Ibaraki

**Key Idea:** a spanning tree contains at least one edge from any cut

- Let  $\hat{\lambda}$  be your current bound for minimum cut
- Want: smaller minimum cut
- Compute  $\hat{\lambda} - 1$  maximal spanning forests (iteratively)
  - ~~ edges not in forests connect vertices with connectivity  $\geq \hat{\lambda}$
  - ~~ contract all of them

Example:  $\hat{\lambda} = 4$  ~~ compute 3 edge-disjoint **spanning forests**



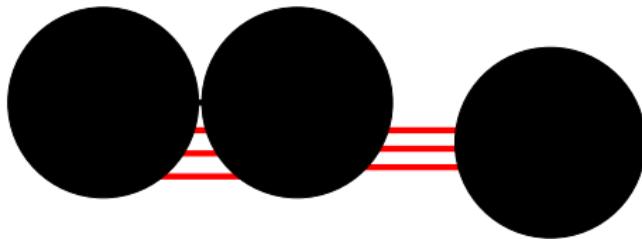
# Kernelization

Nagamochi, Ono, and Ibaraki

**Key Idea:** a spanning tree contains at least one edge from any cut

- Let  $\hat{\lambda}$  be your current bound for minimum cut
- Want: smaller minimum cut
- Compute  $\hat{\lambda} - 1$  maximal spanning forests (iteratively)
  - ~~ edges not in forests connect vertices with connectivity  $\geq \hat{\lambda}$
  - ~~ contract all of them

Example:  $\hat{\lambda} = 4$  ~~ compute 3 edge-disjoint **spanning forests**



# Nagamochi, Ono, Ibaraki

## Details

- $\lambda(x, y)$  capacity of minimum cut separating  $x$  and  $y$
- $\lambda(x, y) \geq \hat{\lambda} \rightsquigarrow \exists$  no cut separating  $x$  and  $y$  with capacity  $\leq \hat{\lambda}$   
 $\rightsquigarrow$  we can contract  $(x, y)$
- but computing  $\lambda(x, y)$  expensive (max-flow algorithm)
- NOI: compute **lower bound**  $q(e)$  on  $\lambda(x, y)$ , i.e.

$$\begin{aligned}\lambda(x, y) &\geq q(e) \geq \hat{\lambda} \\ \rightsquigarrow \text{can contract edge } e\end{aligned}$$

$q(e) = \#$  edge disjoint paths that connect  $x, y$

$q(e)$  via  $k$ -edge-connected subgraph  $\rightsquigarrow$  following algorithm

# $k$ -edge-connected subgraph

**invariant**  $r[v] = i$  smallest  $i$  s.t.  $E_{i+1} \cup \{e\}$  does not contain a cycle

initialize  $r[v] = 0$

all nodes and edges are **non-scanned**

$E_1 = E_2 = \dots = E_{|E|} = \emptyset$

**while**  $\exists$  **non-scanned** node

$u :=$  **non-scanned** node  $v$  with maximal  $r[v]$

**for each** non-scanned edge  $e = (u, v) \in E$  **do**

insert  $e$  into  $E_{r(v)+1}$

$q(e) = r(v) + 1, r(v) = r(v) + 1$

$\rightsquigarrow H_i = (V, E_i)$  is a maximal spanning forest in  $G \setminus E_1 \cup \dots \cup E_{i-1}$

Long story short:

Everything in  $E_{\hat{\lambda}} \cup \dots \cup E_{|E|}$  can be contracted.

$\rightsquigarrow$  contract  $e$  if  $q(e) \geq \hat{\lambda}$

# $k$ -edge-connected subgraph

**invariant**  $r[v] = i$  smallest  $i$  s.t.  $E_{i+1} \cup \{e\}$  does not contain a cycle

**invariant**  $r[v] = i$  incident to first  $i$  trees

initialize  $r[v] = 0$

all nodes and edges are **non-scanned**

$E_1 = E_2 = \dots = E_{|E|} = \emptyset$

**while**  $\exists$  **non-scanned** node

$u :=$  **non-scanned** node  $v$  with maximal  $r[v]$

**for each** non-scanned edge  $e = (u, v) \in E$  **do**

insert  $e$  into  $E_{r(v)+1}$

$q(e) = r(v) + 1, r(v) = r(v) + 1$

$\rightsquigarrow H_i = (V, E_i)$  is a maximal spanning forest in  $G \setminus E_1 \cup \dots \cup E_{i-1}$

Long story short:

Everything in  $E_{\hat{\lambda}} \cup \dots \cup E_{|E|}$  can be contracted.

$\rightsquigarrow$  contract  $e$  if  $q(e) \geq \hat{\lambda}$

# $k$ -edge-connected subgraph

**invariant**  $r[v] = i$  smallest  $i$  s.t.  $E_{i+1} \cup \{e\}$  does not contain a cycle

initialize  $r[v] = 0$

all nodes and edges are **non-scanned**

$E_1 = E_2 = \dots = E_{|E|} = \emptyset$

**while**  $\exists$  **non-scanned** node

$u :=$  **non-scanned** node  $v$  with maximal  $r[v]$

**for each** non-scanned edge  $e = (u, v) \in E$  **do**

insert  $e$  into  $E_{r(v)+1}, \dots, E_{r(v)+c(e)}$

$q(e) = r(v) + c(e)$ ,  $r(v) = r(v) + c(e)$

$\rightsquigarrow H_i = (V, E_i)$  is a maximal spanning forest in  $G \setminus E_1 \cup \dots \cup E_{i-1}$

Long story short:

Everything in  $E_{\hat{\lambda}} \cup \dots \cup E_{|E|}$  can be contracted.

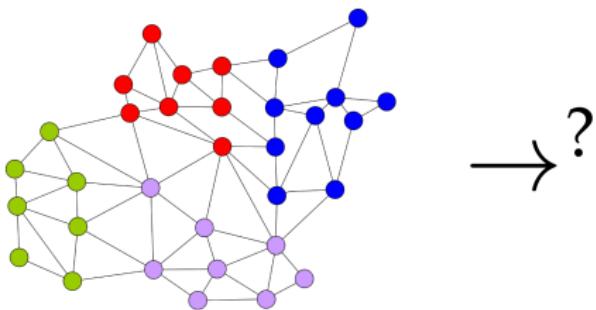
$\rightsquigarrow$  contract  $e$  if  $q(e) \geq \hat{\lambda}$

$c(e)$  replaces one edge by  $c(e)$  edges

# Partitioning

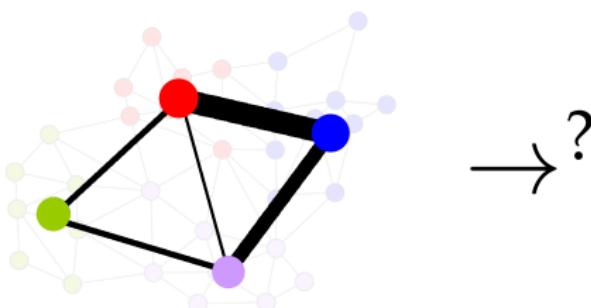
see lecture slides

# Process Mapping and ...



Single blocks represents *process* to be executed by a PE  
 $c_{i,j}$  amount of communication between *process i*, *process j*

# Process Mapping and ...



Single blocks represents *process* to be executed by a PE  
 $c_{i,j}$  amount of communication between *process i*, *process j*

# ... and Quadratic Assignment

[Brandfass et al. and others]

$\mathcal{C} \in \mathbf{R}^{n \times n}$  communication matrix

$\mathcal{D} \in \mathbf{R}^{n \times n}$  distance matrix

**Quadratic assignment problem (QAP):**

Find one-to-one mapping  $\Pi$  (processes to PEs) which minimizes

$$J(\mathcal{C}, \mathcal{D}, \Pi) := \sum_{i,j} \mathcal{C}_{\Pi(i), \Pi(j)} \mathcal{D}_{i,j}$$

$\mathcal{C}_{\Pi(i), \Pi(j)} \mathcal{D}_{i,j}$ : cost for communication between PE  $i$  and  $j$

$\rightarrow J(\mathcal{C}, \mathcal{D}, \Pi)$  overall cost of communication

# Basic Assumptions

... Sparse Quadratic Assignment



- communication matrix  $\mathcal{C}$  is **sparse** → graph  $G_{\mathcal{C}}$
- compute system is **hierarchically structured**

$\mathcal{S} = a_1, \dots, a_k$  system hierarchy

e.g.  $a_1$  cores per processor,  $a_2$  processors per node,  $a_3$  nodes per rack, ...

$D = d_1, \dots, d_k$  distances in hierarchy

$d_i$  distance of cores in same subsystems for  $i' < i$ , different subsystems for  $i' \geq i$

# Initial Solutions

## Intuition:

- place well connected subgraphs in  $G_C$  closely

## Algorithms:

- Identity
- Müller-Merbach (Greedy)
  - iteratively build mapping
  - assign process with largest comm. vol to core with smallest total distance
- Dual-Recursive-Bisection
  - bipartition  $G_C$  and  $\mathcal{D}$  recursively
- Bottom up approach
- Top down approach

# Initial Solutions

## Intuition:

- place well connected subgraphs in  $G_C$  closely

## Top down approach:

*Idea:* partition along the system hierarchy

$\mathcal{S} = a_1, a_2, \dots, a_k$  hierarchy (with  $n = \prod_i a_i$ )

compute *perfectly balanced partition* of  $G_C$  into  $a_k$  blocks

each block has  $n/a_k$  vertices (processes)

each system entity provides  $n/a_k$  PEs

→**assign blocks** to system entities

**recurse** on each block with  $a_{k-1}, \dots$  until  $a_1$  vertices left

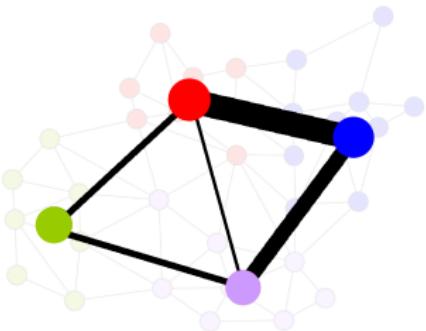
# Initial Solutions

Top Down Illustrated

System configuration  $\mathcal{S}$

racks	4
nodes per rack	4
PEs per node	4
cores per PE	4

Model  $G_{\mathcal{C}}$ :



# Initial Solutions

Top Down Illustrated

System configuration  $\mathcal{S}$

racks 4

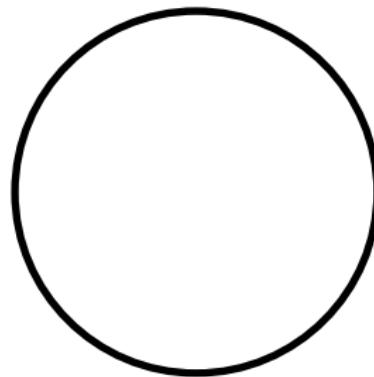
nodes per rack 4

PEs per node 4

cores per PE 4



Model  $G_{\mathcal{C}}$ :



# Initial Solutions

Top Down Illustrated

System configuration  $\mathcal{S}$

racks 4

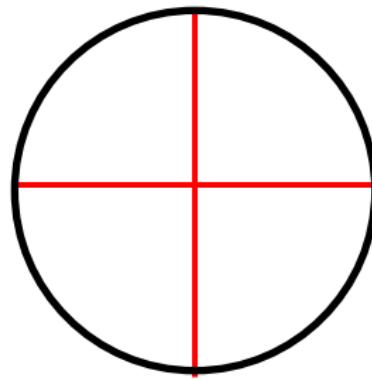
nodes per rack 4

PEs per node 4

cores per PE 4



Model  $G_{\mathcal{C}}$ :



# Initial Solutions

Top Down Illustrated

System configuration  $\mathcal{S}$

racks 4

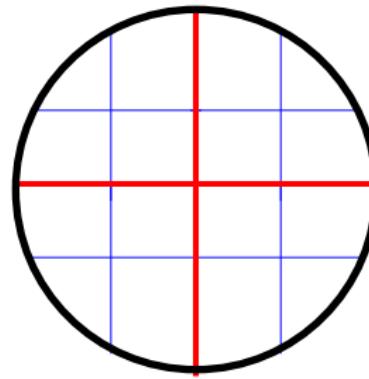
nodes per rack 4

PEs per node 4

cores per PE 4



Model  $G_{\mathcal{C}}$ :



# Initial Solutions

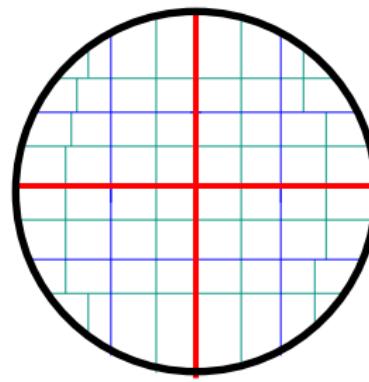
Top Down Illustrated

System configuration  $\mathcal{S}$

racks	4
nodes per rack	4
PEs per node	4
cores per PE	4



Model  $G_{\mathcal{C}}$ :



# Local Search

[Heider'72]

**Goal:** perform swaps to improve mapping

**Swaps** in pair-exchange neighborhood  $N(\Pi)$ :

all permutations reachable by swapping two elements in  $\Pi$

- swap  $i, j$ :  $\Pi(i)$  assigned to PE  $j$ ,  $\Pi(j)$  is assigned to PE  $i$
- check  $N(\Pi)$  in cyclic manner
- keep swap if objective reduced (else undo)

**Overall runtime**  $O(n^3)$

- quadratic amount of pairs
- each swap takes  $O(n)$  time (update objective)

# Local Search

Modifications by [Brandfass et al.'13]

## $\mathcal{N}_p$ (*pruned neighborhood*)

- remove pairs  $(i, j)$  if objective cannot change
- partition search space into  $s$  consecutive index blocks  
→ only perform swaps inside those blocks
- number of pairs  $O(ns)$  → overall runtime  $O(n^2s)$

$$\Pi : \quad \boxed{\phantom{s} \quad \boxed{s} \quad \boxed{s} \quad \boxed{s} \quad \boxed{s}}$$

## Considered initial solutions:

- Identity
- Müller-Merbach  
→ improves runtime of local search and objective

# Local Search

## Runtime Problem

Computing and updating objective is expensive!

$$J(\mathcal{C}, \mathcal{D}, \Pi) := \sum_{i,k} \mathcal{C}_{\Pi(i), \Pi(k)} \mathcal{D}_{i,k}$$

- matrices have quadratic number of elements  
→ initial computation of objective  $O(n^2)$
- update objective using old value:
  - swap  $(i, j)$  needs access to all elements in two columns of **communication and distance matrix**
    - update step takes  $O(n)$  time
    - bottleneck for sparse  $G_C$



# Faster Local Search

**Observation:** initial objective can be computed in  $O(n + m)$

rewrite objective to work with inverse of permutation:

$$\begin{aligned} J(\mathcal{C}, \mathcal{D}, \Pi) &= \sum_{i,k} \mathcal{C}_{\Pi(i), \Pi(k)} \mathcal{D}_{i,k} \\ &= \sum_{u,v} \mathcal{C}_{u,v} \mathcal{D}_{\Pi^{-1}(u), \Pi^{-1}(v)} \end{aligned}$$

with the interpretation that task  $u$  is assigned to PE  $\Pi^{-1}(u)$

rewrite objective to work with graph representation:

$$J(\mathcal{C}, \mathcal{D}, \Pi) := \sum_{(u,v) \in E[\mathcal{C}]} \mathcal{C}_{u,v} \mathcal{D}_{\Pi^{-1}(u), \Pi^{-1}(v)}$$

# Faster Local Search

**Next goal:** make update of objective fast

Vertex contribution to objective:

$$\Gamma_{\Pi^{-1}}(u) := \sum_{v \in N(u)} \mathcal{C}_{u,v} \mathcal{D}_{\Pi^{-1}(u), \Pi^{-1}(v)}$$

Rewrite objective using vertex contribution

$$J(\mathcal{C}, \mathcal{D}, \Pi) := \sum_{u \in V} \Gamma_{\Pi^{-1}}(u)$$

**Throughout the algorithm:**

- keep vertex contributions  $\Gamma$  up to date
- swap affects nodes  $u, v$  and their neighborhood
  - update node contributions of those nodes
  - update objective accordingly
- $O(d_u + d_v)$  time

# Alternative Local Search Spaces

**Simple version  $N_{\mathcal{C}}$ :**

swap only if there is an edge in  $G_{\mathcal{C}}$

- $|N_{\mathcal{C}}| = m$  (contains exactly  $m$  pairs)
- swaps are performed in random order
- terminate after  $m$  unsuccessful swaps
- assumes that swaps with positive gain are close in  $G_{\mathcal{C}}$

**Augmented version  $N_{\mathcal{C}}^d$ :**

swaps only if processes have distance less than  $d$  in  $G_{\mathcal{C}}$

# Benchmark Instances

Two different configurations:

$\mathcal{S} = a_1, \dots, a_k$  system hierarchy

$D = d_1, \dots, d_k$  distances in hierarchy

$d_i$  distance of two cores in same subsystems for  $i' < i$ ,  
and different subsystems for  $i' \geq i$

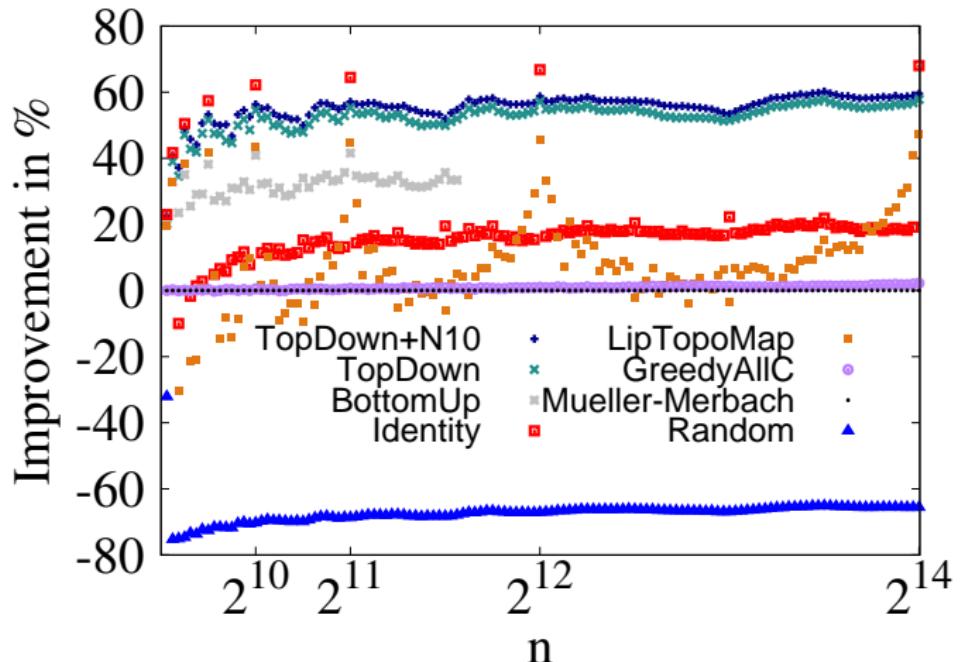
$n = \prod_i a_i$  total number of cores

Process to create instances:

- partition graph into  $n$  blocks
- compute the communication graph
- use edge cut as communication volume
- compute mapping to specified system

Graph	$n$	$m$
UF Graphs		
cop20k_A	99843	1262244
2cubes_sphere	101492	772886
thermomech_TC	102158	304700
cfd2	123440	1482229
boneS01	127224	3293964
DubcovA3	146689	1744980
bmwcra_1	148770	5247616
G2_circuit	150102	288286
shipsec5	179860	4966618
cont-300	180895	448799
Large Walshaw Graphs		
598a	110971	741934
fe_ocean	143437	409593
144	144649	1074393
wave	156317	1059331
m14b	214765	1679018
auto	448695	3314611
Large Other Graphs		
del23	≈8.4M	≈25.2M
del24	≈16.7M	≈50.3M
rgg23	≈8.4M	≈63.5M
rgg24	≈16.7M	≈132.6M
deu	≈4.4M	≈5.5M
eur	≈18.0M	≈22.2M
af_shell9	≈504K	≈8.5M
thermal2	≈1.2M	≈3.7M
nlr	≈4.2M	≈12.5M

# Initial Heuristics



baseline: Müller-Merbach

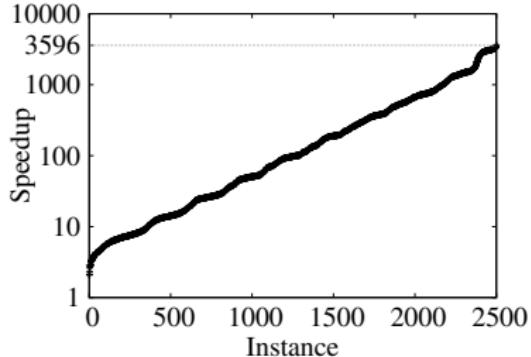
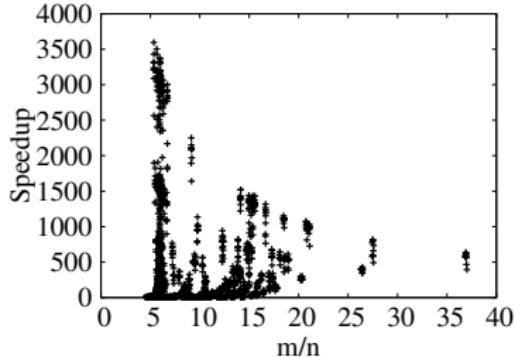
# Speedup of Local Search

Running Time and Speedup of Local Search for  $N_p$

$n$	$m/n$	$t_{\text{LS}}[\text{s}]$	$t_{\text{fastLS}}[\text{s}]$	speedup
64	6.7	0.016	0.003	5.3
128	7.3	0.064	0.006	10.7
256	7.9	0.268	0.014	19.1
512	8.3	1.073	0.029	37.0
1K	8.8	4.263	0.059	72.3
2K	9.2	17.083	0.124	137.8
4K	9.7	68.360	0.260	262.9
8K	10.3	268.907	0.540	498.0
16K	11.2	1 075.107	1.158	928.4
32K	12.5	4 348.374	2.472	1 759.1

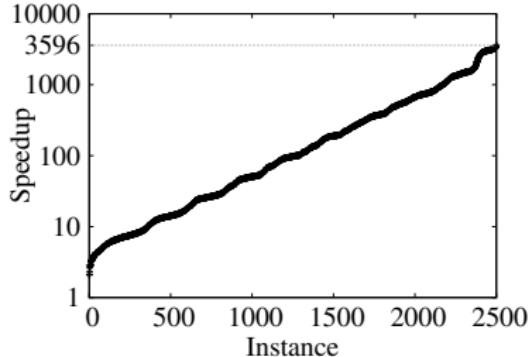
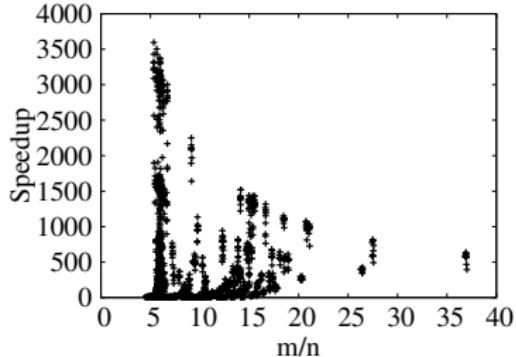
- $t_{\text{LS}}$  running time using slow gain computations
- $t_{\text{fastLS}}$  running time using fast gain computations

# Faster Local Search



- algorithmic speedup as a function of graph density
- algorithmic speedup for the different instances

# Faster Local Search



- algorithmic speedup as a function of graph density
- algorithmic speedup for the different instances

Vienna Mapping and Sparse Quadratic Assignment:  
[viem.taa.univie.ac.at](http://viem.taa.univie.ac.at)

# Concrete Example

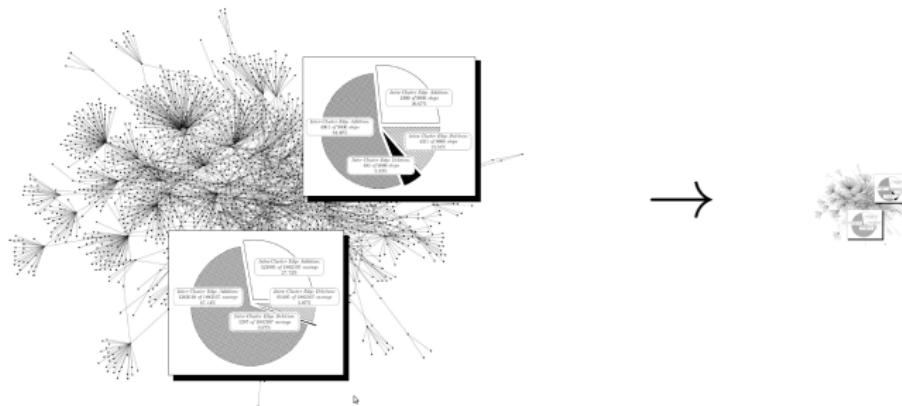
## Applied Kernelization for Independent Sets

# Kernelization

## General Idea

### Reductions:

rules to decrease graph size, while maintaining optimality



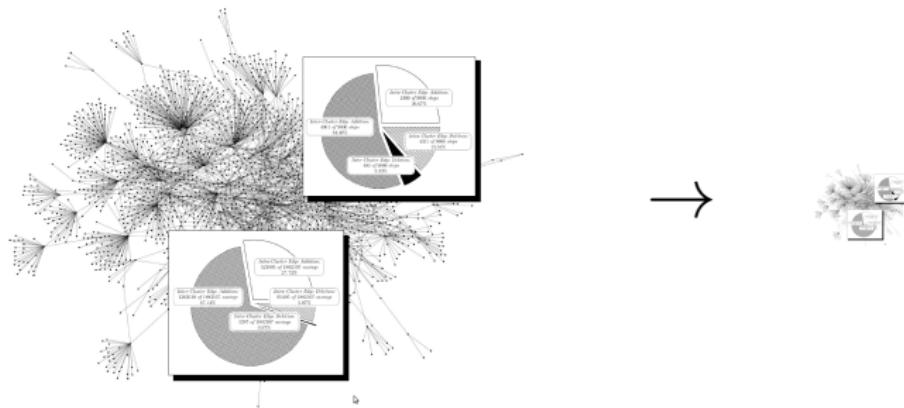
solve problem on **problem kernel**  
→ obtain solution on input graph

# Kernelization

## General Idea

### Reductions:

rules to decrease graph size, while maintaining optimality



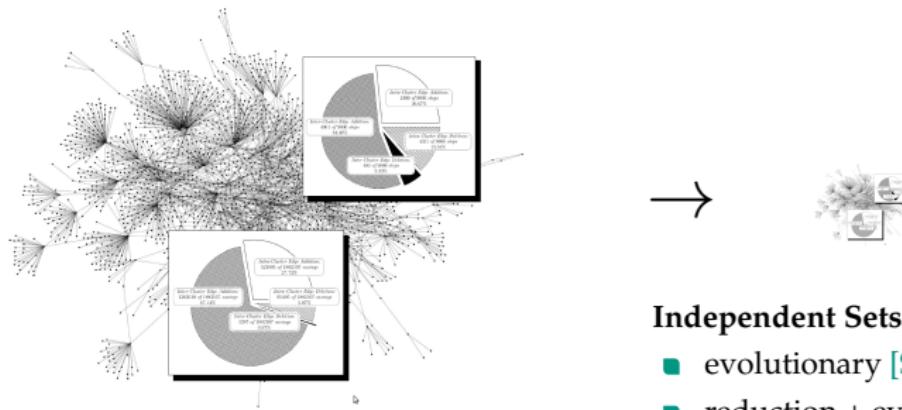
solve problem on **problem kernel (using a heuristic)**  
→ obtain solution on input graph **quickly**

# Kernelization

## General Idea

### Reductions:

rules to decrease graph size, while maintaining optimality



solve problem on **problem kernel**  
→ obtain solution on input graph

### Independent Sets

- evolutionary [SEA'15]
- reduction + evolutionary [ALX'16]
- online reductions + LS [SEA'16]
- shared-mem parallel [ALX'18]
- weighted exact [ALX'19]

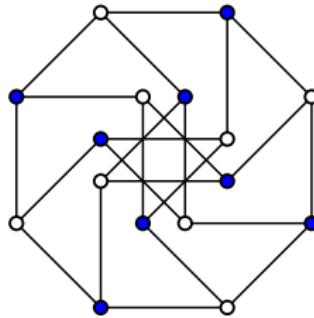
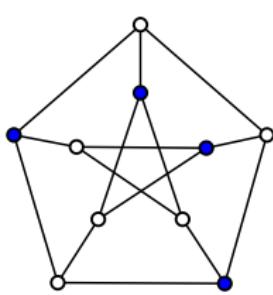
# Definitions

## Independent Set:

- subset  $S \subseteq V$  such that there are no adjacent nodes in  $S$

## Maximum Independent Set (MIS):

- maximum cardinality set  $S$



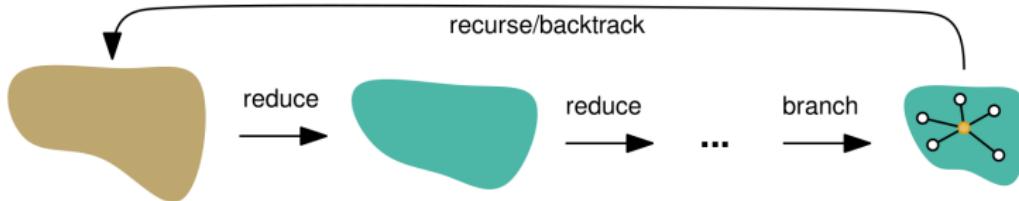
- related to **maximum clique** and **minimum vertex cover**
- finding a MIS is **NP-hard** and hard to approximate

# Exact Algorithms

## Independent Sets

Exact solution using **branch and bound**:

- Modify and remove subgraphs during recursions → **reductions**
- Branch when the graph can no longer be reduced



→ Running time  $O(1.211^n)$  [Akiba, Iwata'16]

Works well, except when it doesn't:

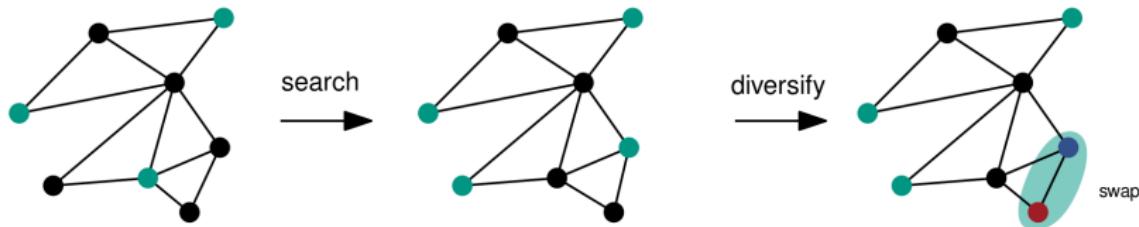
- Many networks with **small kernels** are easy to solve
- Similar instances with **large kernels** remain unsolved

# Heuristic Algorithms

Huge real-world networks infeasible for exact algorithms  
→ finding **high-quality** approximations in short time

## Recent Approaches:

- Local search based on swaps [Andrade et al. 2012]

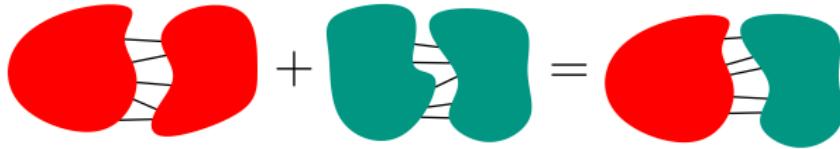


# Heuristic Algorithms

Huge real-world networks infeasible for exact algorithms  
→ finding **high-quality** approximations in short time

## Recent Approaches:

- Local search based on swaps [Andrade et al. 2012]
- Evolutionary algorithms (EvoMIS) [SEA'15]

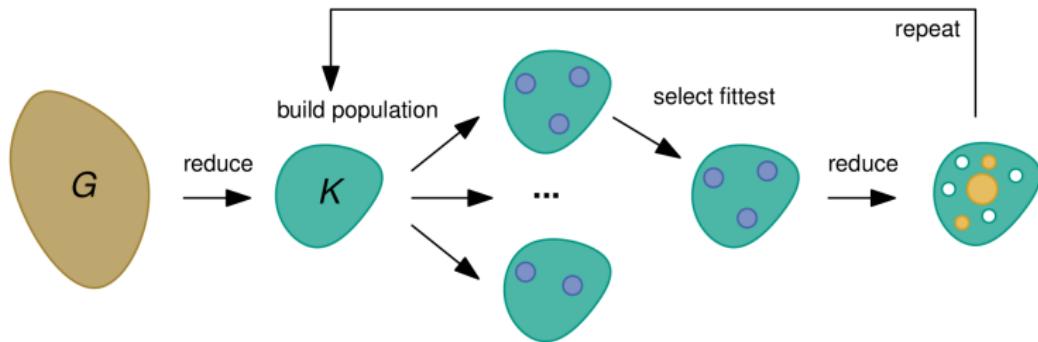


# Heuristic Algorithms

Huge real-world networks infeasible for exact algorithms  
→ finding **high-quality** approximations in short time

## Recent Approaches:

- Local search based on swaps [Andrade et al. 2012]
- Evolutionary algorithms (EvoMIS) [SEA'15]
- Combine with reductions (ReduMIS) [ALENEX'16]

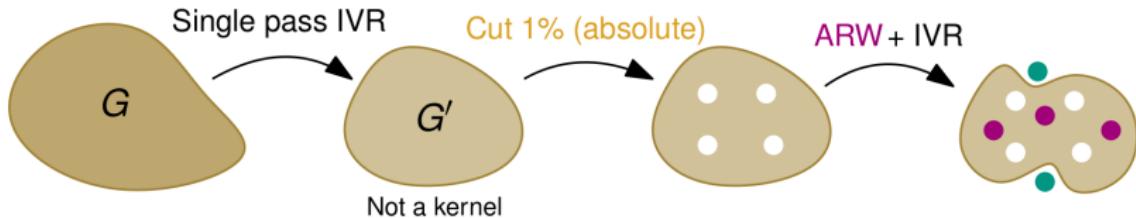


# Heuristic Algorithms

Huge real-world networks infeasible for exact algorithms  
→ finding **high-quality** approximations in short time

## Recent Approaches:

- Local search based on swaps [Andrade et al. 2012]
- Evolutionary algorithms (EvoMIS) [SEA'15]
- Combine with reductions (ReduMIS) [ALENEX'16]
- Online reductions (OnlineMIS) [SEA'16]

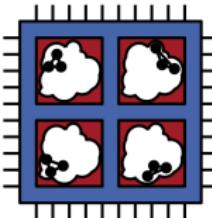


# Heuristic Algorithms

Huge real-world networks infeasible for exact algorithms  
→ finding **high-quality** approximations in short time

## Recent Approaches:

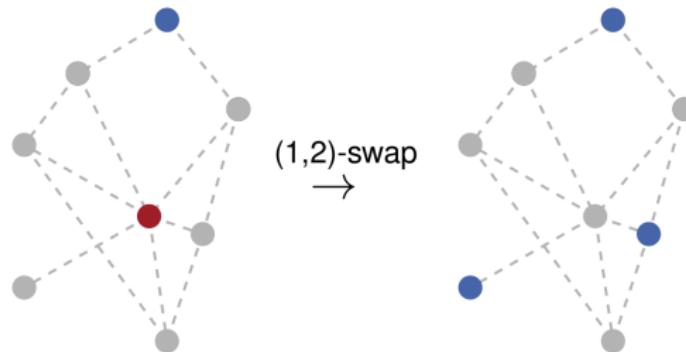
- Local search based on swaps [Andrade et al. 2012]
- Evolutionary algorithms (EvoMIS) [SEA'15]
- Combine with reductions (ReduMIS) [ALENEX'16]
- Online reductions (OnlineMIS) [SEA'16]
- Shared-Memory parallel (ParMIS) [ALENEX'17]



# Iterated Local Search

[Andrade et al.'12]

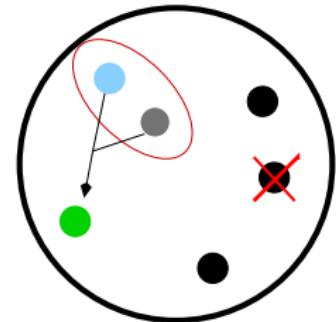
- Compute initial maximal independent set
- Improve using **(1, 2)-swaps**:
  - remove **single** solution node and insert **two** new ones
  - search for (1, 2)-swaps in time  $O(m)$
- **Tabu mechanism** for “recently” swapped vertices
- No (1,2)-swap → **perturbation step**



# Evolutionary Algorithm

[SEA'15]

```
create initial population  $\mathcal{P}$ 
while stopping criterion not fulfilled
  select parents  $\mathcal{P}_1, \mathcal{P}_2$  from  $\mathcal{P}$ 
  recombine  $\mathcal{P}_1$  with  $\mathcal{P}_2$  to create offspring  $o$ 
  mutate offspring  $o$ 
  evict individual in population using  $o$ 
return fittest individual
```



## Recombine Using Graph Partitioning:

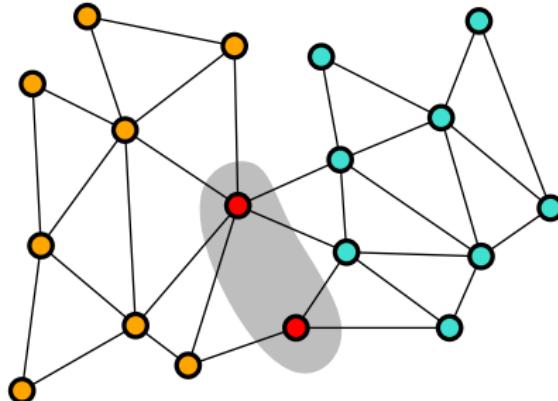
- exchange **whole blocks** of solutions
- small **objective** vital for efficiency



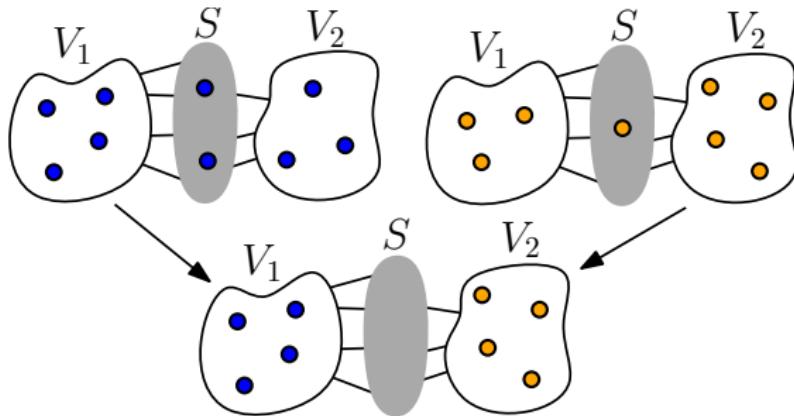
# Separators

Partition graph  $G = (V, E, c : V \rightarrow \mathbf{R}_{>0}, \omega : E \rightarrow \mathbf{R}_{>0})$   
into  $k$  disjoint blocks + **node separator** s.t.

- total node weight of each block  $\leq \frac{1+\epsilon}{k}$  total node weight
- total size of **node separator** as small as possible

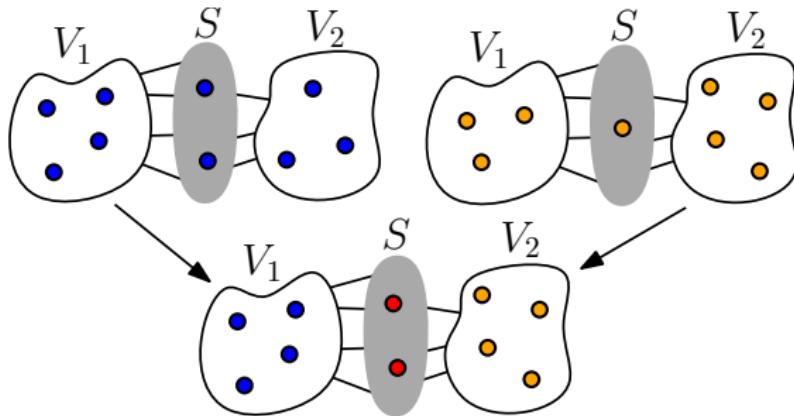


# Separator Combine



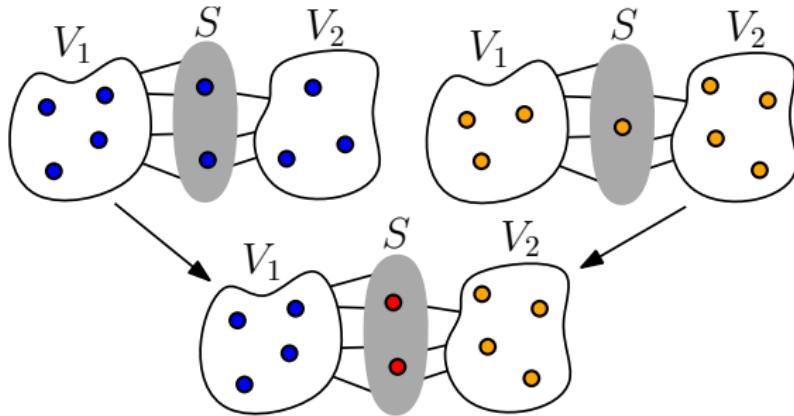
- build **node separator**  $V = V_1 \cup V_2 \cup S$
- use node separator as **crossover point**
- combination takes **linear time**  $O(n)$
- maximize with local search

# Separator Combine



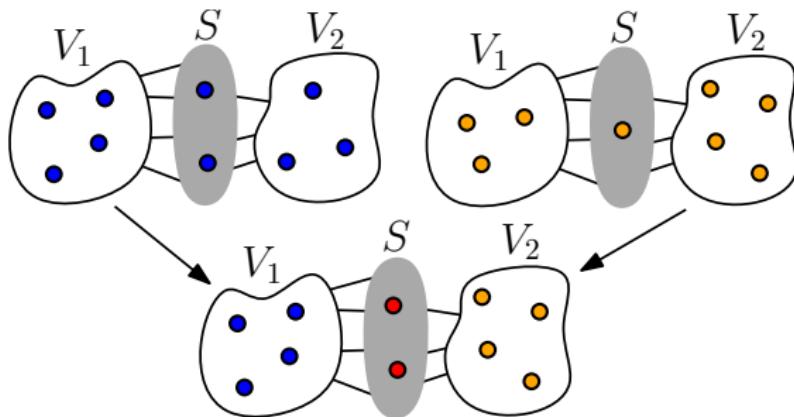
- build **node separator**  $V = V_1 \cup V_2 \cup S$
- use node separator as **crossover point**
- combination takes **linear time**  $O(n)$
- **maximize** with local search

# Separator Combine



Good Solutions, but Slow

# Separator Combine



Good Solutions, but Slow



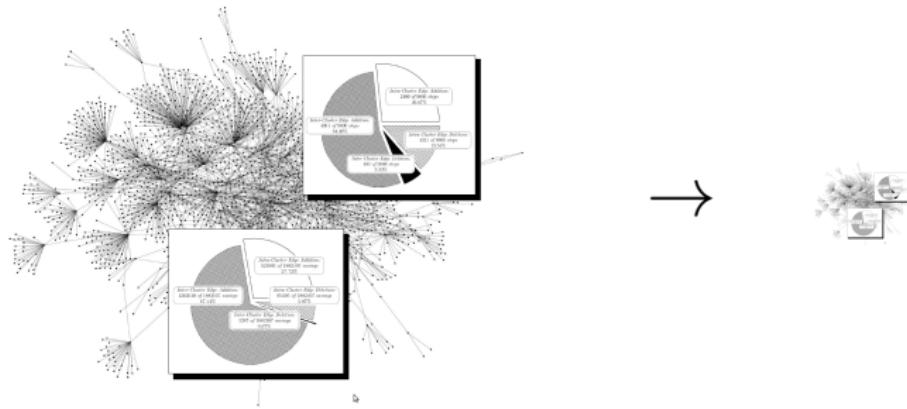
Apply EA on Kernel

# Kernelization

[Akiba, Iwata'15]

## Reductions:

rules to decrease graph size, while maintaining optimality



solve problem on **problem kernel** (using EA)  
→ obtain solution on input graph

# Kernelization

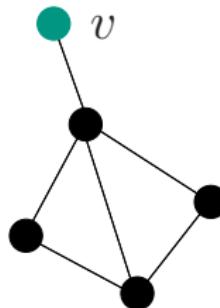
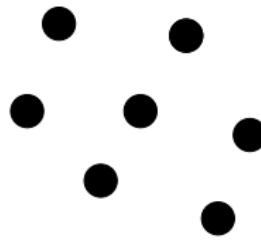
[Akiba, Iwata'15]

## Reductions:

rules to decrease graph size, while maintaining optimality

## Example:

remove degree 0 or 1 vertices recursively



there is always a MIS that contains  $v$

if neighbor of  $v$  in MIS choose  $v$  instead; otherwise add  $v$

# Kernelization

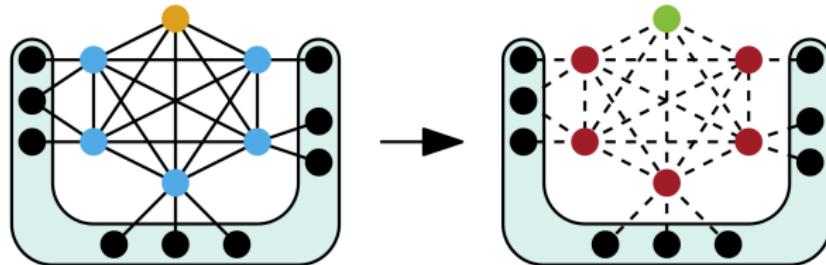
[Akiba, Iwata'15]

## Reductions:

rules to decrease graph size, while maintaining optimality

## Example:

isolated clique reduction



# Kernelization

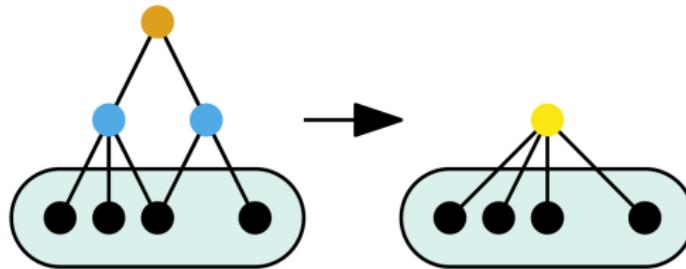
[Akiba, Iwata'15]

## Reductions:

rules to decrease graph size, while maintaining optimality

## Example:

vertex folding



# Kernelization

[Akiba, Iwata'15]

## Reductions:

rules to decrease graph size, while maintaining optimality

## Example:

linear programming relaxation for MIS

$$\begin{aligned} & \max \sum_{v \in V} x_v \text{ s.t.} \\ & \forall (u, v) \in E : x_u + x_v \leq 1 \\ & \forall v \in V : x_v \geq 0 \end{aligned}$$

- $\exists$  half integral solution (i.e., using only values 0, 1/2, and 1)
- solve via maximum bipartite matching (also parallel, stay tuned)
- vertices with value 1, MUST be in the MIS

# Kernelization

[Akiba, Iwata'15]

## Reductions:

rules to decrease graph size, while maintaining optimality

## Example:

linear programming relaxation for MIS

$$\begin{aligned} & \max \sum_{v \in V} x_v \text{ s.t.} \\ & \forall (u, v) \in E : x_u + x_v \leq 1 \\ & \forall v \in V : x_v \geq 0 \end{aligned}$$

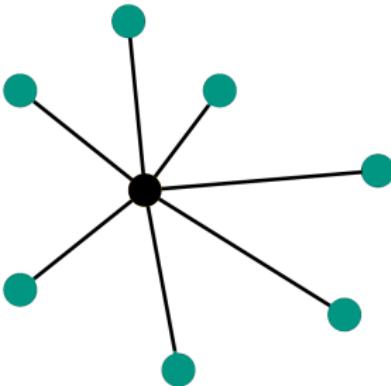
- $\exists$  half integral solution (i.e., using only values 0, 1/2, and 1)
- solve via maximum bipartite matching (also parallel, stay tuned)
- vertices with value 1, MUST be in the MIS

more reductions used in practice → [ALENEX'16]

# Guess “likely candidates”

can we **guess** vertices that are in MIS?

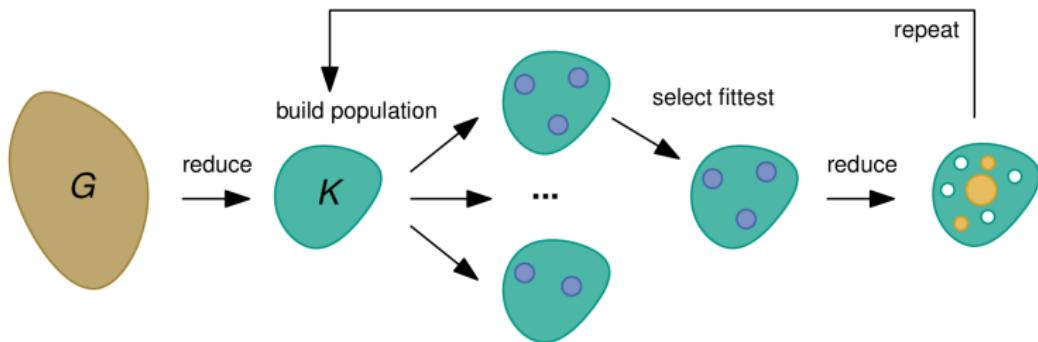
**idea:** select small-degree vertices from “fittest” independent set  
apply more reductions and recurse!



# Guess “likely candidates”

can we **guess** vertices that are in MIS?

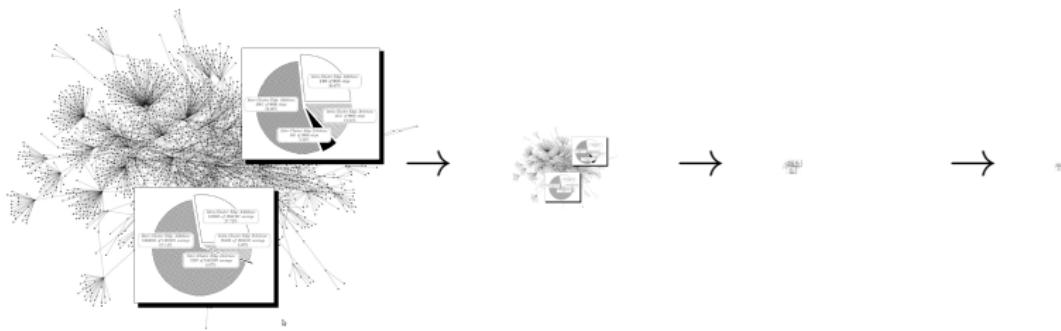
**idea:** select small-degree vertices from “fittest” independent set  
apply more reductions and recurse!



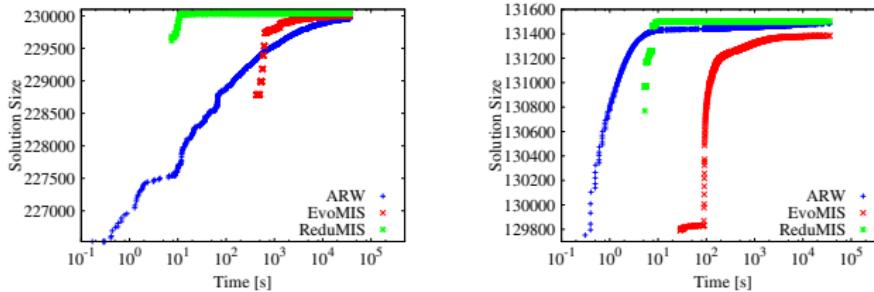
# Guess “likely candidates”

can we **guess** vertices that are in MIS?

**idea:** select small-degree vertices from “fittest” independent set  
apply more reductions and recurse!



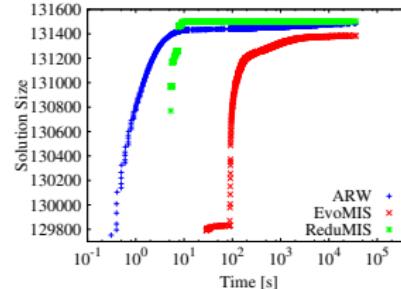
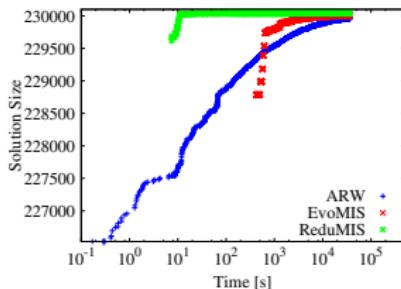
# Near-optimal on “difficult networks”



## Results:

- finds exact MIS faster, when exact algorithm is slow:
  - Skitter **48 min** → **21 min**
  - Stanford **13 hours** → **5 min**
  - bcsstk30 **8.6 hours** → **2.4 sec**
  - Skitter **2 hours** → **28 sec**, ...
- finds exact MIS, for large networks with known MIS size
- consistently finds larger solutions on social and road networks
- even as we scale to graphs to **10M** to **100M** nodes

# Near-optimal on “difficult networks”



## Problems:

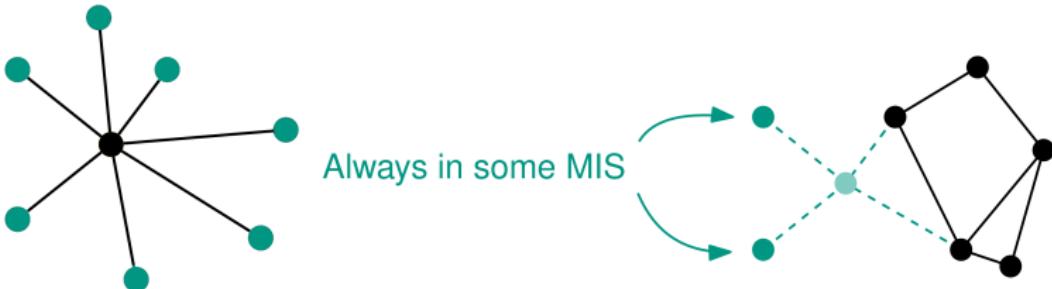
- Kernelization gives **high-quality** at cost of preprocessing time
- ARW is **fast** but struggles with complex scale-free networks

→ accelerate local search? parallelization?

# ARW on Complex Networks

## Issues

- High degree vertices take long to process ...  
... and are unlikely to be in MIS → High-degree cutting
- Many vertices are always in some solution → Kernelization



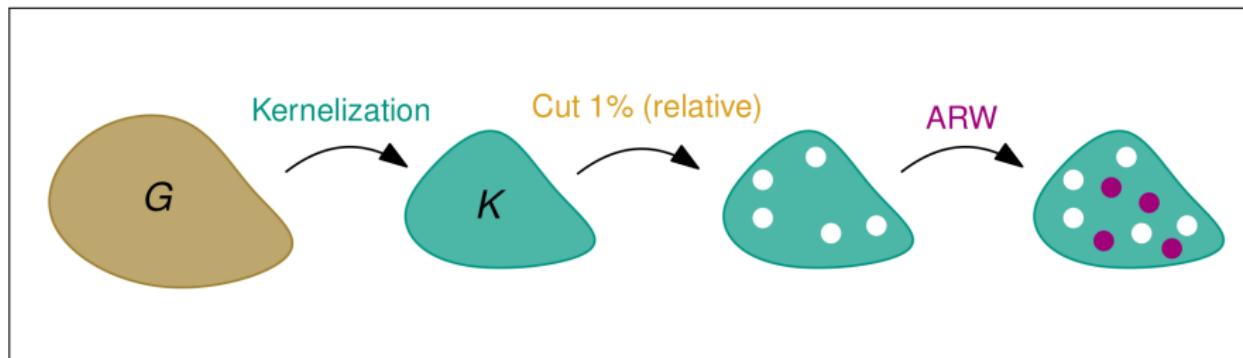
# Strategy I

## Full Kernelization

- Repeatedly apply reductions → kernel graph

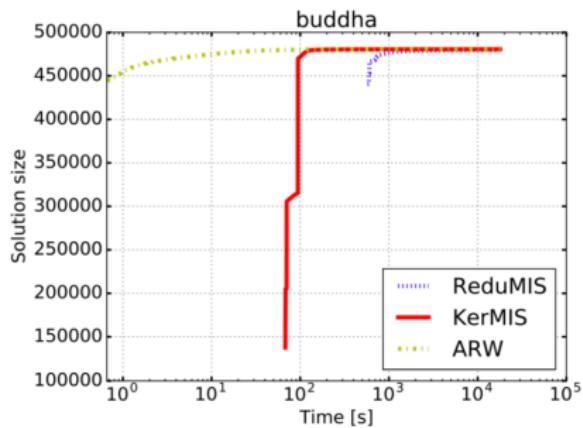
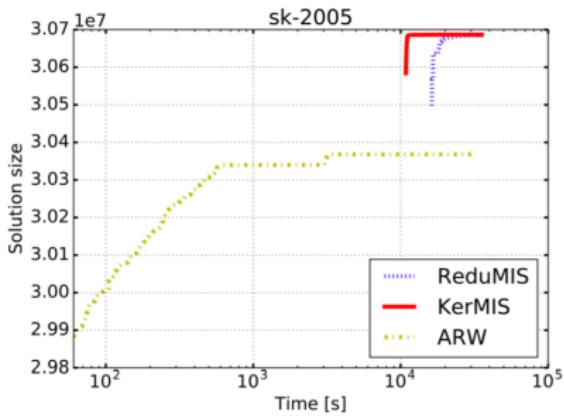


## KerMIS:



# KerMIS evaluation

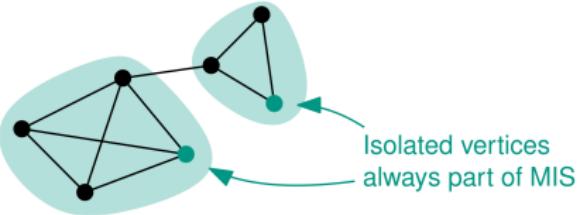
- Still long **preprocessing time** to compute kernel
- Inexact reductions **improve performance** of local search
- Quality comparable to **ReduMIS**
  - → Inexact reductions **don't worsen solution quality**



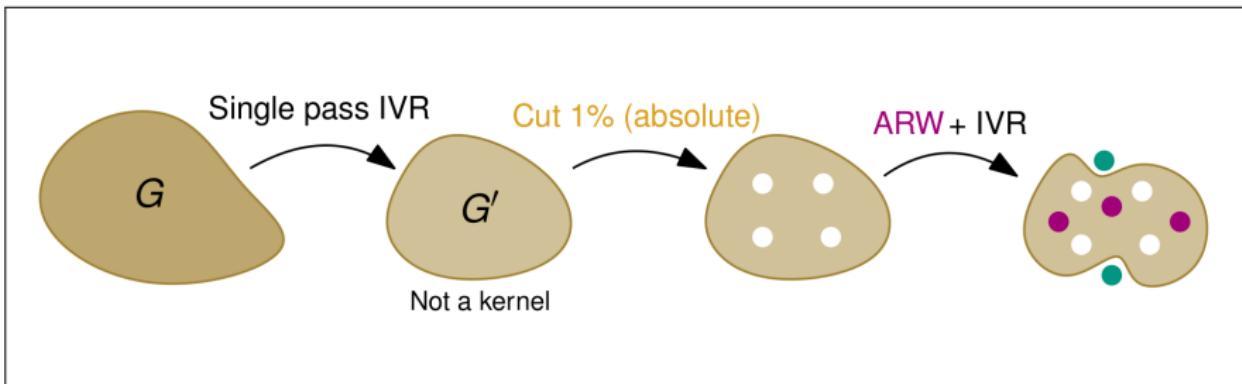
# Strategy II

## Online “Removal”

- Isolated vertex removal (IVR)

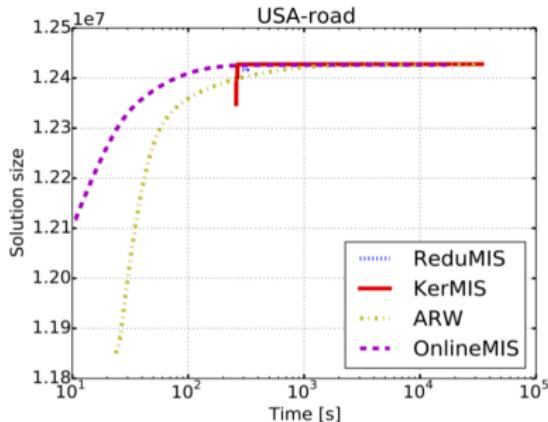
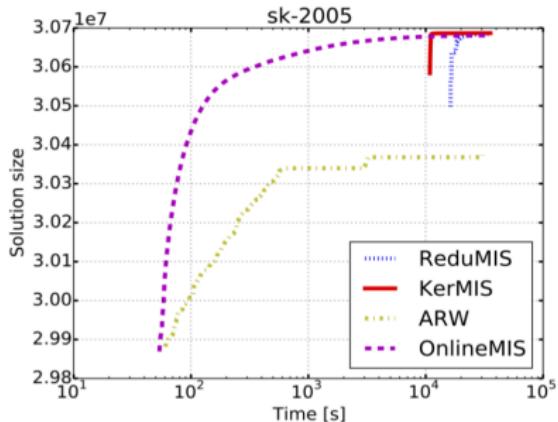


## OnlineMIS:



# OnlineMIS evaluation

- Significant improvements in **speed and quality**
- Huge scale-free networks can be processed efficiently

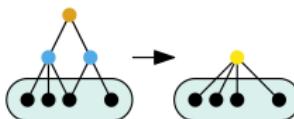


- Maximum speedup  $\geq 300$  over ReduMIS for 14/24 instances

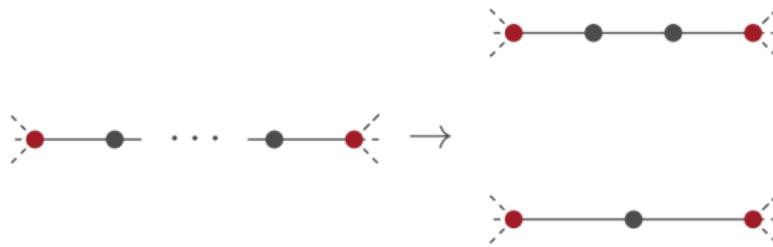
# Linear Time Reductions

[Chang et al.'17]

~~ vertex folding is slow with high-degree neighbors



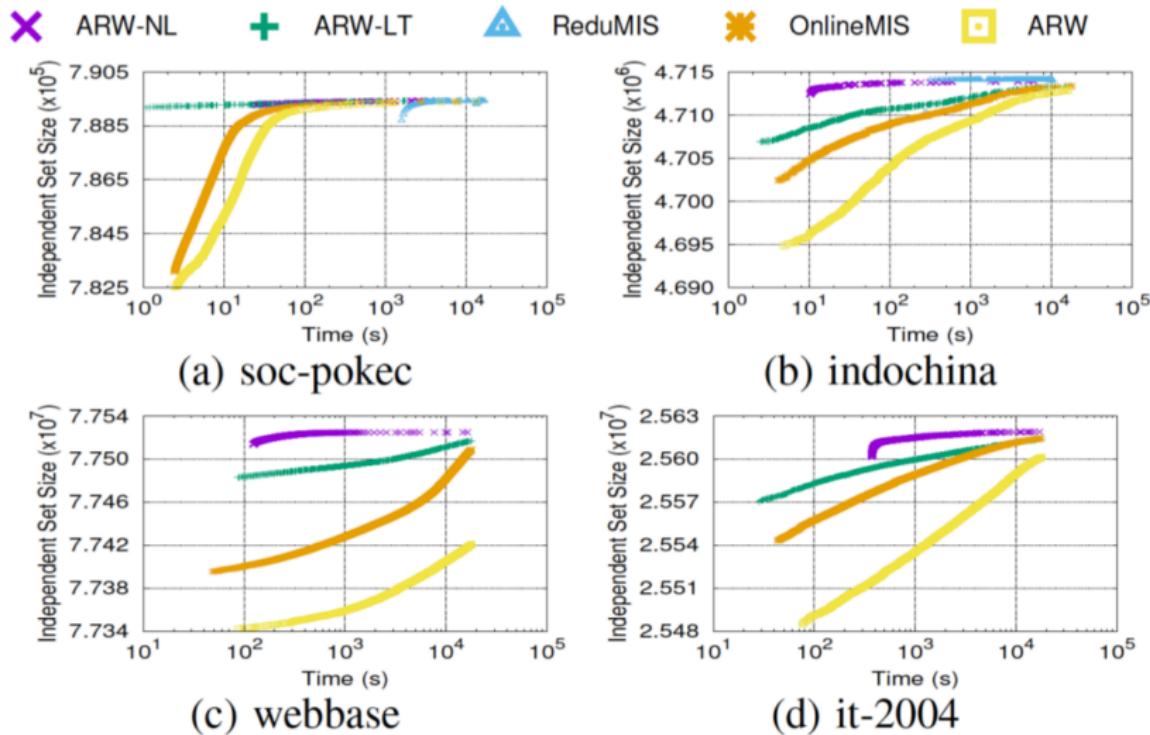
~~ avoid it



**Repeat:** add small degree vertices to solution + reduce

# Linear Time Reductions

[Chang et al.'17]



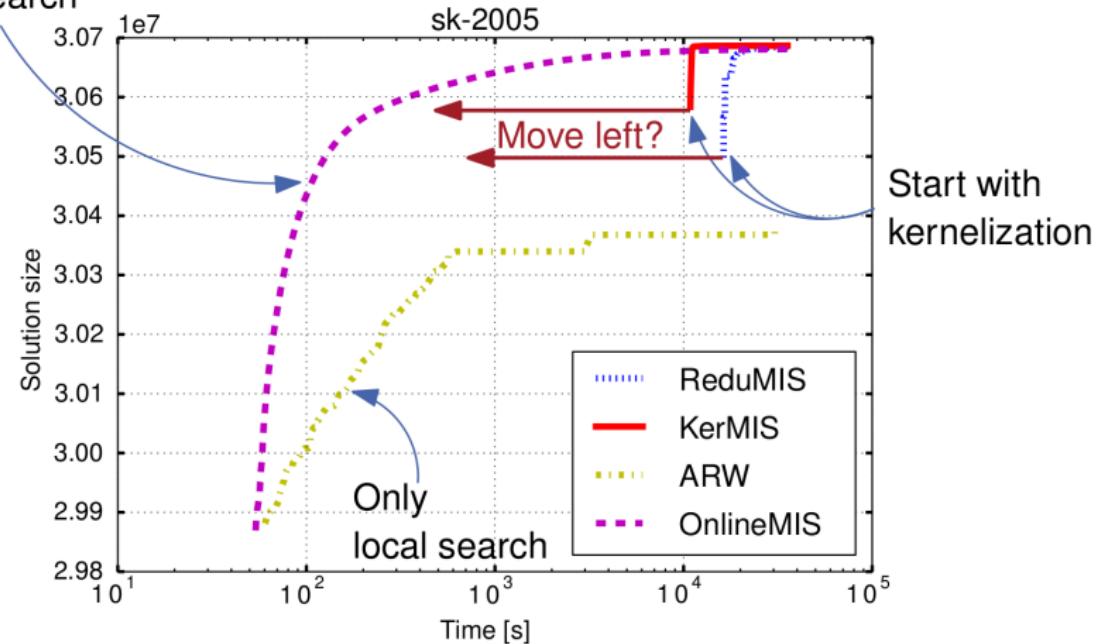
# Scalable Reductions

Effective Reductions Are Slow!

Graph		LinearTime		NearLinear		VCSolver	
name	n	K	time	K	time	K	time
uk-2002	19M	11.7M	1.5	4.0M	28.0	0.2M	336.9
arabic-2005	23M	15.6M	2.6	6.7M	246.1	0.6M	1 033.2
gsh-2015-tpd	31M	2.0M	11.6	1.2M	97.4	0.4M	372.3
uk-2005	39M	28.2M	2.5	5.9M	60.5	0.8M	541.4
it-2004	41M	27.1M	3.3	11.3M	1 544.6	1.6M	6 749.0
sk-2005	51M	*	*	*	*	3.2M	10 010.5
uk-2007-05	106M	*	*	*	*	3.5M	18 829.4
webbase-2001	118M	51.7M	13.0	17.3M	121.1	0.7M	4 207.8
asia.osm	12M	626.7K	0.8	594.4K	1.4	15.2K	204.7
road_usa	24M	2.5M	2.5	2.4M	4.1	0.2M	310.0
europe.osm	51M	1 500.0K	4.1	1 329.9K	6.1	8.4K	302.4
rgg26	67M	67.1M	1.0	51.3M	172.6	49.6M	9 887.7
rhg	100M	*	*	*	*	0	124.0
del24	17M	16.8M	0.2	15.6M	12.7	12.4M	4 789.5
del26	67M	67.1M	0.7	62.5M	53.3	49.9M	20 728.7

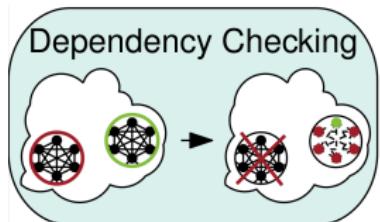
# Even Faster??

Reductions during local search

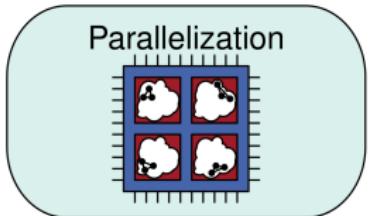


# Remaining Ingredients

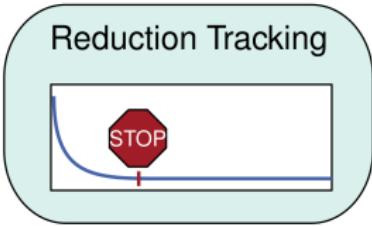
## Dependency Checking



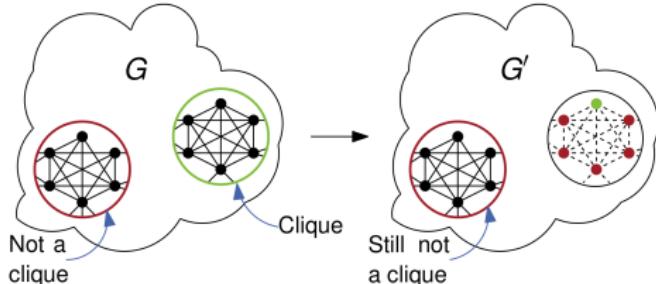
## Parallelization



## Reduction Tracking



## Dependency Checking

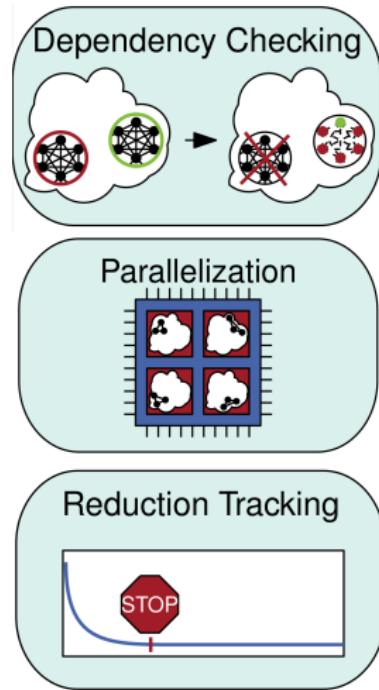


No reduction in  $G$  and  $N_G(v) = N_{G'}(v) \Rightarrow$   
No reduction in  $G'$

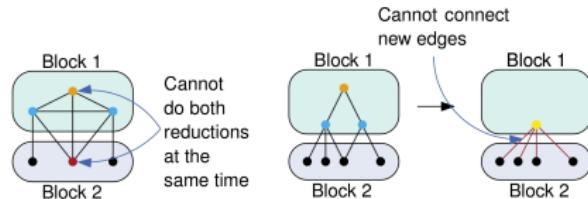
## Used for

- Isolated Clique Reduction
- Degree 2 Fold
- Twin Reduction

# Remaining Ingredients



**Parallelization**  
partition graph in blocks, reduce separately



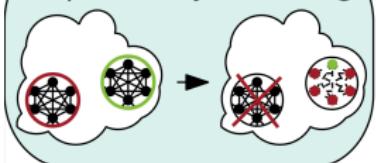
→ deal with conflicts

## Ingredients:

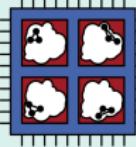
- ParHIP for partitioning
- Parallelize LP reduction with parallel maximum bipartite matching  
[Azad et al.'17]

# Remaining Ingredients

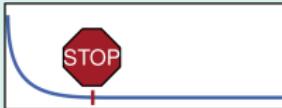
## Dependency Checking



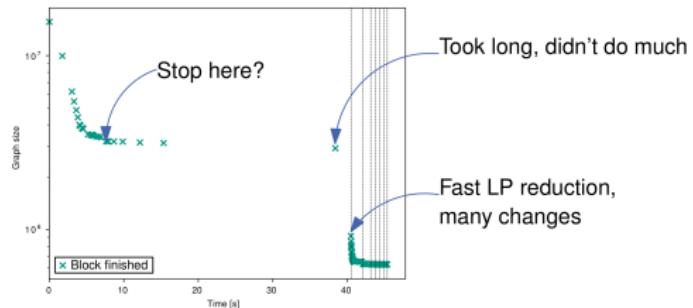
## Parallelization



## Reduction Tracking



## Reduction Tracking stop kernelization early

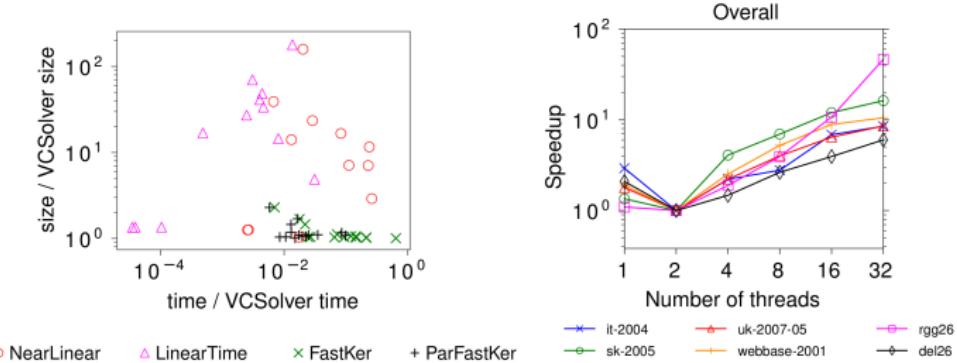


- some blocks take longer than others
- few changes after a while
  - stop “local” reductions early
  - run LP reduction

# Scalable Reductions

NearLinear		VCSSolver		ParFastKer		
$ \mathcal{K} $	time	$ \mathcal{K} $	time	$ \mathcal{K} $	time	su
4.0M	28.0	0.2M	336.9	<b>0.3M</b>	11.8	28.4
6.7M	246.1	0.6M	1 033.2	<b>0.6M</b>	25.7	40.2
1.2M	97.4	0.4M	372.3	<b>0.5M</b>	32.0	11.7
<b>5.9M</b>	<b>60.5</b>	<b>0.8M</b>	<b>541.4</b>	<b>0.9M</b>	<b>53.3</b>	10.1
11.3M	1 544.6	1.6M	6 749.0	<b>1.7M</b>	151.8	44.4
*	*	3.2M	10 010.5	3.5M	178.3	56.1
*	*	3.5M	18 829.4	<b>3.7M</b>	372.4	50.6
<b>17.3M</b>	<b>121.1</b>	<b>0.7M</b>	<b>4 207.8</b>	<b>0.9M</b>	<b>54.9</b>	76.6
594.4K	1.4	15.2K	204.7	<b>34.9K</b>	1.2	169.8
2.4M	4.1	0.2M	310.0	<b>0.2M</b>	4.1	76.0
1 329.9K	6.1	8.4K	302.4	<b>14.2K</b>	4.9	61.3
51.3M	172.6	49.6M	9 887.7	<b>49.8M</b>	150.3	65.8
*	*	0	124.0	<b>16</b>	64.6	1.9
15.6M	12.7	12.4M	4 789.5	12.9M	51.5	93.1
62.5M	53.3	49.9M	20 728.7	51.7M	179.0	115.8

# Experimental Results



- VCSolver [Akiba and Iwata'16]: slow but small kernels
- LinearTime and NearLinear [Chang et al.'17]: fast but large kernels  
→ integrate as preprocessing

## Summary:

- Orders of magnitude smaller than fast methods
- Orders of magnitude faster than algorithm with similar-sized kernels
- Integration into local search (small kernels matter):
  - larger independent sets faster

# PACE Challenge 2019

## Results

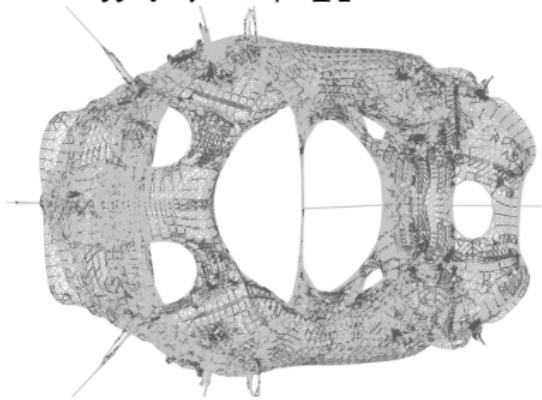
- Exact vertex cover solver  
→based on reductions + local search + portfolio of exact solvers
- **Exact Solvers:**
  - Branch-and-Reduce for vertex cover
  - Clique Solver by Li et al.
  - Get the easy instances **fast**, then get the hard instances.
  - 24 competing algorithms
  - **Results: FIRST place**

# Problem

$$G = (V, E, d)$$
$$d : E \rightarrow \mathbf{R}$$

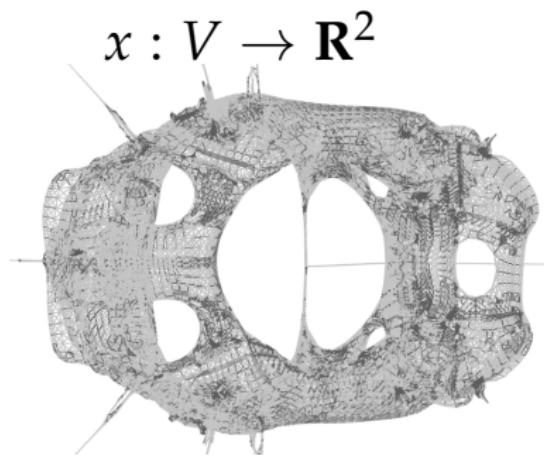
$\Rightarrow$

$$x : V \rightarrow \mathbf{R}^2$$



# Problem

$$\begin{aligned} G &= (V, E, d) \\ d : E &\rightarrow \mathbf{R} \\ d &\equiv 1 \end{aligned} \quad \Rightarrow$$



# Maximal Entropy Stress Model

[Gansner et al.'13]

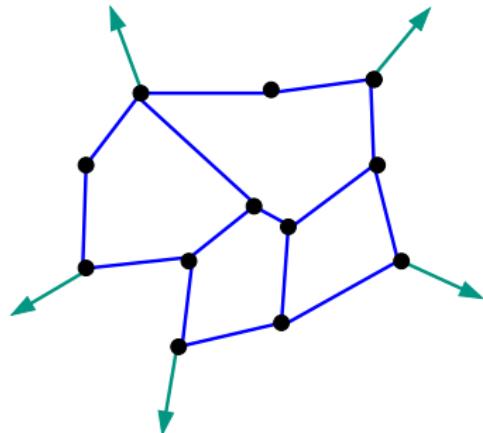
Entropy  $H(x)$ :

- physics: nodes evenly dispersed  
→ nodes as far away as possible
- some nodes have predefined **distance!**

Maximal Entropy Stress Model:

$$\max H(x) := \sum_{\{u,v\} \notin E} \ln \|x_u - x_v\|$$

subject to  $\|x_u - x_v\| = d_{uv}, \{u,v\} \in E$



# Maximal Entropy Stress Model

[Gansner et al.'13]

Entropy  $H(x)$ :

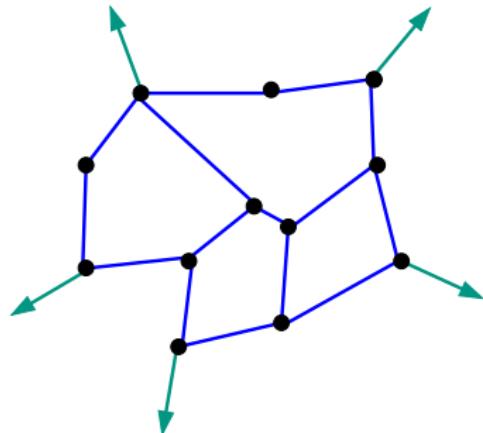
- physics: nodes evenly dispersed  
→ nodes as far away as possible
- some nodes have predefined **distance!**

Maximal Entropy Stress Model:

$$\max H(x) := \sum_{\{u,v\} \notin E} \ln ||x_u - x_v||$$

subject to  $||x_u - x_v|| = d_{uv}, \{u,v\} \in E$

- not possible to satisfy all constraints!  
→ model may be **infeasible**



# Maximal Entropy Stress Model

[Gansner et al.'13]

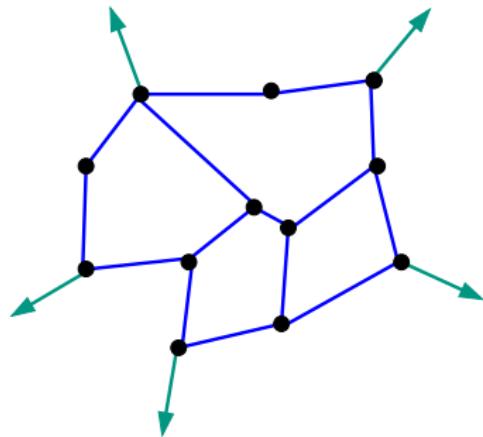
Compromise: min error, max entropy

$$\min \sum_{u,v \in E} w_{uv} (||x_u - x_v|| - d_{uv})^2 - \alpha H(x)$$

$\alpha$  trade-off parameter

Solve optimization problem by

- repeatedly solving Laplacian systems
- or iterative scheme ...



# Maximal Entropy Stress Model

[Gansner et al.'13]

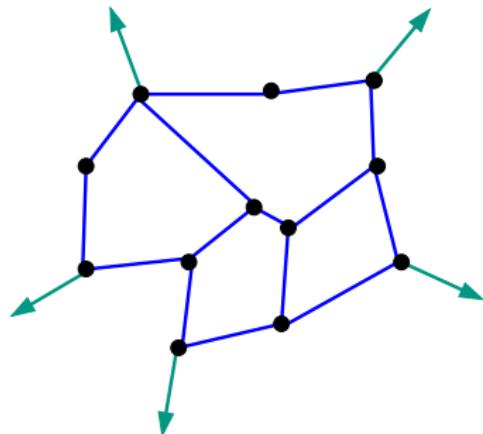
Compromise: min error, max entropy

$$\min \sum_{u,v \in E} w_{uv} (||x_u - x_v|| - d_{uv})^2 - \alpha H(x)$$

$\alpha$  trade-off parameter

Solve optimization problem by

- repeatedly solving Laplacian systems
- or **iterative scheme** ...



# Maximal Entropy Stress Model

[Gansner et al.'13]

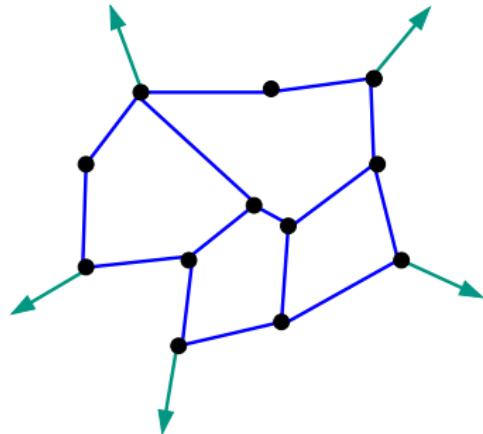
... or iterative scheme:

$$x_u \leftarrow \frac{1}{\rho_u} \sum_{\{u,v\} \in E} w_{uv} \left( x_v + d_{uv} \frac{x_u - x_v}{\|x_u - x_v\|} \right) + \frac{\alpha}{\rho_u} \sum_{\{u,v\} \notin E} \frac{x_u - x_v}{\|x_u - x_v\|^2}$$

→ overall update costs  $O(n \log n)$  per iteration

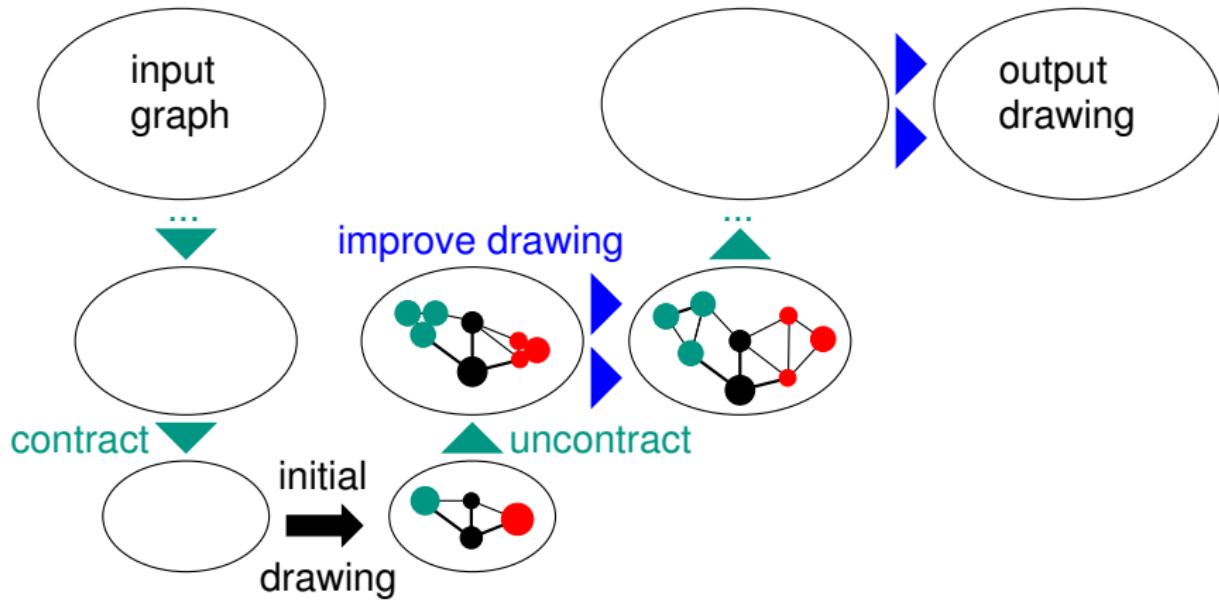
Our contributions:

- make this usable and fast in practice
- multilevel integration
- approximate long-range forces
- employ parallelism



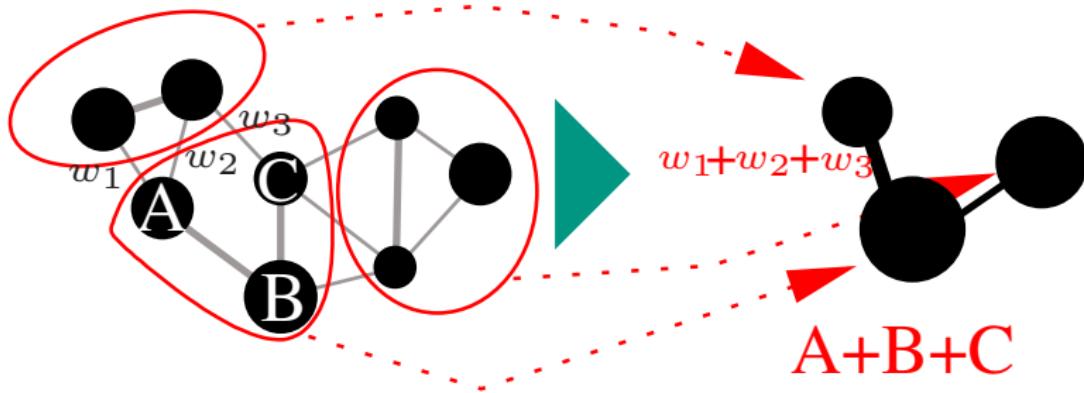
# Multilevel Graph Drawing

[Hadany, Harel'99]



# Basic Idea

## Contraction of Clusterings

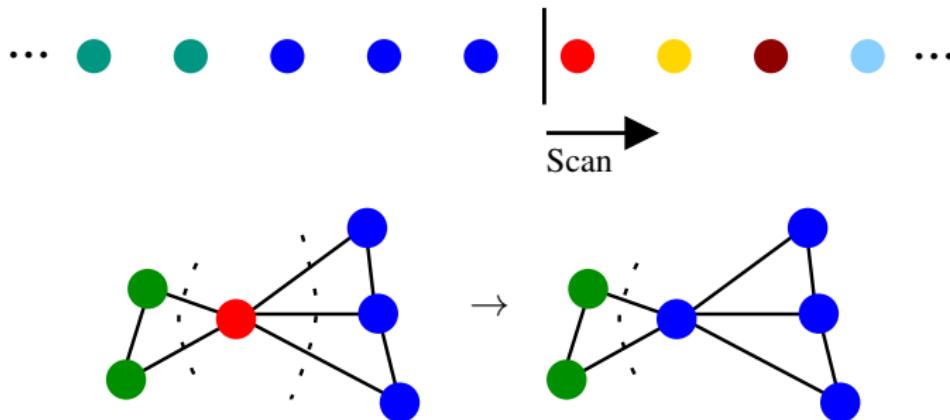


- control block size: size-constraint  $U$
- recurse until graph is small
- $c : V_i \rightarrow \mathbb{R}$  number of nodes on finest level

# Label Propagation

Cut-based, Linear Time Clustering Algorithm [Raghavan et. al]

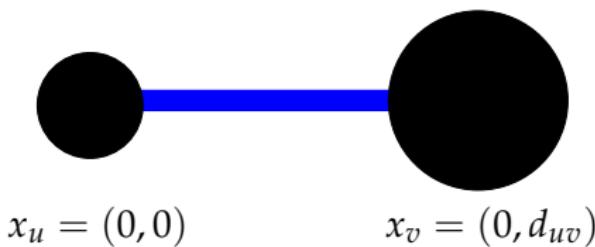
- **cut-based** clustering using size-constraint label propagation
  - start with **singletons**
  - traverse nodes in random order or **smallest degree first**
  - move node to cluster having **strongest eligible** connection
  - modification *eligible*: w.r.t size constraint  $U$



# Multilevel Graph Drawing

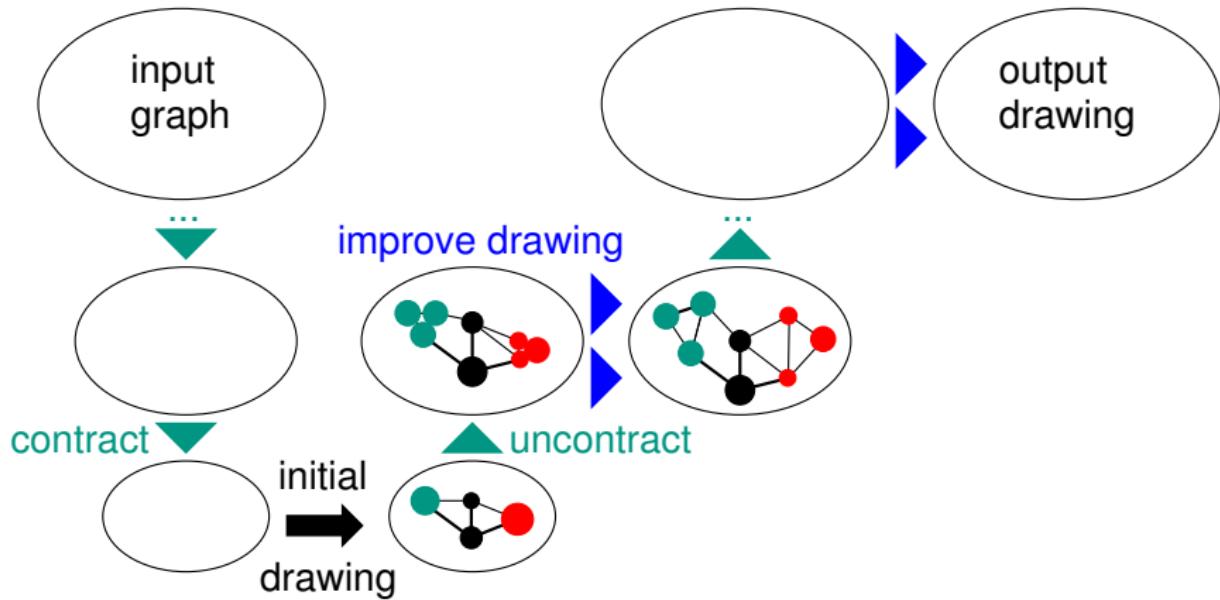
## Initial Drawing

- coarsen until only two nodes left
- place them at optimal distance
- define distances on coarse graphs, stay tuned



# Multilevel Graph Drawing

[Hadany, Harel'99]



# Uncoarsening

## Uncontraction/Transfer

- $v$  fine vertex of coarse representative  $v'$  at  $P = (x, y)$
- put  $v$  at random position in disk around  $P$  with radius  $r = \sqrt{c(v')}$

## Local Improvement

- minimize maxent-stress on each level of hierarchy
- assume disk with radius  $\sqrt{c(u)}$  to draw  $c(u)$  vertices  
→ define distance  $d_{uv} := \frac{\sqrt{c(u)} + \sqrt{c(v)}}{2}$  on current level

## Iterative scheme

$$x_u \leftarrow \frac{1}{\rho_u} \sum_{\{u,v\} \in E} w_{uv} \left( x_v + d_{uv} \frac{x_u - x_v}{\|x_u - x_v\|} \right) + \frac{\alpha}{\rho_u} \sum_{\{u,v\} \notin E} \frac{x_u - x_v}{\|x_u - x_v\|^2}$$

# Uncoarsening

## Uncontraction/Transfer

- $v$  fine vertex of coarse representative  $v'$  at  $P = (x, y)$
- put  $v$  at random position in disk around  $P$  with radius  $r = \sqrt{c(v')}$

## Local Improvement

- minimize maxent-stress on each level of hierarchy
- assume disk with radius  $\sqrt{c(u)}$  to draw  $c(u)$  vertices  
→ define distance  $d_{uv} := \frac{\sqrt{c(u)} + \sqrt{c(v)}}{2}$  on current level

## Iterative scheme

$$x_u \leftarrow \frac{1}{\rho_u} \sum_{\{u,v\} \in E} w_{uv} \left( x_v + d_{uv} \frac{x_u - x_v}{\|x_u - x_v\|} \right) + \frac{\alpha}{\rho_u} \sum_{\{u,v\} \notin E} \frac{x_u - x_v}{\|x_u - x_v\|^2}$$

# Local Improvement

## Iterative scheme

$$x_u \leftarrow \dots \sum_{\{u,v\} \notin E} \underbrace{\frac{x_u - x_v}{\|x_u - x_v\|^2}}_{=:r(u,v)}$$

## Approximation

$$x_u \leftarrow \dots \sum_{\substack{u \neq v \\ M(u) = M(v)}} r(u, v) + \sum_{\substack{v' \in V' \\ v' \neq M(u)}} \nu(v') \frac{x_u - x'_{v'}}{\|x_u - x'_{v'}\|^2} - \sum_{\{u,v\} \in E} r(u, v)$$

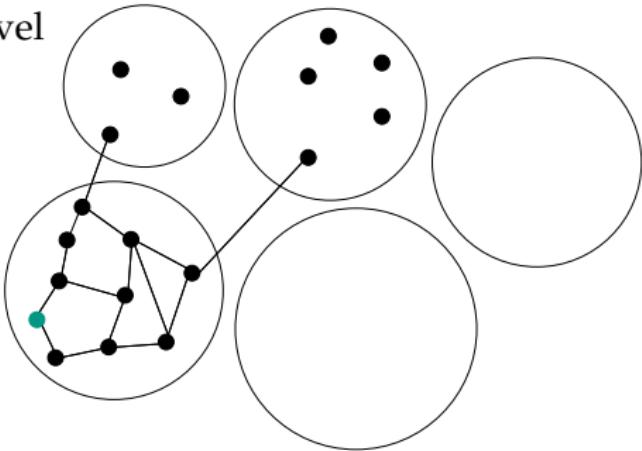
- $M(u)$  cluster of  $u$
- $\nu(v')$  number of finer vertices of  $v'$  on current level
- $V'$  vertex set of next coarser level
- $x'_{v'}$  coordinate of  $v'$

# Local Improvement

## Approximation

$$x_u \leftarrow \cdots + \sum_{\substack{u \neq v \\ M(u) = M(v)}} r(u, v) + \sum_{\substack{v' \in V' \\ v' \neq M(u)}} v(v') \frac{x_u - x'_{v'}}{\|x_u - x'_{v'}\|^2} - \sum_{\{u, v\} \in E} r(u, v)$$

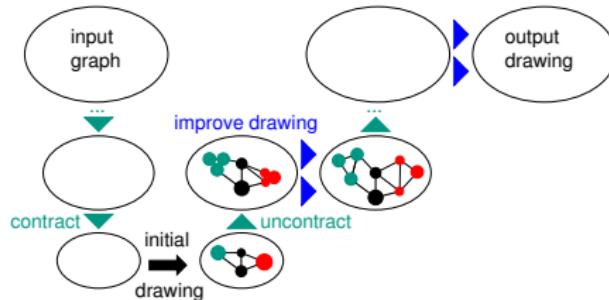
- $M(u)$  cluster of  $u$
- $v(v')$  number of finer vertices of  $v'$  on current level
- $V'$  vertex set of next coarser level
- $x'_{v'}$  coordinate of  $v'$



# Local Improvement

## Additional Enhancements

- after each iteration → update barycenter of coarse nodes
- vertex computations independent → add parallelism
- use approximation multiple –  $h$  – levels beneath in hierarchy

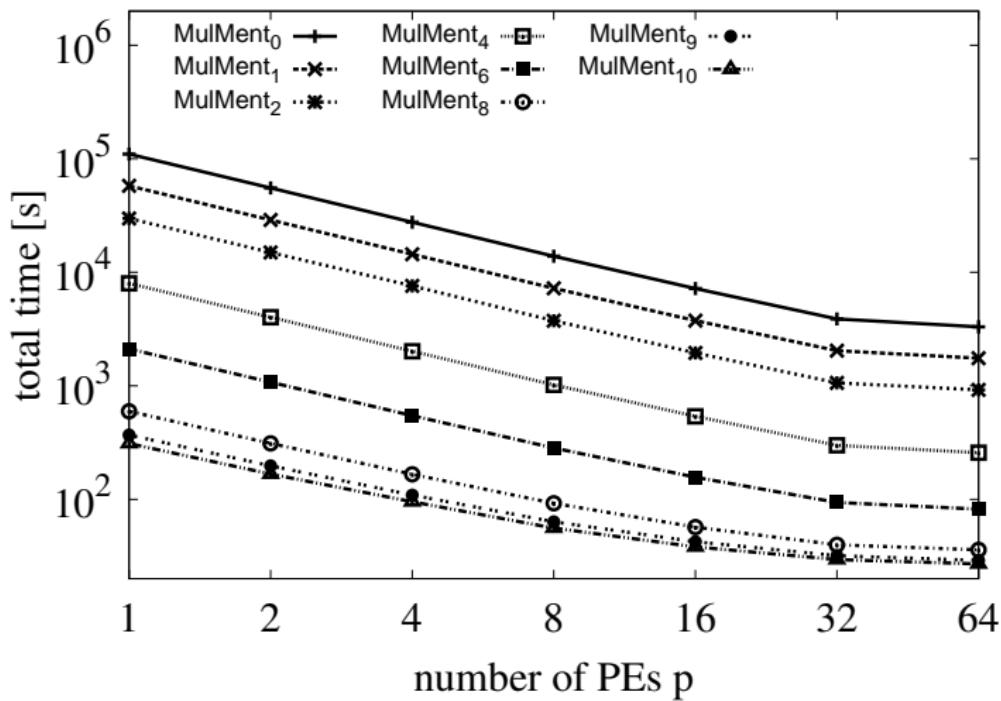


**Proposition:** Assume equal cluster sizes. The running time of one iteration of  $\text{MulMent}_h$ ,  $h \geq 0$ , is  $\mathcal{O}(m + n^{\frac{h+2}{h+1}})$

# Experimental Results

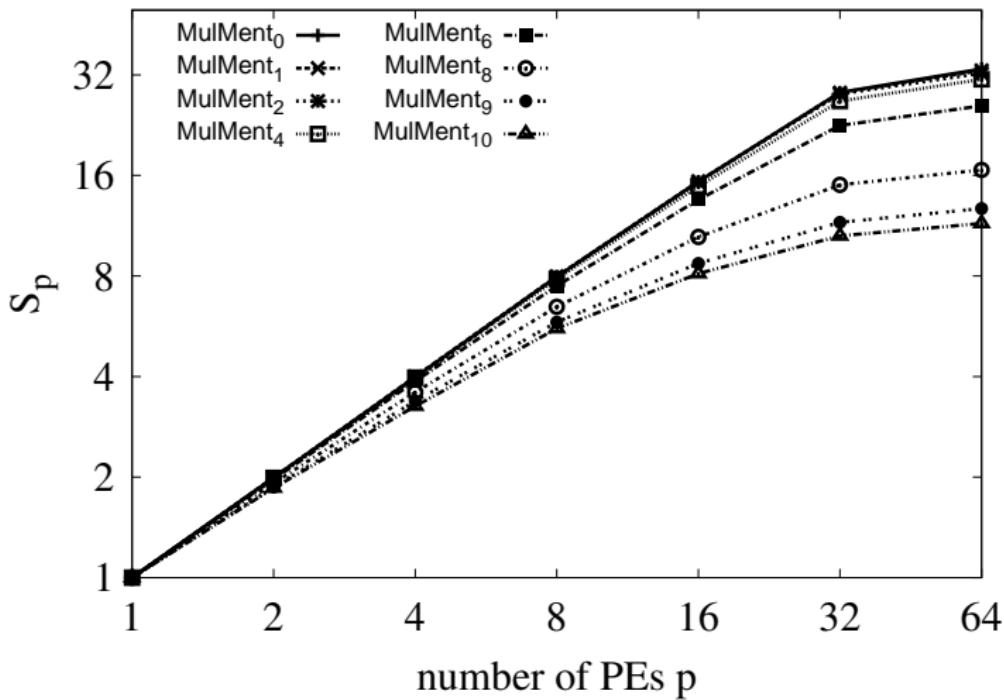
# Scalability

Running Time [Delaunay  $n = 2^{20}$ ]



# Scalability

Speedup [Delaunay  $n = 2^{20}$ ]



# Running Times

graph	PMDS	MaxEnt	MulMent <sub>0</sub>	MulMent <sub>1</sub>	MulMent <sub>10</sub>
btree	0.02	<b>1.14</b>	0.10	0.17	<b>0.11</b>
1138_bus	0.04	<b>1.41</b>	0.15	0.13	<b>0.11</b>
USpowerG	0.14	<b>3.82</b>	0.29	0.23	<b>0.18</b>
3elt	0.16	<b>3.45</b>	0.21	0.19	<b>0.17</b>
commanche	0.24	<b>5.42</b>	0.32	0.27	<b>0.18</b>
bcsstk31	3.44	<b>48.48</b>	5.63	3.97	<b>1.82</b>
fe_pwt	1.49	<b>31.60</b>	4.46	2.67	<b>0.66</b>
del16	2.86	<b>61.42</b>	13.63	7.75	<b>1.10</b>
luxembourg	3.10	<b>96.10</b>	40.94	22.87	<b>1.39</b>
nyc	9.03	<b>233.94</b>	216.27	119.33	<b>3.70</b>
auto	41.80	<b>665.67</b>	613.51	329.08	<b>20.70</b>
del20	53.80	<b>1125.03</b>	3303.82	1749.77	<b>27.01</b>

Table: Running times in seconds per graph. Smaller is better. PivotMDS and MaxEnt use one thread (sequential codes), the MulMent<sub>\*</sub> algorithms use 32 cores (64 threads). Running times of MaxEnt are without the time of PMDS (which yields input coordinates to MaxEnt)

# Experimental Results

## Summary

### Influence of $h$

- increasing  $h$  not a large impact on solution quality
- maxent-stress remains comparable
- slight increase in FSM on 2 instances

### Comparison

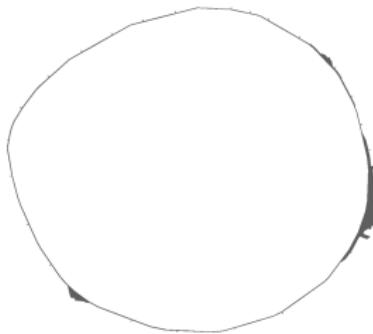
- maxent-stress of MaxEnt and MulMent more or less similar
- four out of nine full stress comparable to MaxEnt
- largest three instances worse full stress (not astonishing!)

### Dynamic Networks

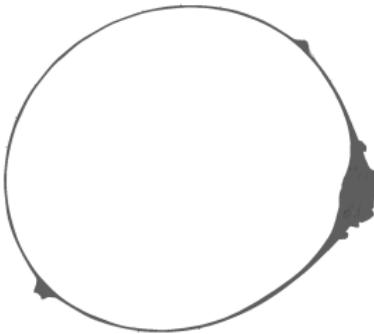
- $\approx$ model: remove  $x\%$  random edges, insert  $x\%$  edges (distance  $\leq \mathcal{D}$ )
- $4\times$  faster ( $h = 0$ ), save 50% time ( $h = 7$ )
- 9% worse full-stress, 1% better maxent-stress

# Example Drawings

**fe\_pwt**



**PMDS**



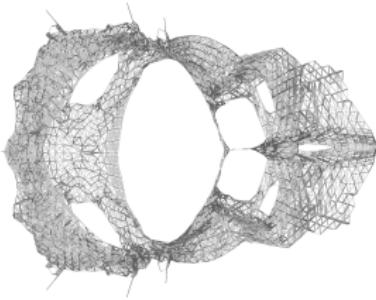
**MaxEnt**



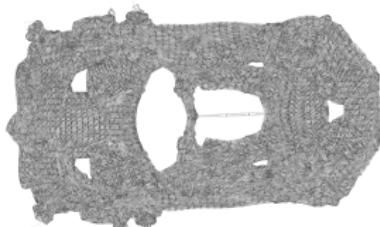
**MulMent**

# Example Drawings

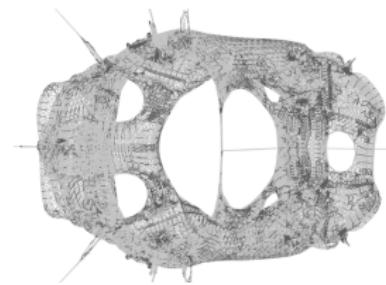
bcsstk31



PMDS



MaxEnt



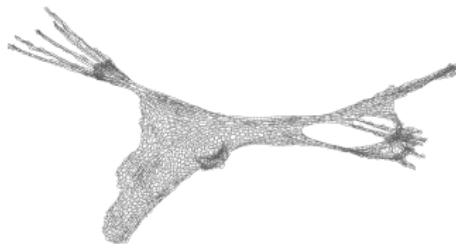
MulMent

# Example Drawings

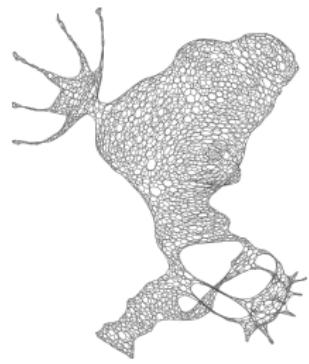
commanche



PMDS



MaxEnt



MulMent